

# HOST-HOST Communication, 1970

- Paul Barran, [On Distributed Communication Networks](#), IEEE Transactions on Communication Systems, Volume 12, No. 1, March 1964
- Introduced Concepts of
  - Distributed Networks
  - Routing  
(hot-potato-routing)
  - Packet-Switching  
(message-block)
- Name worth mentioning: Bob Kahn

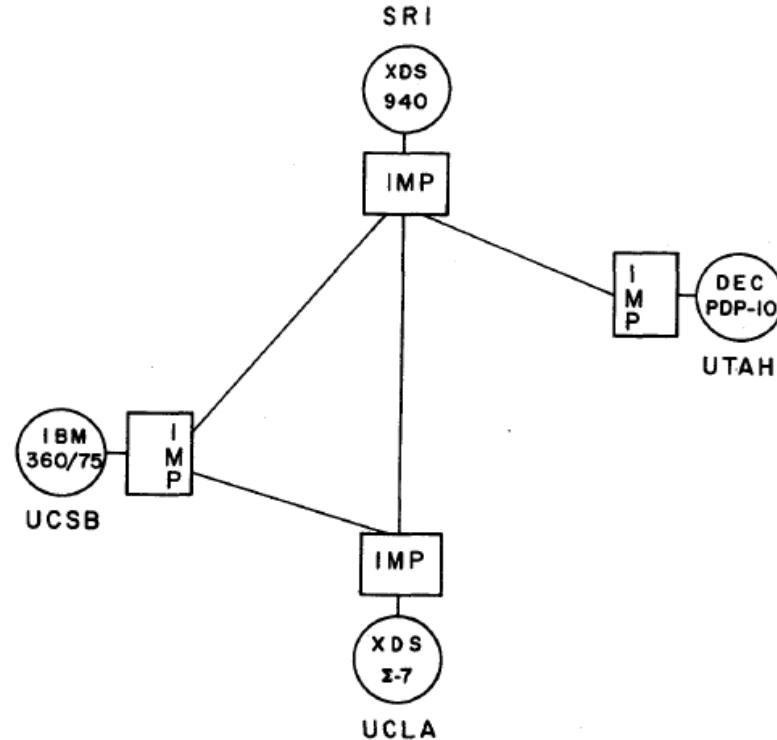
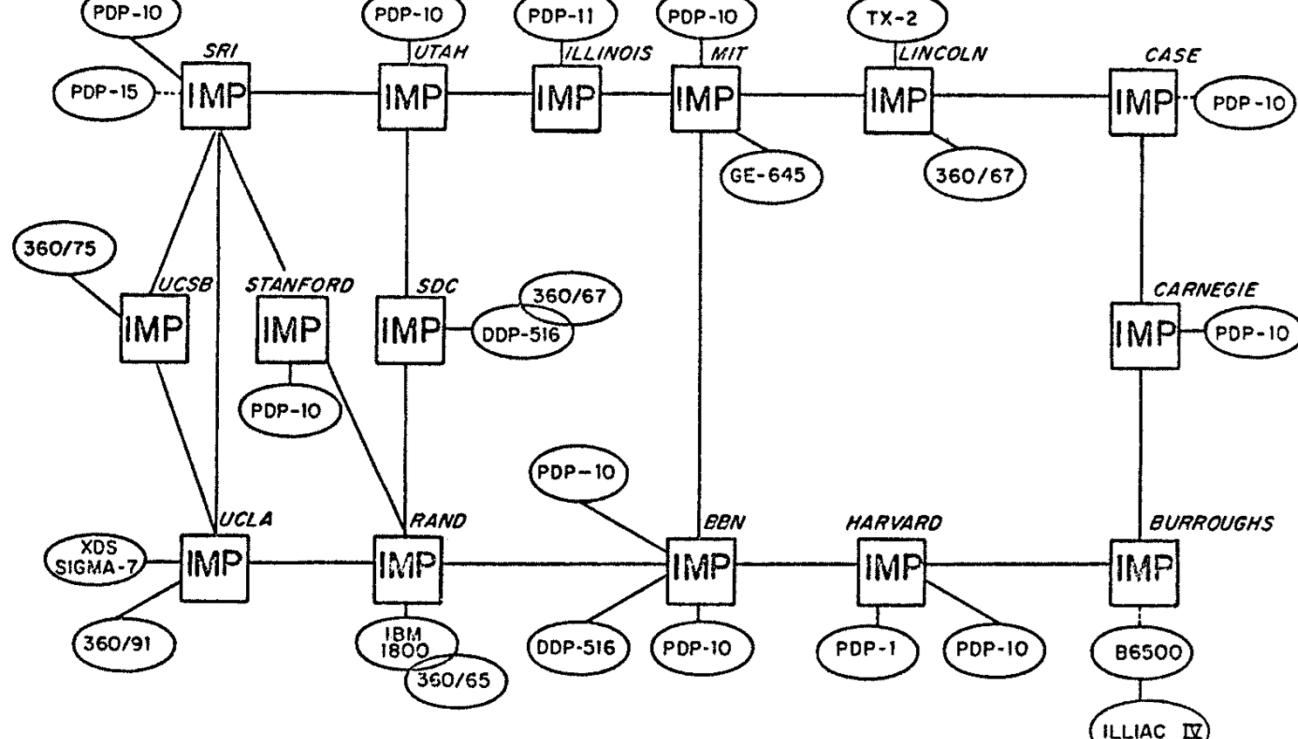


Figure 1—Initial network configuration

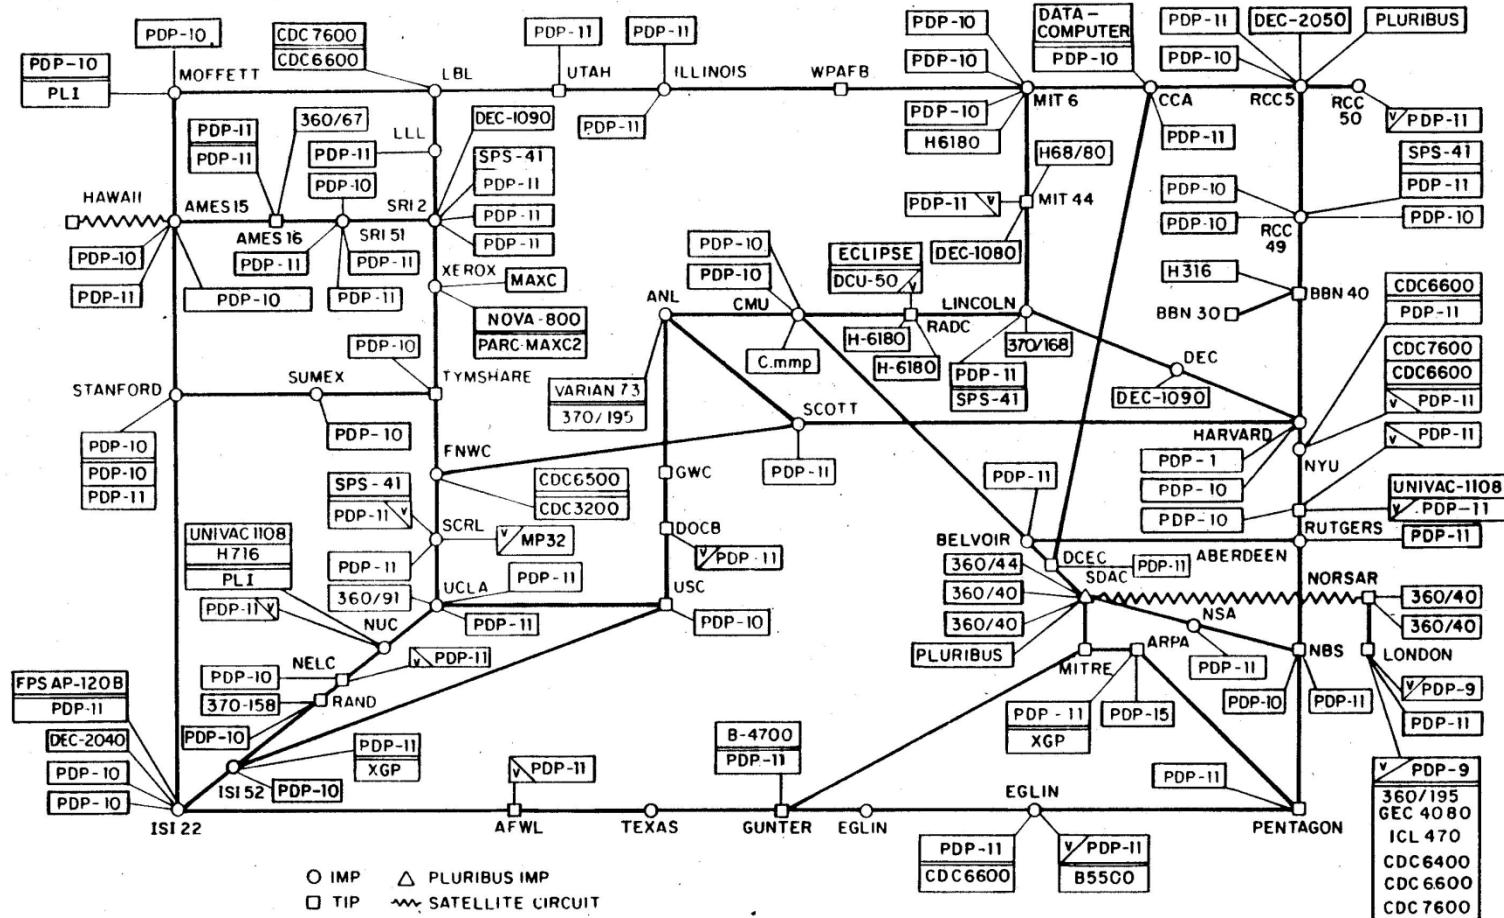
\* C. Stephen Carr, Stephen D. Crocker, and Vinton G. Cerf, [HOST-HOST Communication Protocol in the ARPA Network](#), Spring Joint Computer Conference, pages 589-597, Atlantic City, NJ, USA, May 1970

AT&T's comment was that packet-switching will never be useful.

# ARPANET 1971



# ARPANET 1977



# IP Addresses

10000000 00001011 00000011 00011111

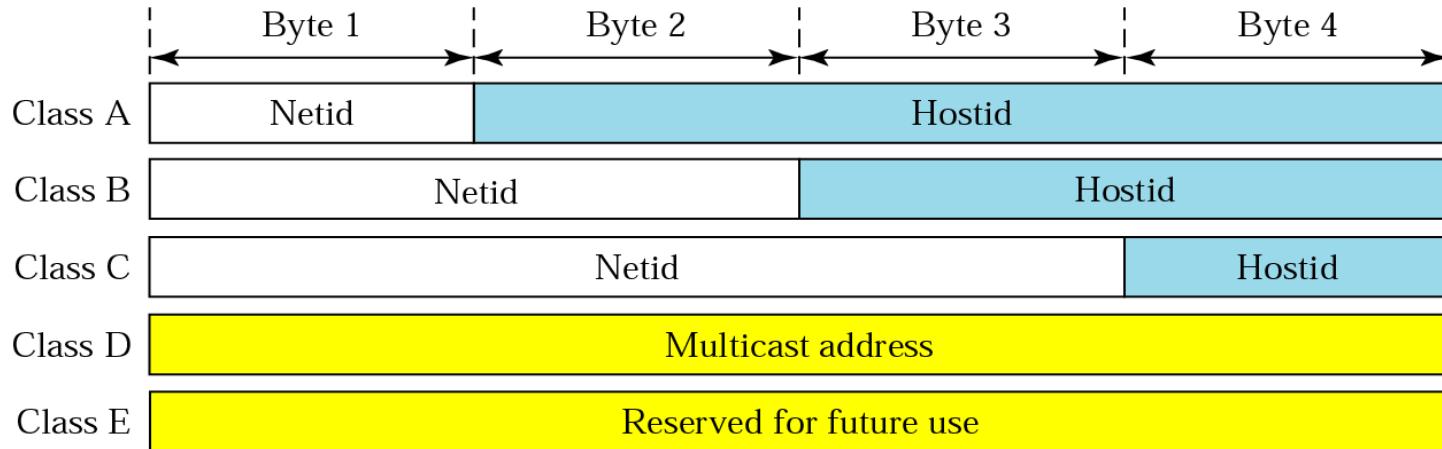
128.11.3.31

- 32-bit number
  - 4.294.967.296 addresses
- IP addresses are unique and universal
  - except private addresses

Range	Total
10.0.0.0 to 10.255.255.255	$2^{24}$
172.16.0.0 to 172.31.255.255	$2^{20}$
192.168.0.0 to 192.168.255.255	$2^{16}$

- Dotted decimal notation:
  - Bytes of binary notation represented as decimal separated by dot

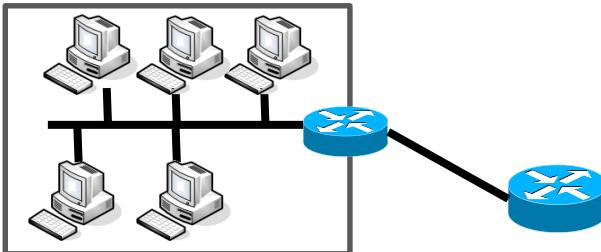
# Classful Addresses



- Class A (international organisations)
  - 126 networks with 16,277,214 hosts each
- Class B (large companies)
  - 16,384 networks with 65,354 hosts each
- Class C (smaller companies)
  - 2,097,152 networks with 254 hosts each
- Inefficient use of hierarchical address space
  - Class C with 2 hosts ( $2/254 = 0.78\%$  efficient)
  - Class B with 256 hosts ( $256/65534 = 0.39\%$  efficient)

# Network Layer

Trinity College



134.226.0.0

← Class B



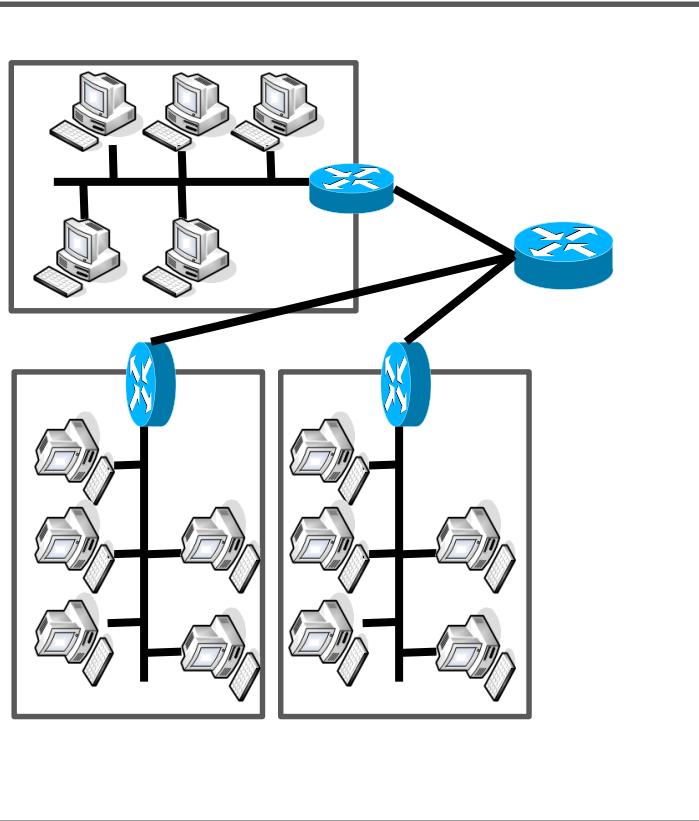
TRINITY COLLEGE DUBLIN  
COLÁISTE NA TRÍONÓIDE, BAILE ÁTHA CLIATH

THE  
UNIVERSITY  
OF DUBLIN

CS2031 Telecommunications II

# Network Layer

Trinity College



134.226.0.0

←  
Class B



TRINITY COLLEGE DUBLIN  
COLÁISTE NA TRÍONÓIDE, BAILE ÁTHA CLIATH

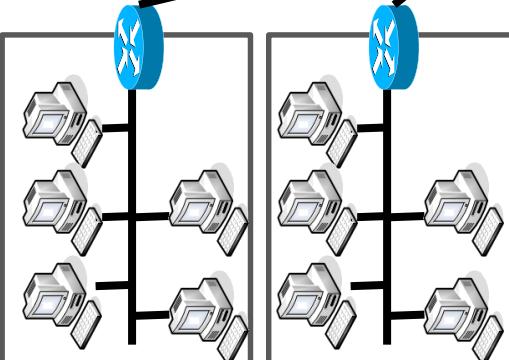
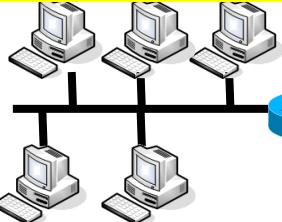
THE  
UNIVERSITY  
OF DUBLIN

CS2031 Telecommunications II

# Network Layer

Trinity College

134.226.36.0



134.226.52.0

134.226.120.0

134.226.0.0

Class B ← Classful Addresses

CS2031 Telecommunications II

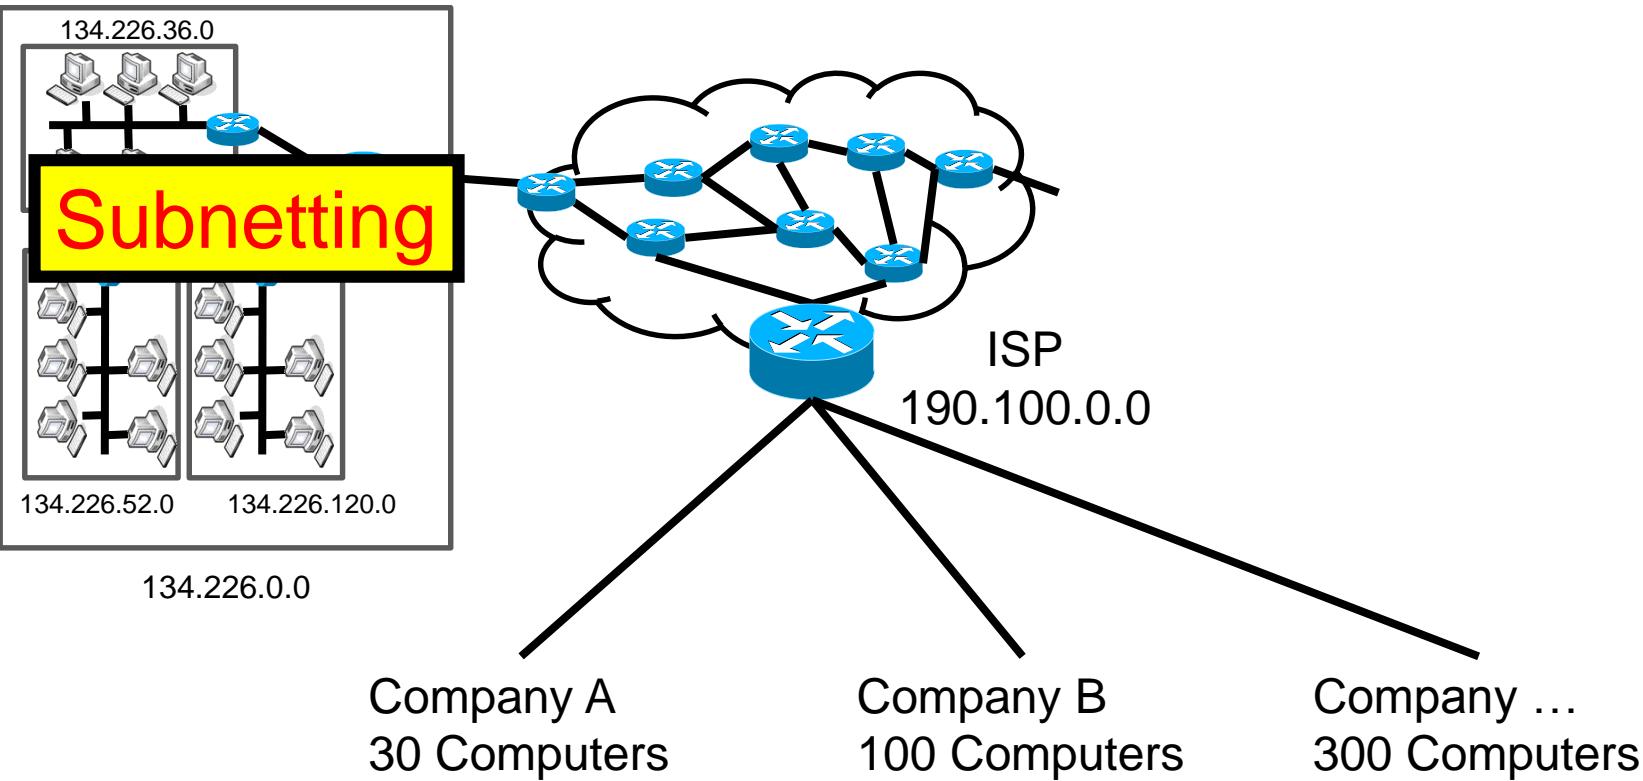


TRINITY COLLEGE DUBLIN  
COLÁISTE NA TRÍONÓIDE, BAILE ÁTHA CLIATH

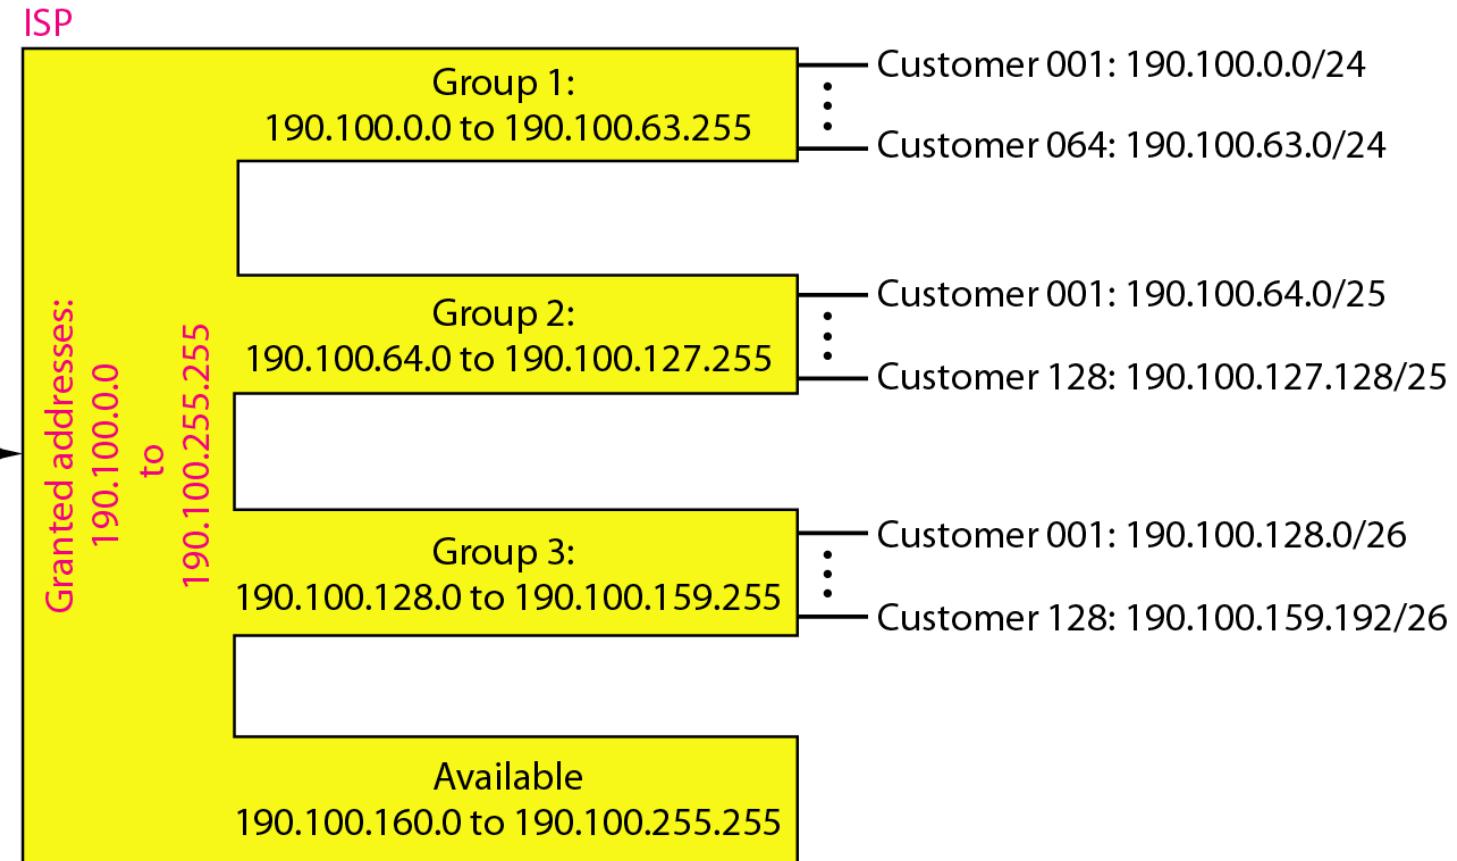
THE  
UNIVERSITY  
OF DUBLIN

# Network Layer

Trinity College



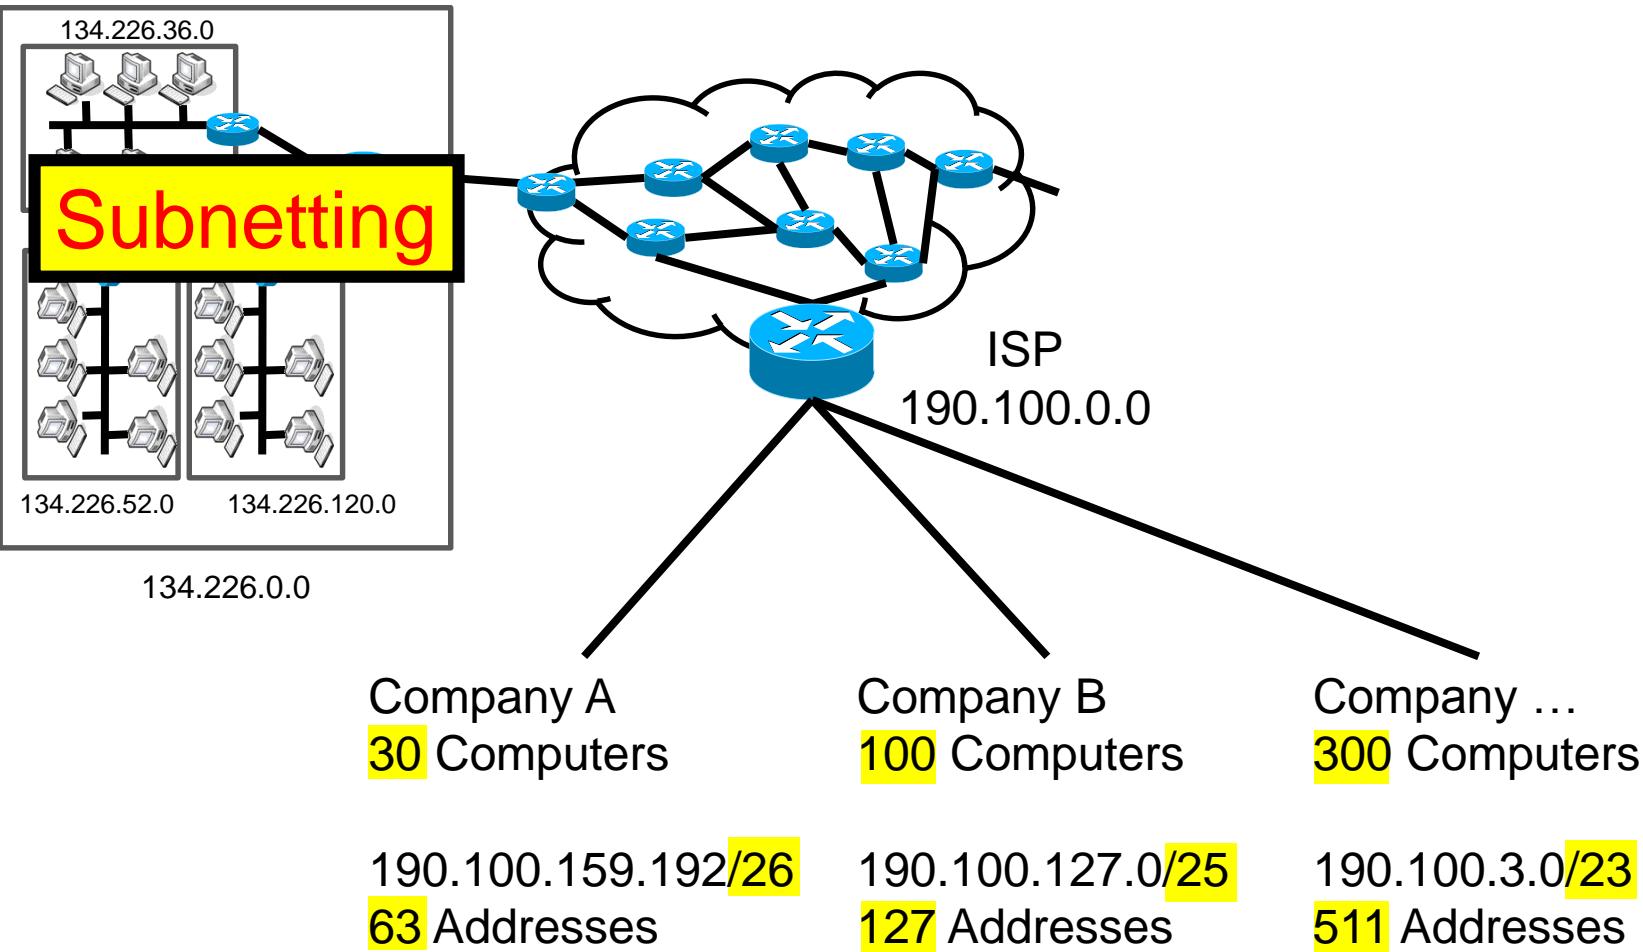
# ISPs & Classless Addresses



\* Figure is courtesy of B. Forouzan

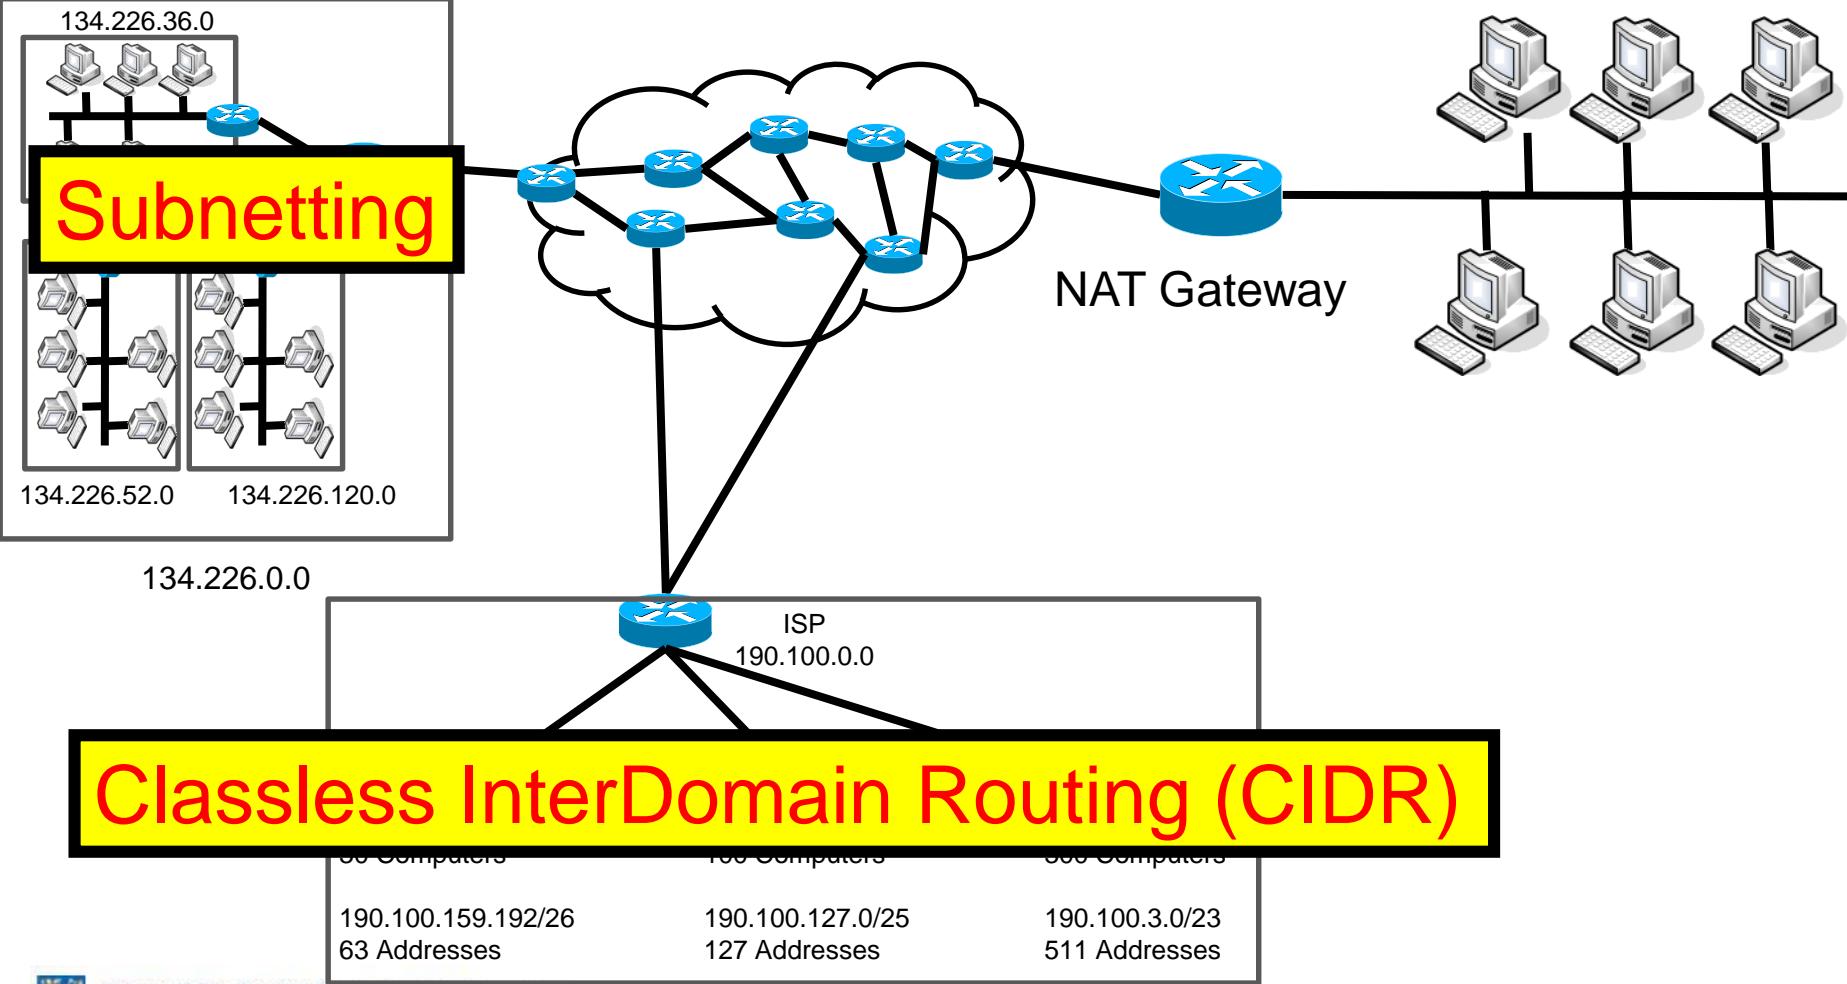
# Network Layer

Trinity College



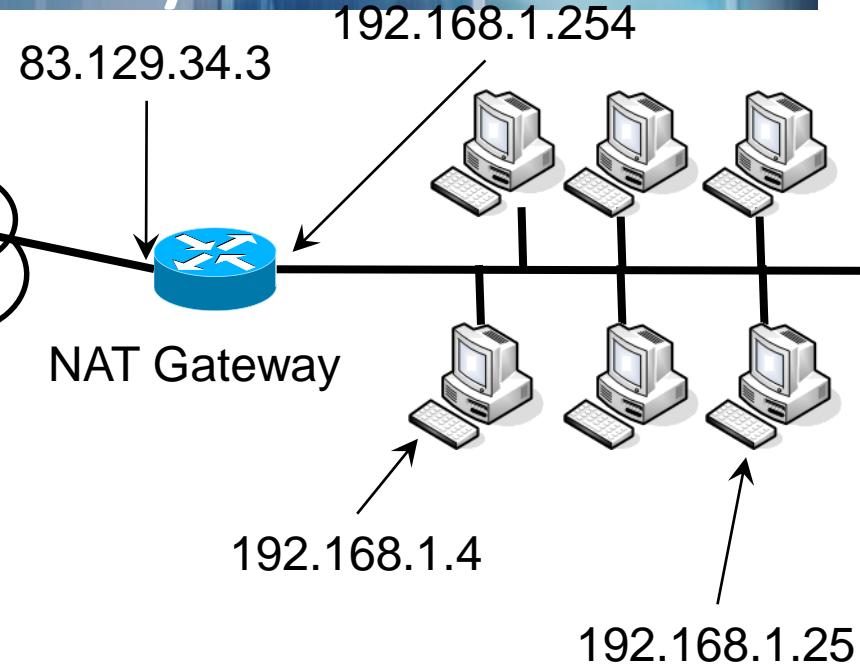
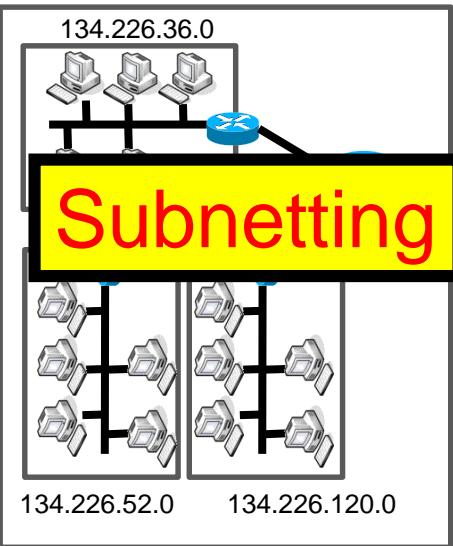
# Network Layer

Trinity College



# Network Layer

Trinity College



Classless InterDomain Routing (CIDR)

60 Computers

190.100.159.192/26  
63 Addresses

100 Computers

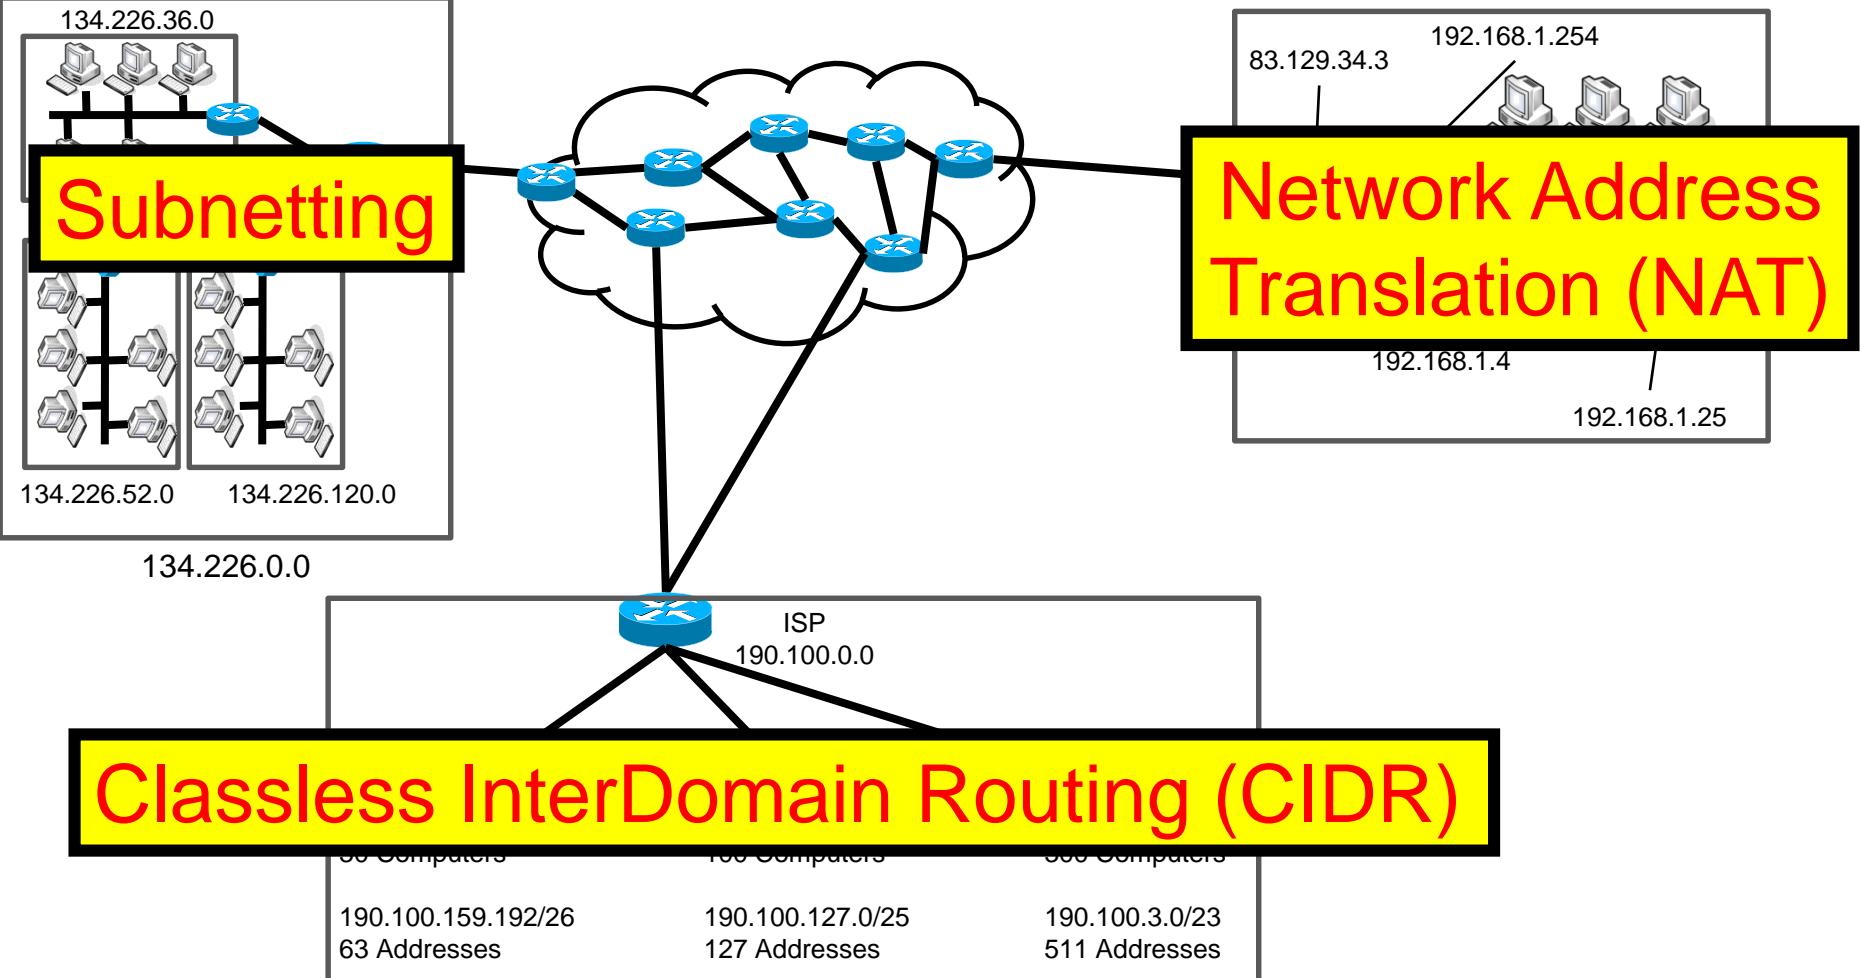
190.100.127.0/25  
127 Addresses

600 Computers

190.100.3.0/23  
511 Addresses



# Network Layer





# CS2031

## Telecommunications II

### Routing

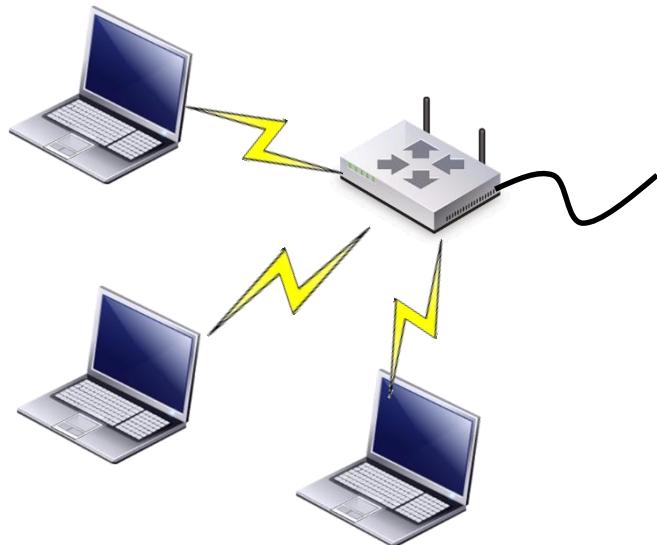


# Outcome for Today

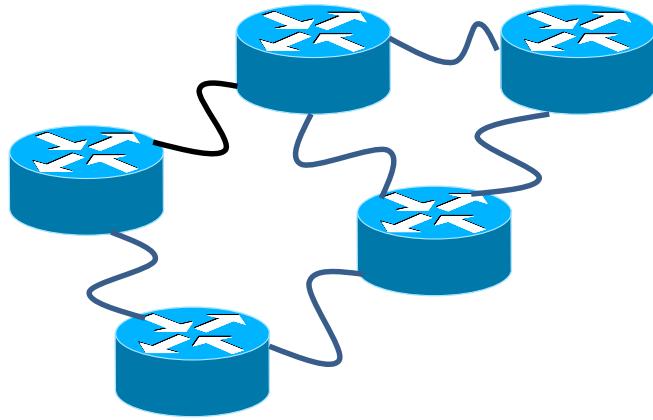
- Distance Vector Routing
- Link State Routing
- Multicast Routing
  - Dense Mode (DM)
  - Sparse Mode (SM)



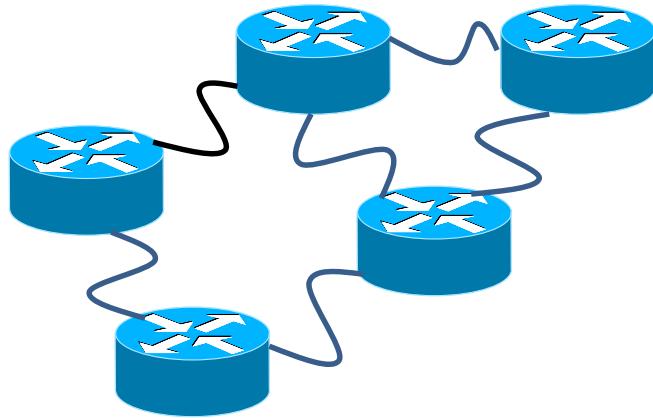
# Home Network



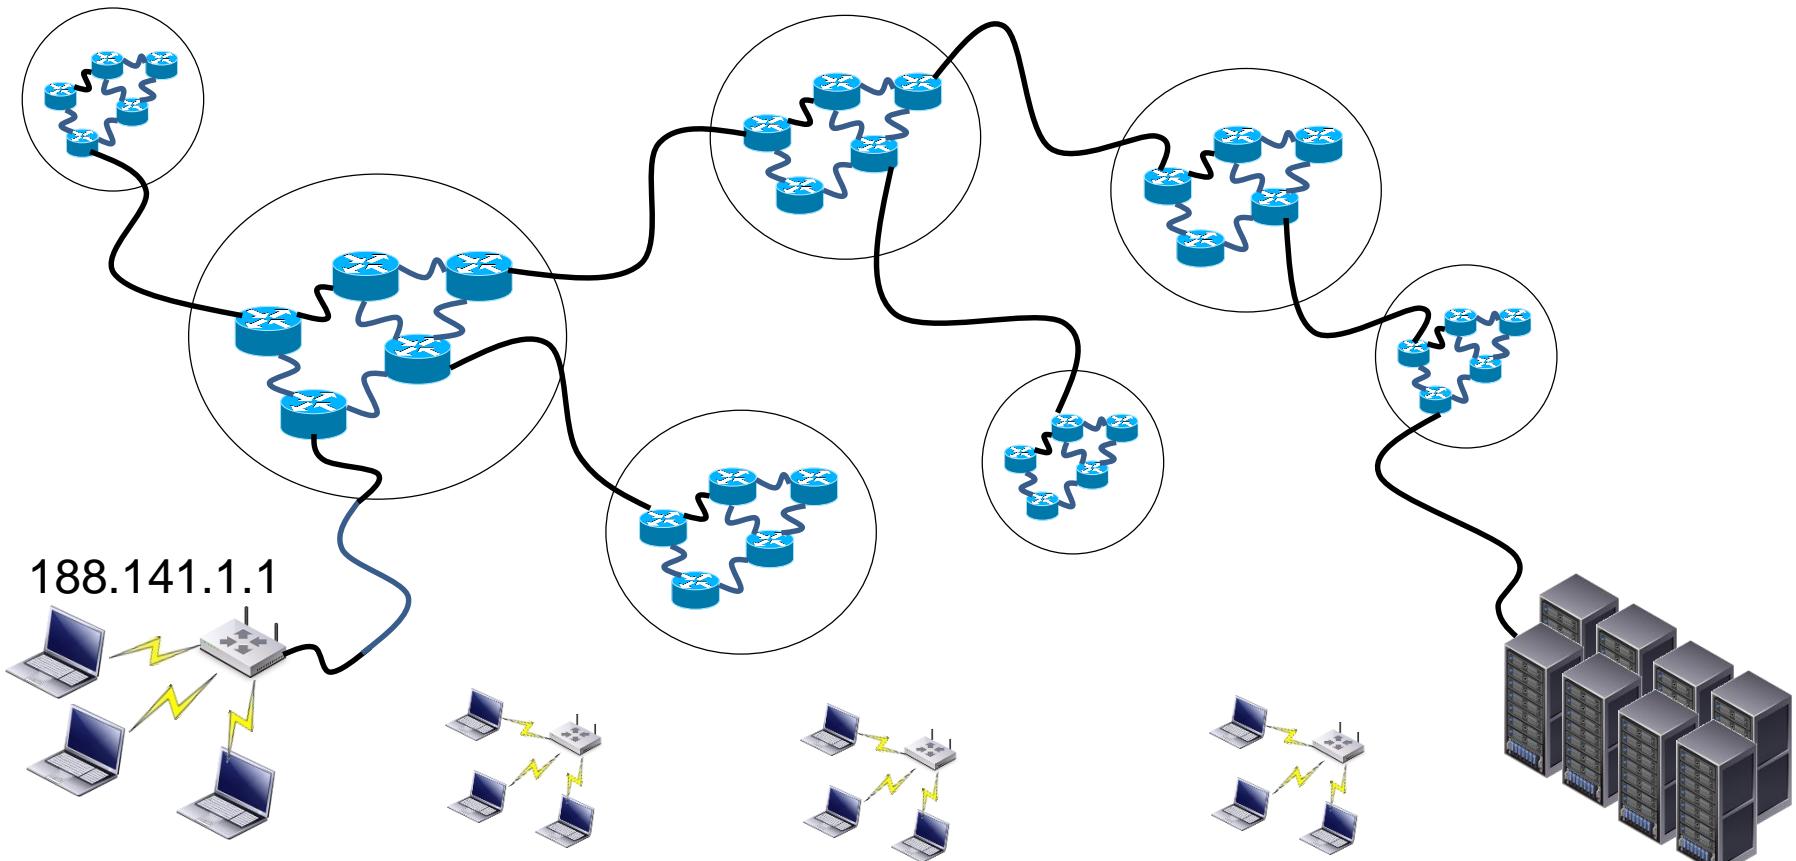
# ISPs



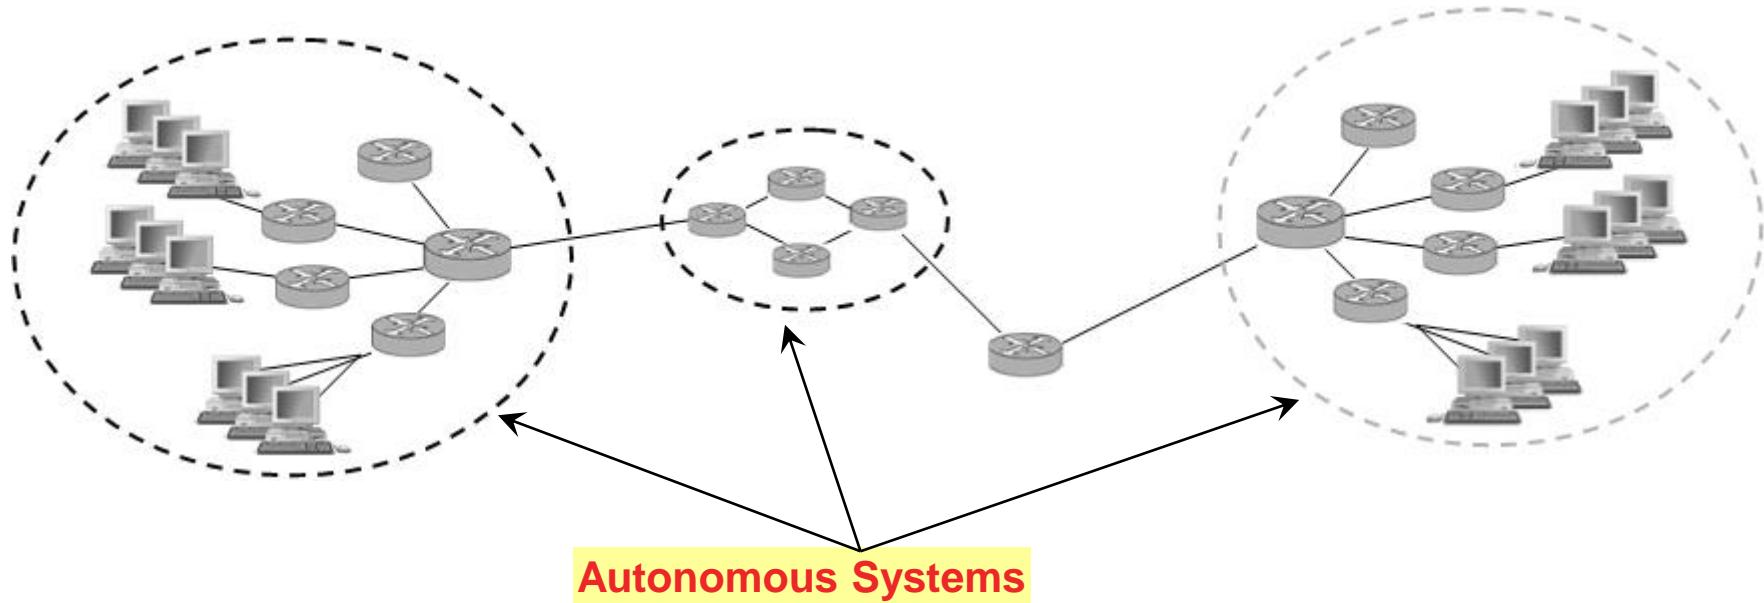
# ISPs



# Internet = Network of Networks

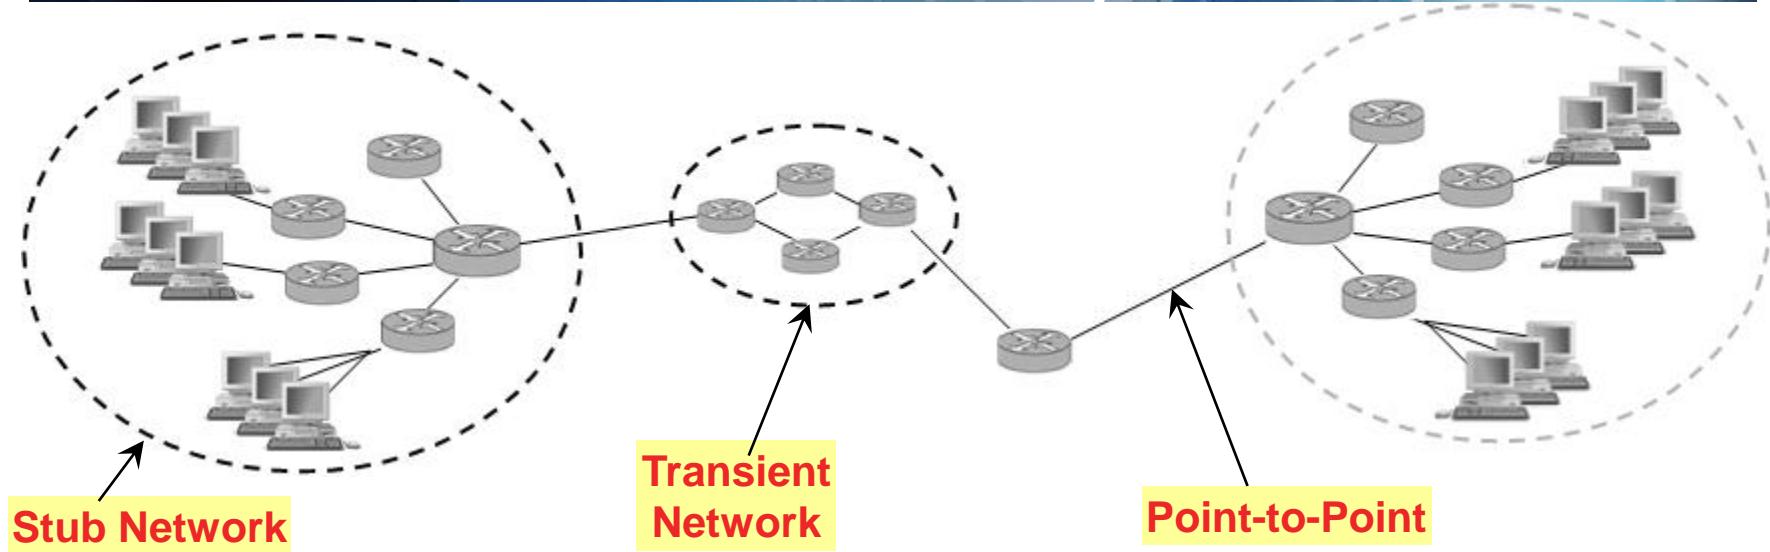


# Structure of the Internet



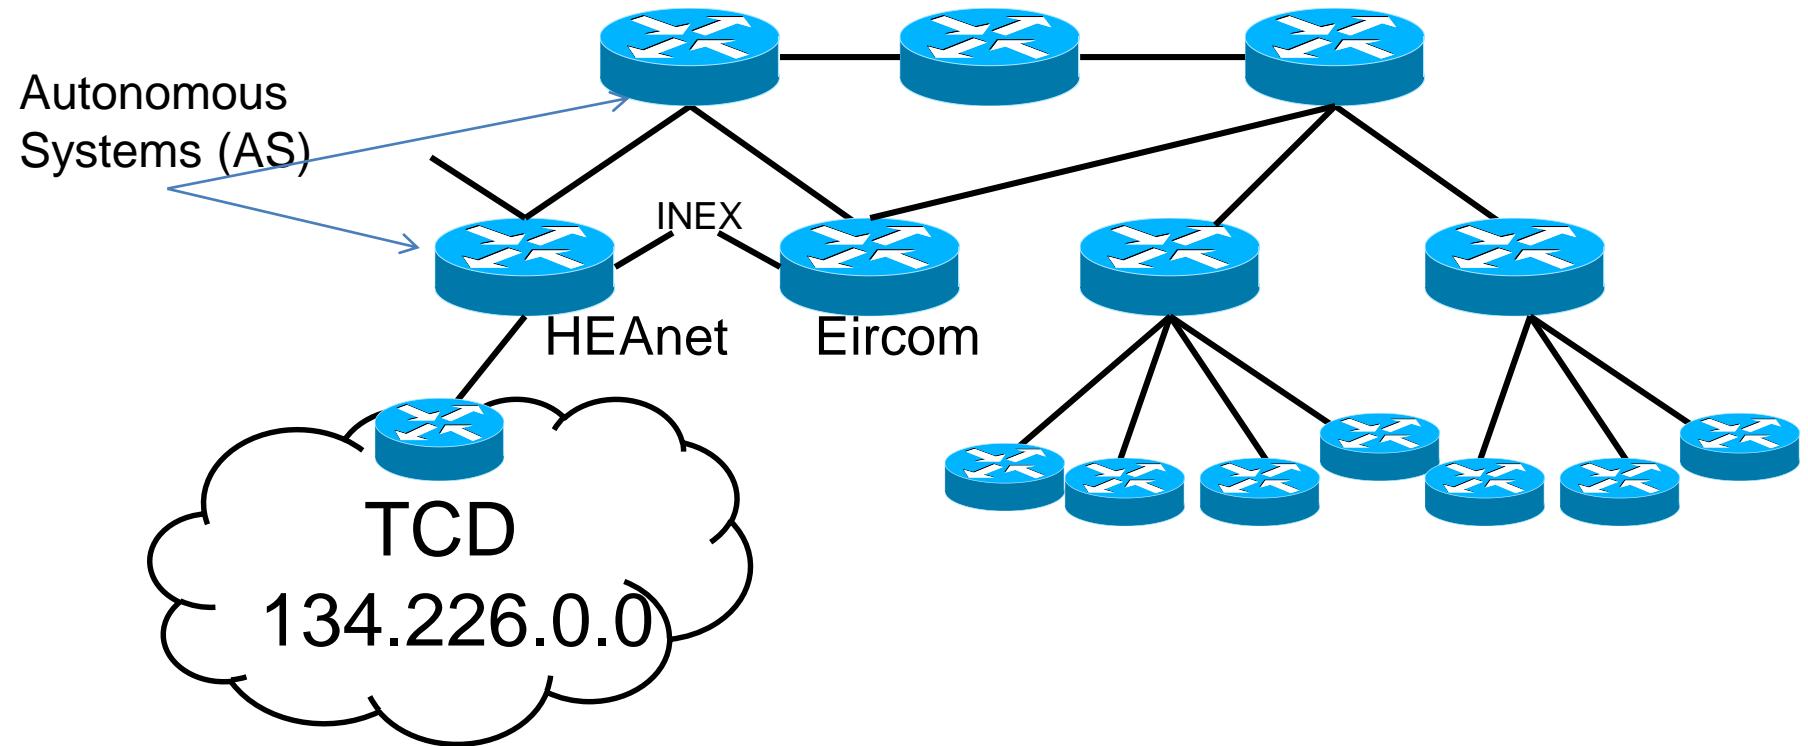
- Autonomous Systems
  - e.g. Companies, ISPs, 3<sup>rd</sup>-level Institutions

# Autonomous Systems

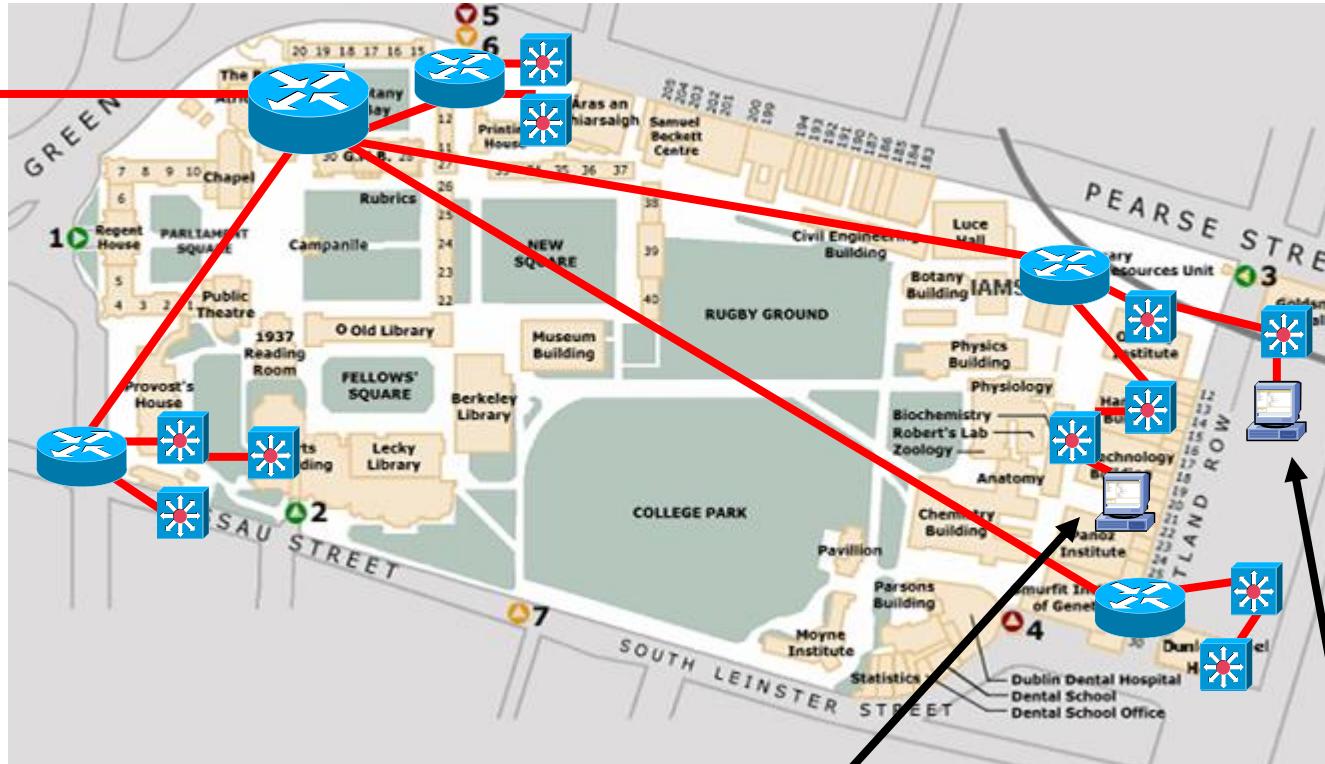


- Stub network
  - Network that does not forward to other network
- Transient network
  - Network that forwards traffic between other networks
- Point-to-point link

# Trinity's Connection



# Trinity Network with Routers



Switch



Router

**Address: 134.226.38.55  
Subnet: 134.226.38.0**

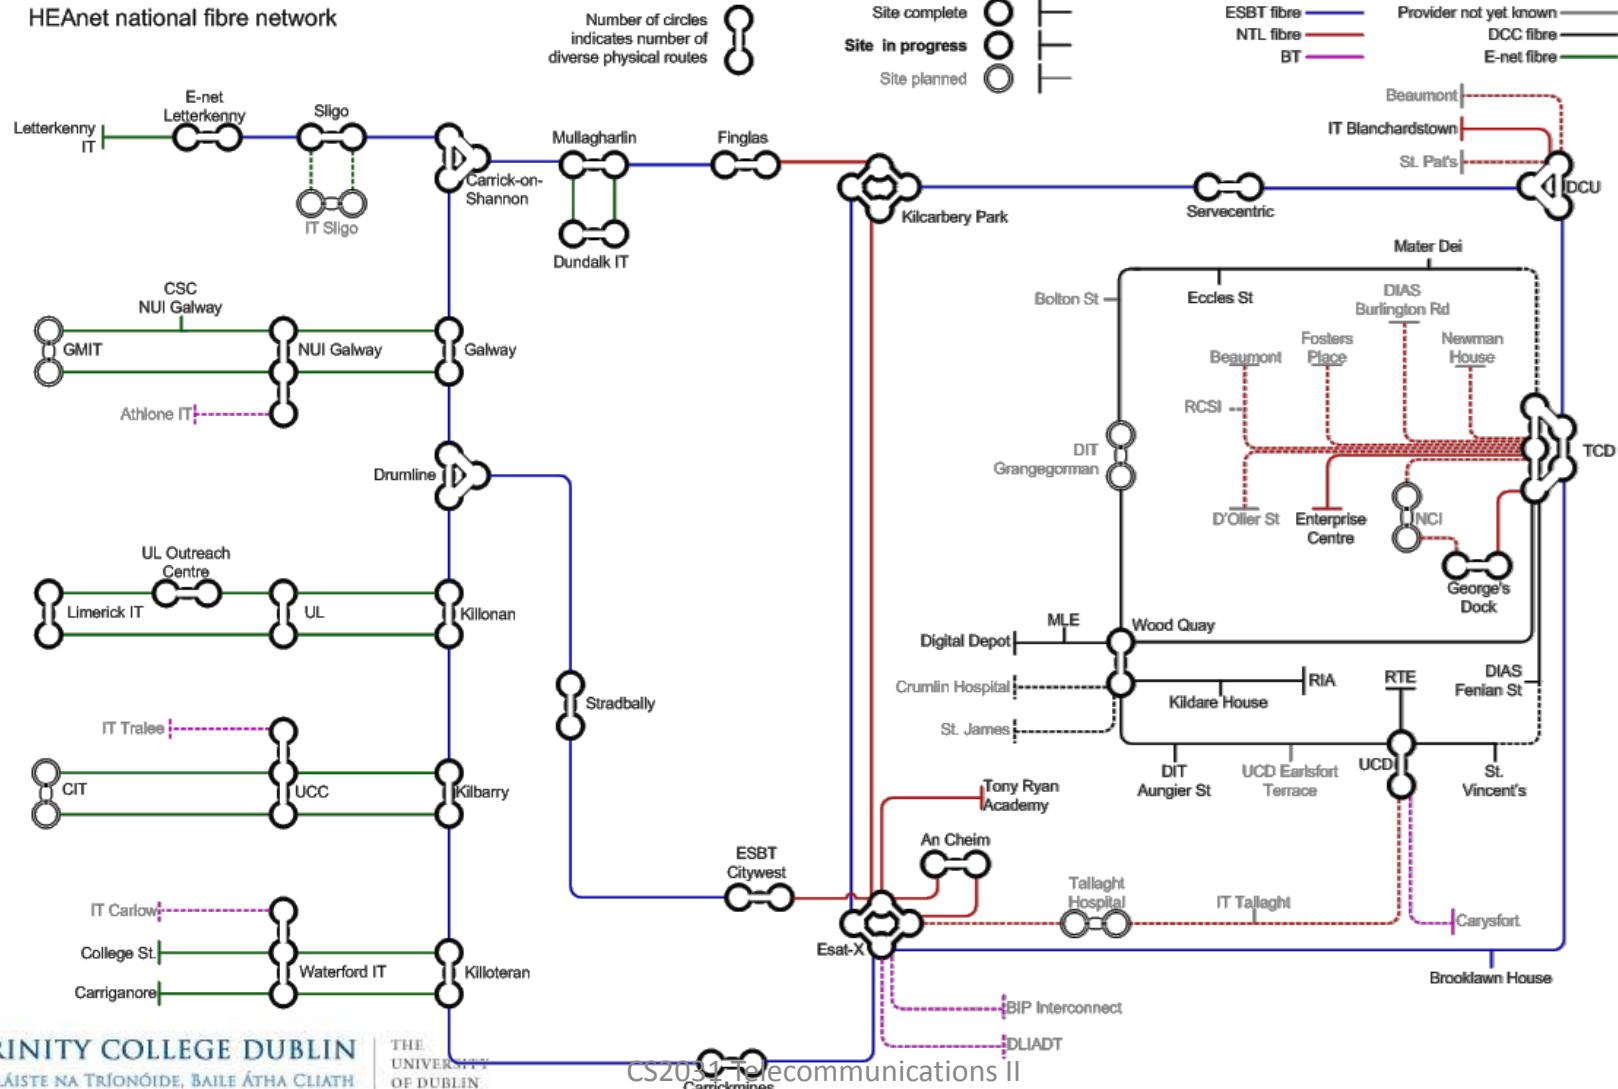
**Address: 134.226.32.55  
Subnet: 134.226.32.0**



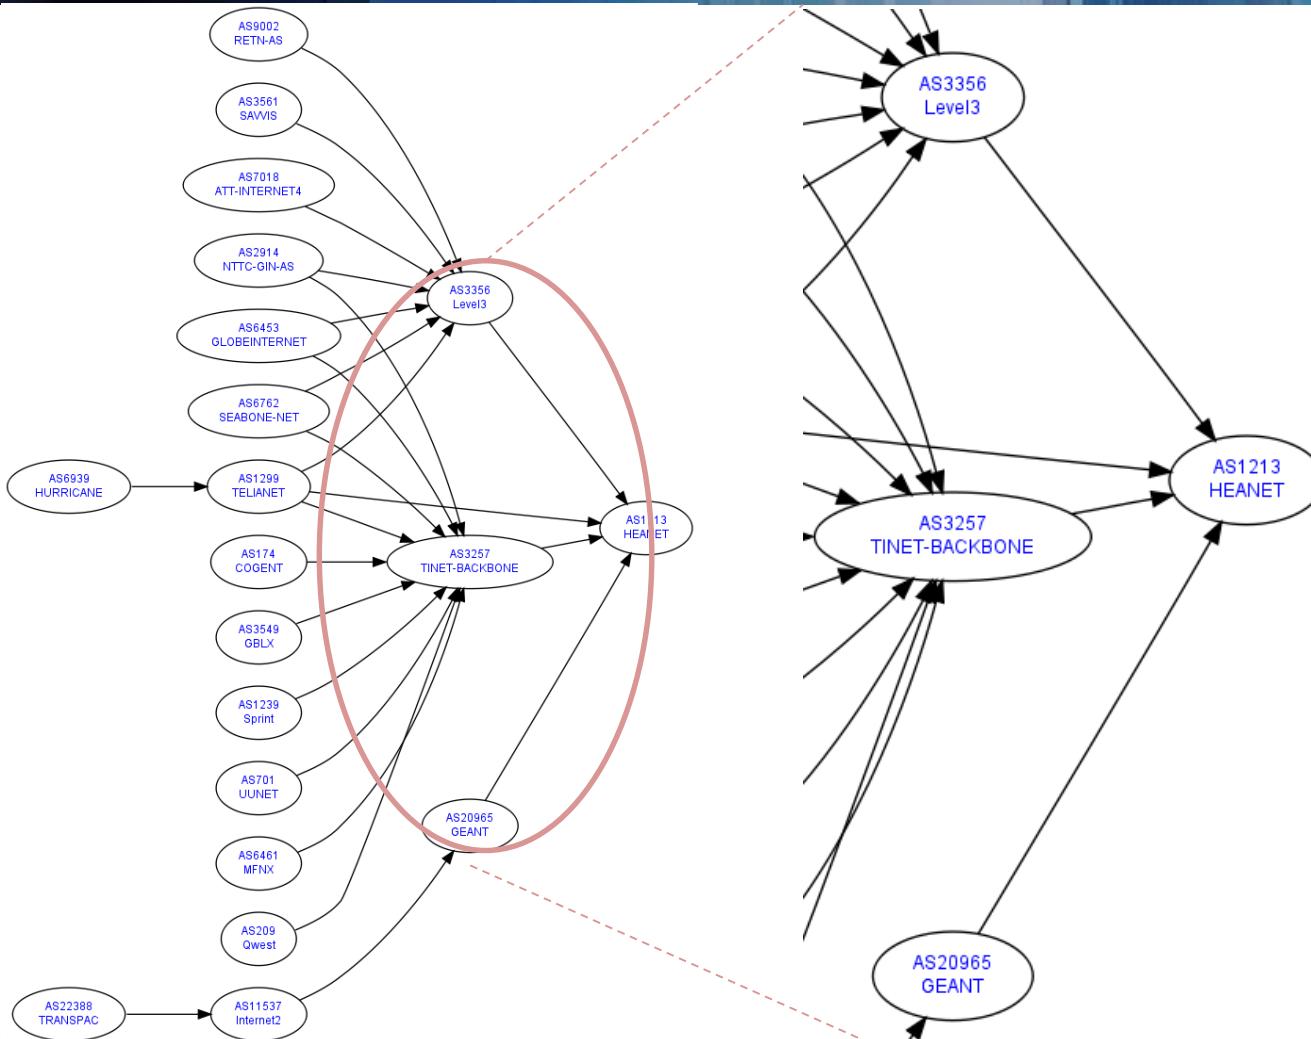
TRINITY COLLEGE DUBLIN  
COLÁISTE NA TRÍONÓIDE, BAILE ÁTHA CLIATH

THE  
UNIVERSITY  
OF DUBLIN

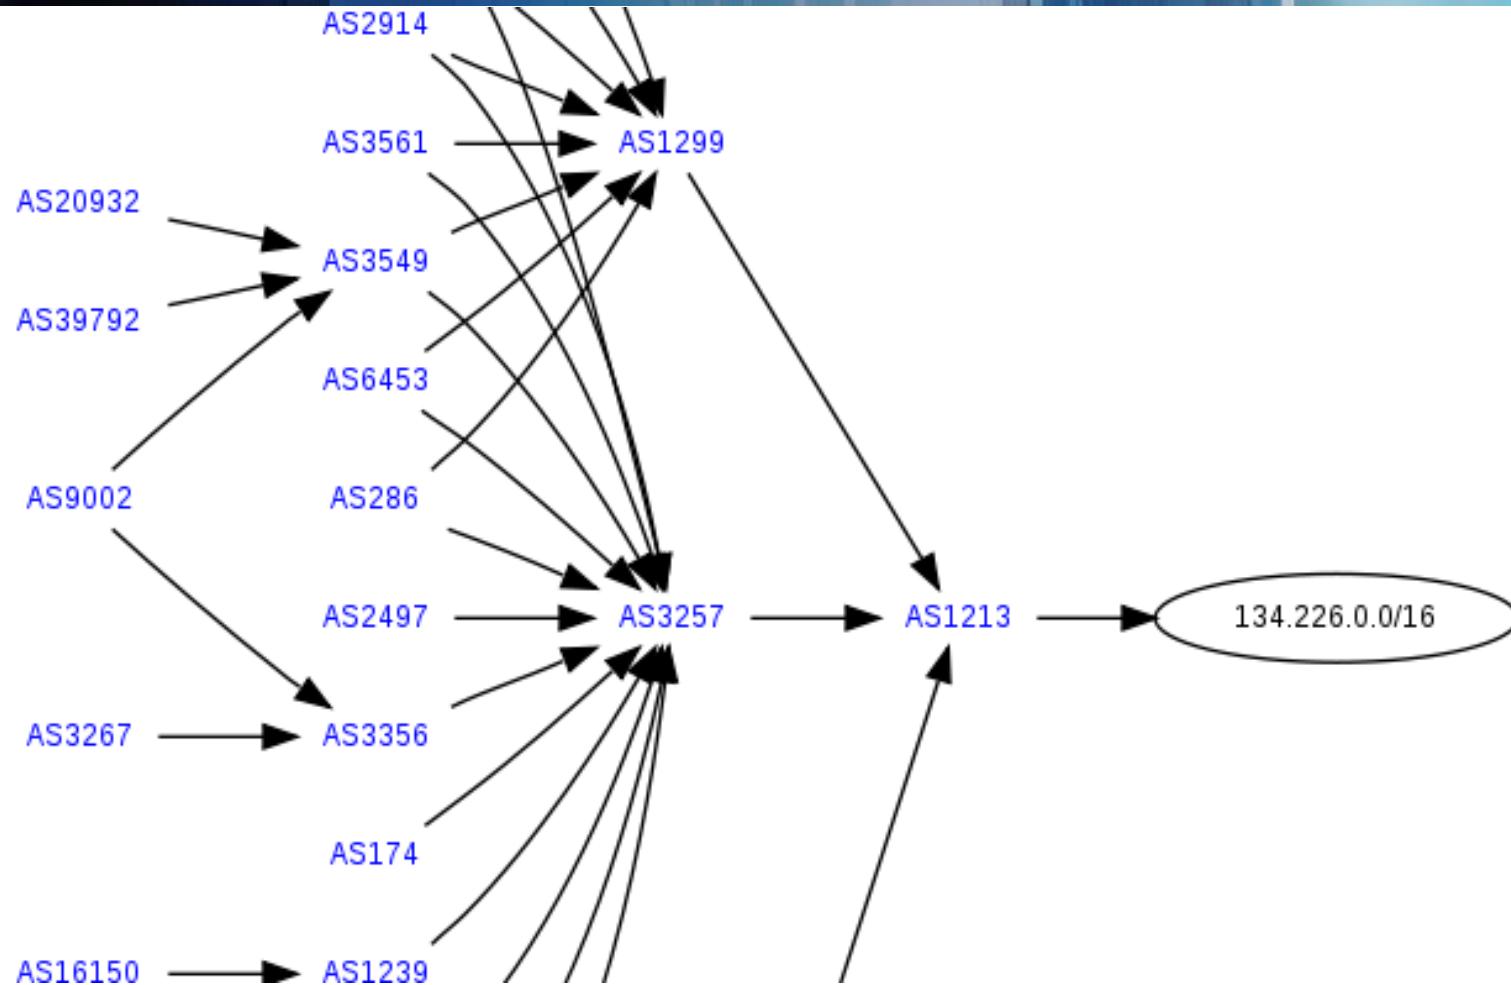
# HEAnet Fibre Network



# AS1213 - HEANET



# AS1213 - HEANET

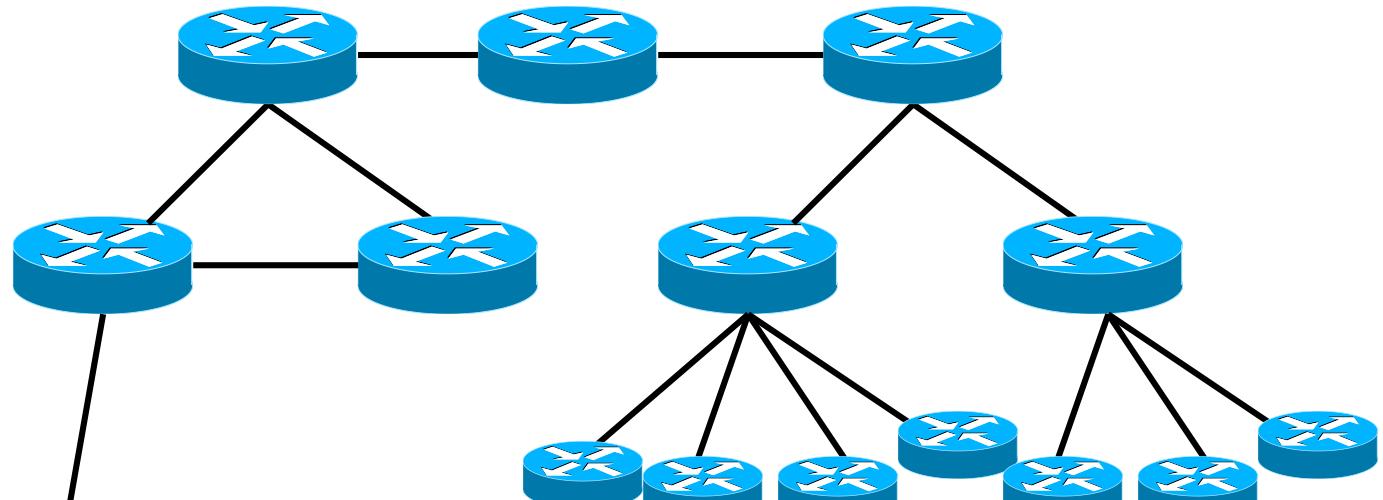


<http://www.robtex.com/route/134.226.0.0-16.html>

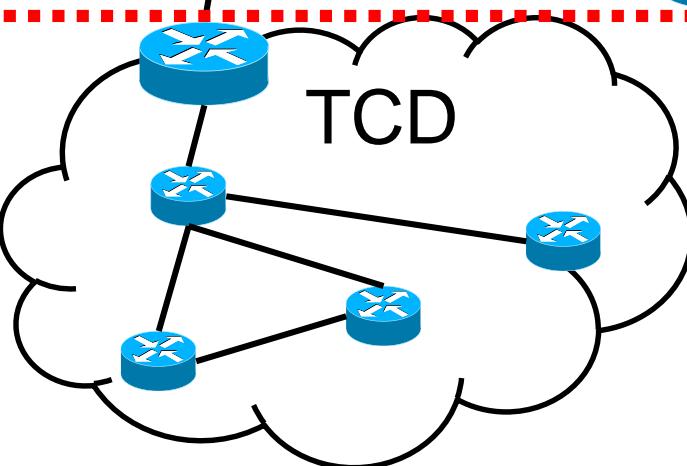


# InterAS vs IntraAS Routing

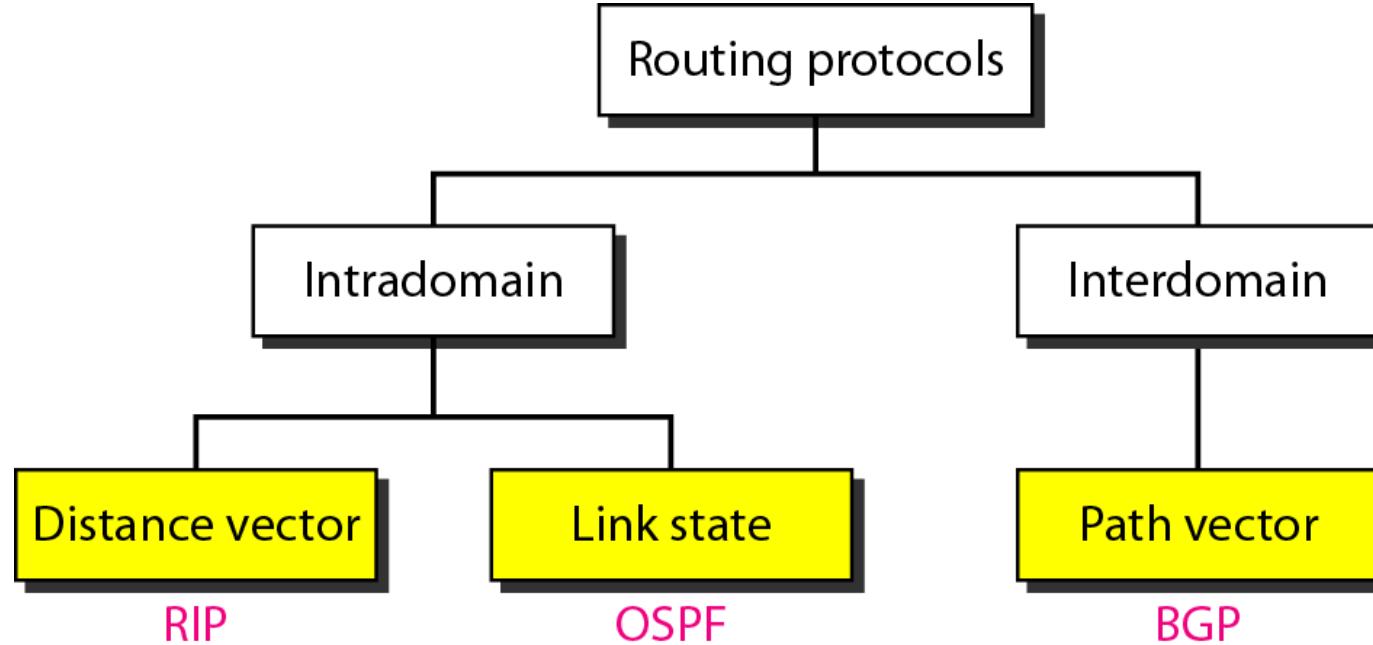
InterAS  
Routing



IntraAS  
Routing



# Routing Protocols



- Interior Routing Protocols
  - Routing within ASs
- Exterior Routing Protocols
  - Routing between ASs

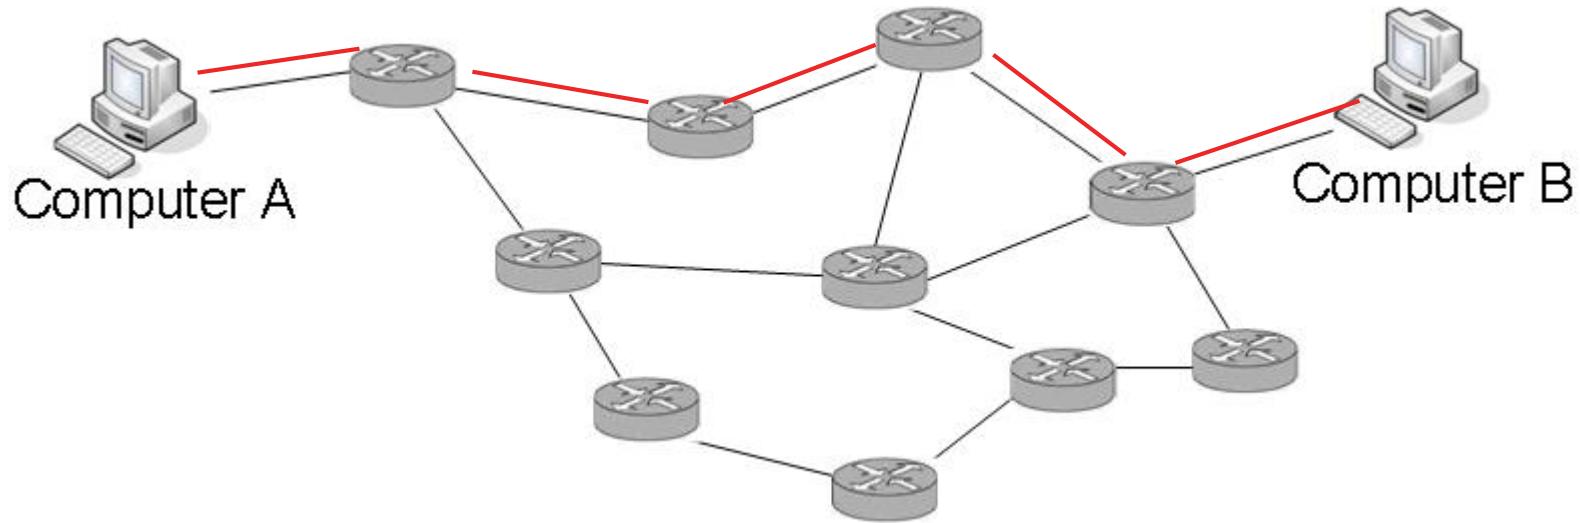
# Routers

- One Main Interest  
**Forwarding Packets**
- Important Aspects  
**Queue Length**  
**Routing Table**

Destination	Gateway	Interface
IP Range <sub>1</sub>	G <sub>1</sub>	IF <sub>1</sub>
IP Range <sub>2</sub>	G <sub>2</sub>	IF <sub>2</sub>

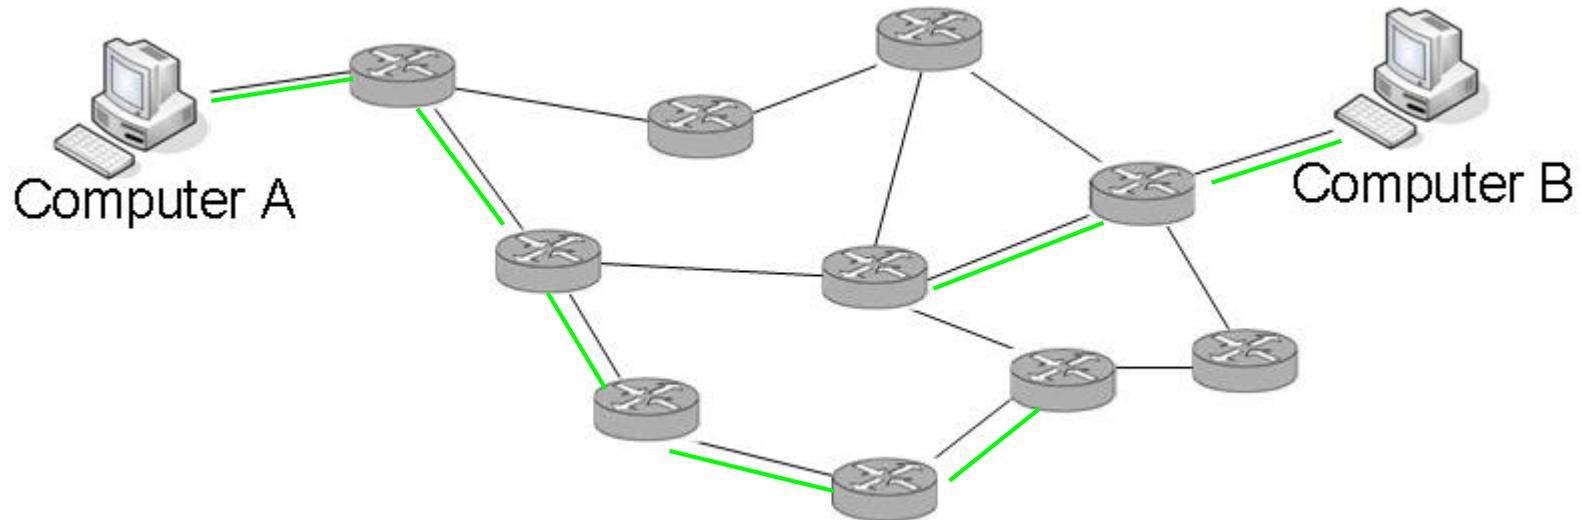


# The Scenario

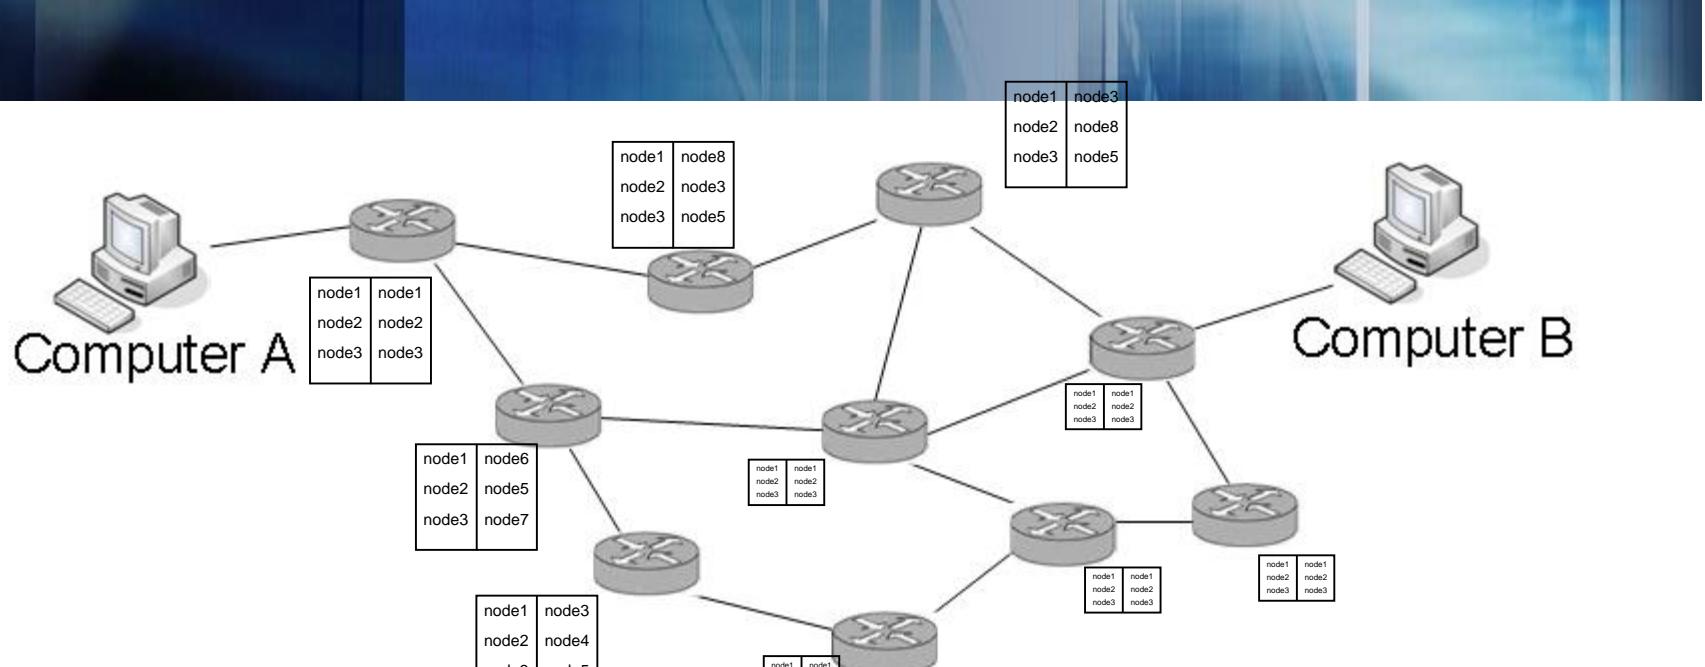


- Best route: Smallest number of hops?

# The Scenario



- Best route:
  - Fastest round-trip time?
  - Highest Bandwidth?



- Routing Tables
  - Creating tables
    - Dynamic vs. Static
  - Maintaining tables
    - Periodic vs. Aperiodic

# Routing Approaches

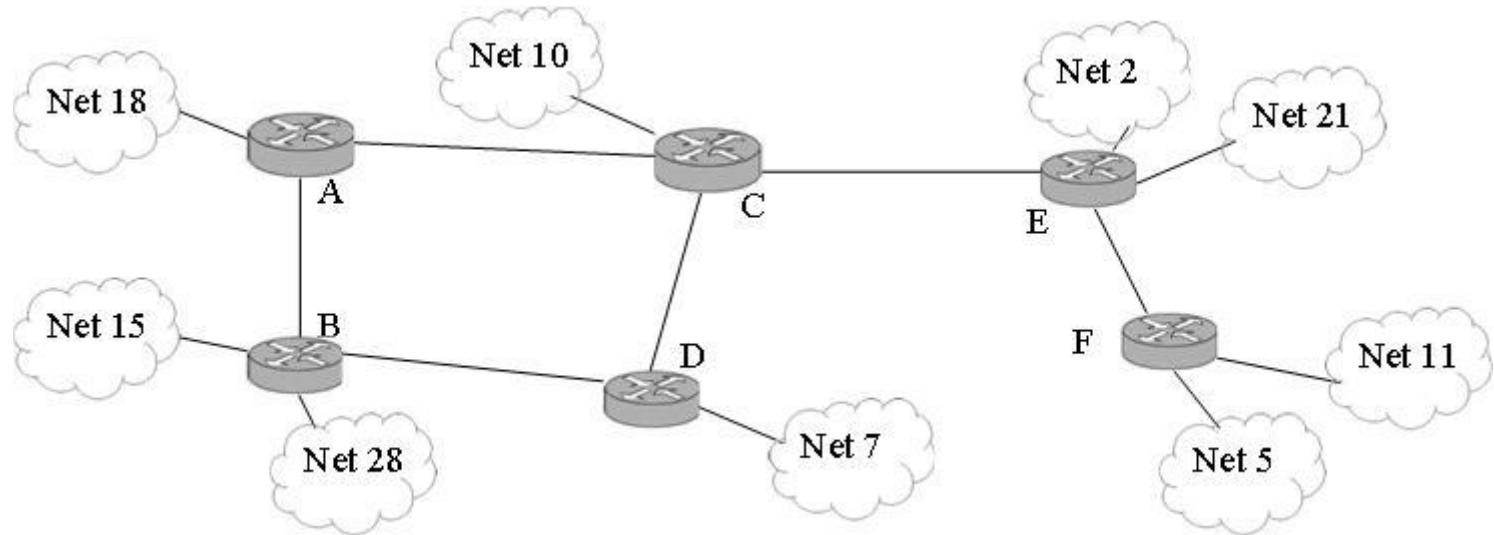
- Distance Vector routing
  - Routes propagate through exchange of routing tables
  - Based on communication with neighbours
- Link State routing
  - Establishing view of complete topology
  - Makes use of Dijkstra's Shortest-Path Algorithm

# Distance Vector Routing

- Each node maintains a set of triples
  - (**Destination**, **Cost**, **NextHop**)
- Exchange updates with directly connected neighbors
  - periodically (on the order of several seconds)
  - whenever table changes (called *triggered update*)
- Each update is a list of all pairs in the routing table:
  - (**Destination**, **Cost**)

Destination	Hop Count	Next Router	Other info
163.5.0.0	7	172.6.23.4	
197.5.13.0	5	176.3.6.17	
189.45.0.0	4	200.5.1.6	
115.0.0.0	6	131.4.7.19	

# Distance Vector: Example



A	Net18	1	-

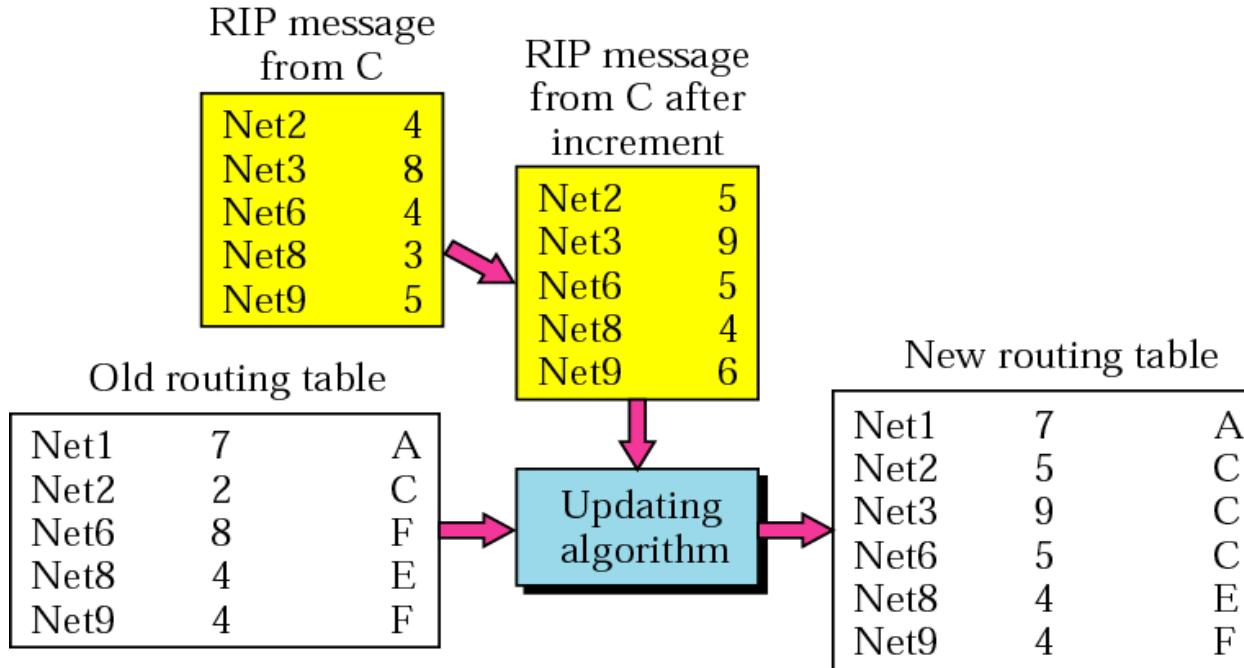
B	Net15	1	-

C	Net10	1	-

D	Net7	1	-

- Each router know the networks that are immediately connected to it

# Receiving an Update



Net1: No news, do not change

Net2: Same next hop, replace

Net3: A new router, add

Net6: Different next hop, new hop count smaller, replace

Net8: Different next hop, new hop count the same, do not change

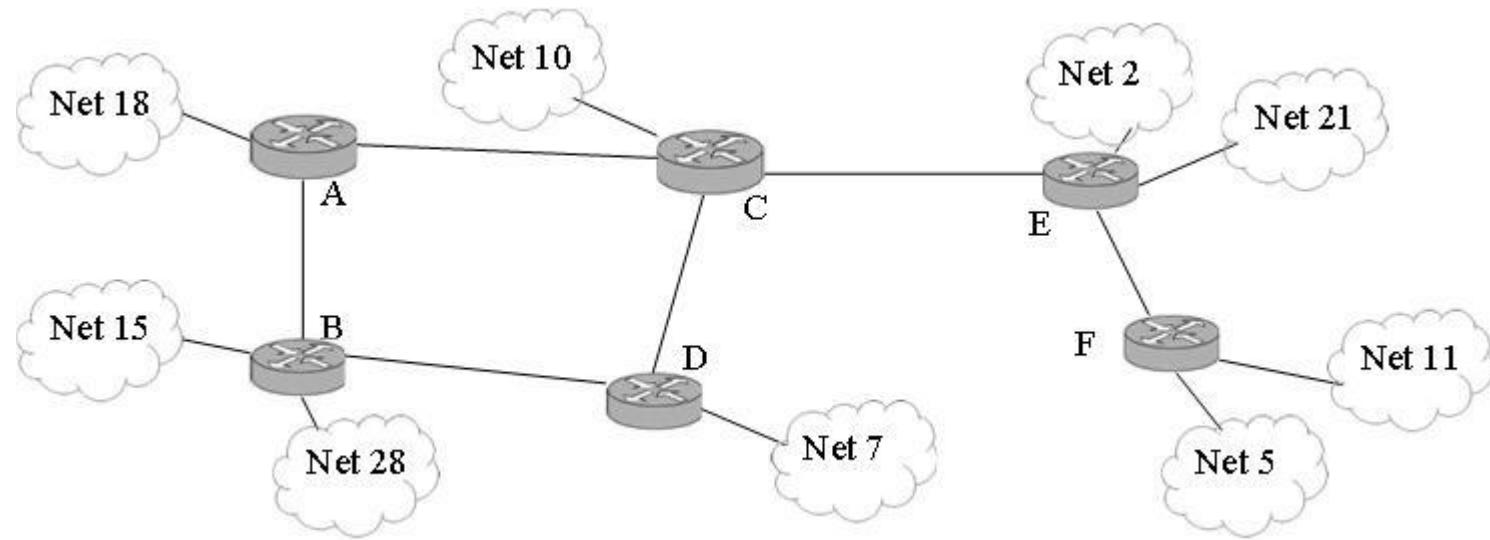
Net9: Different next hop, new hop count larger, do not change

# RIP Updating Algorithm

Receive: A response RIP message

- Add one hop to the hop count for each advertised destination.
- Repeat the following steps for each advertised destination:
  - a) If (destination not in the routing table)  
Add the advertised information to the table.
  - b) Else  
If (next-hop field is the same)  
Replace entry in the table with the advertised one.  
Else  
If (advertised hop count smaller than one in the table)  
Replace entry in the routing table.

# Distance Vector: Example II



A		
Net18	1	-
Net15	2	B
Net28	2	B
Net10	2	C

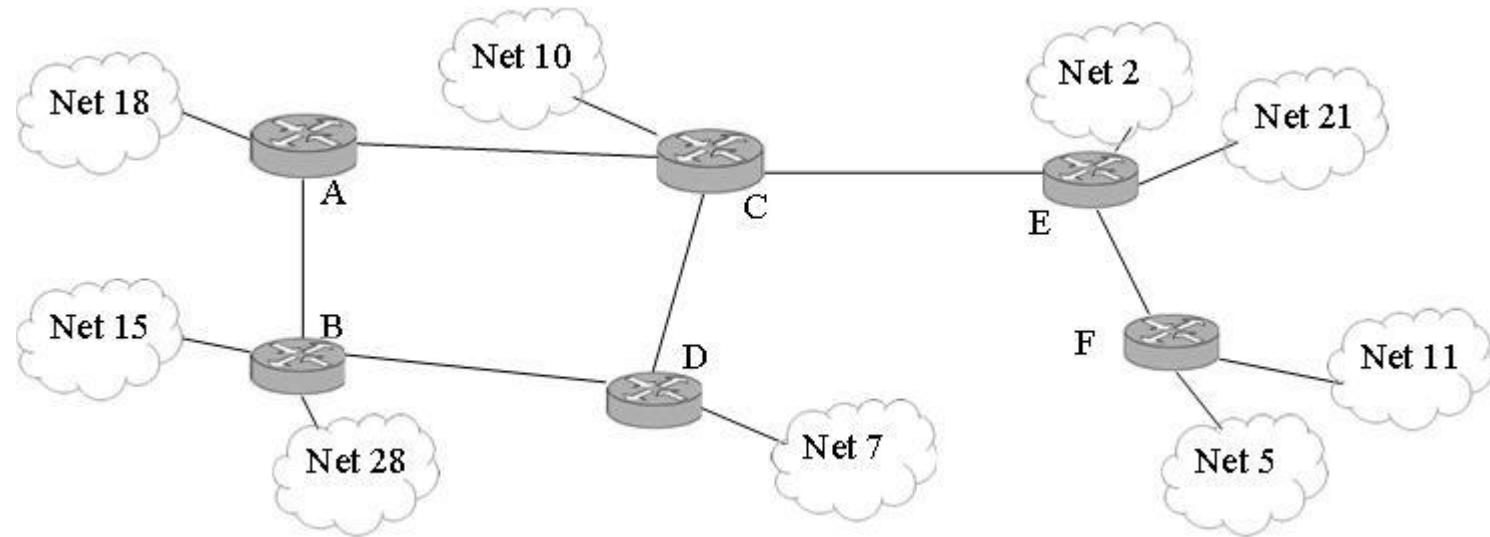
B		
Net15	1	-
Net28	1	-
Net18	2	A
Net7	2	D

C		
Net10	1	-
Net7	2	D
Net2	2	E
Net21	2	E

D		
Net7	1	-
Net15	2	B
Net28	2	B
Net10	2	C

- e.g. Node A incorporates information from Node B

# Distance Vector: Example III



A		
Net18	1	-
Net15	2	B
Net28	2	B
Net10	2	C
Net7	3	B

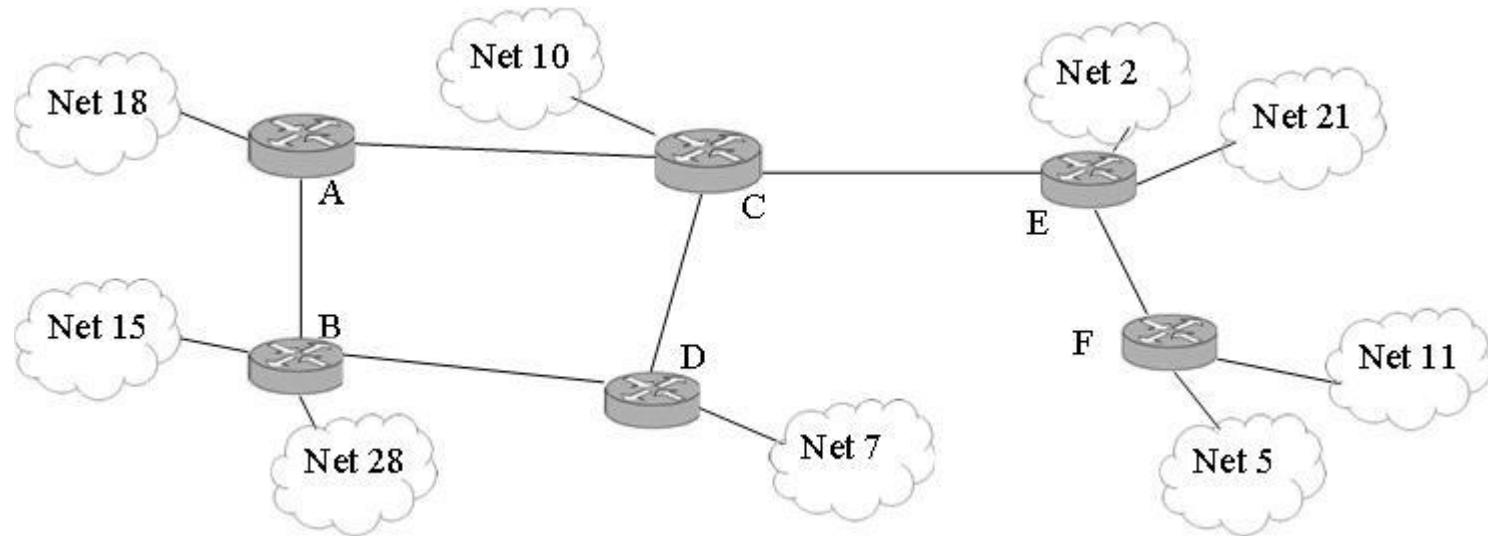
B		
Net15	1	-
Net28	1	-
Net18	2	A
Net7	2	D
Net10	3	D

C		
Net10	1	-
Net7	2	D
Net2	2	E
Net21	2	E
Net11	3	E
Net5	3	E

D		
Net7	1	-
Net15	2	B
Net28	2	B
Net10	2	C
Net2	3	C
Net21	3	C

- e.g. Node A receives information Node B received during the last round

# Distance Vector: Example IV



A		
Net18	1	-
Net15	2	B
Net28	2	B
Net10	2	C
Net7	3	B
...	...	...
Net11	5	C
Net5	5	C

- Time require to build complete routing tables for all nodes
  - this is known as convergence

# Count-to-Infinity Problem



Initially

1	•	•	•	•
1	2	•	•	•
1	2	3	•	•
1	2	3	4	

(a)



Initially

1	2	3	4
3	2	3	4
3	4	3	4
5	4	5	4
5	6	5	6
7	6	7	6
7	8	7	8
⋮	⋮	⋮	⋮

(b)

- Link between A and B breaks
- B receives update from C advertising route to A

- Routers exchange updates
- After 4 steps the network converges
- Every router knows how to get to router A

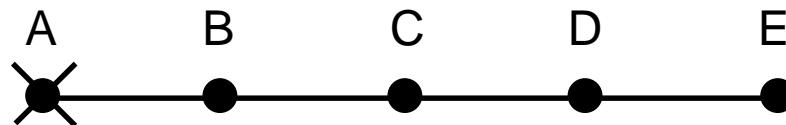


# Solutions for Count-to-Infinity

- Possible solutions:
  1. Impose upper bound on maximum distance
  2. *Split horizon*
    - C should not send to node B its new distance to node A, if node B is C's next-hop towards A
  3. *Split horizon with poisoned reverse*
    - C should tell node B that its distance to node A is  $\infty$ , when node B is C's next-hop towards A
- Unfortunately, none of these solutions can deal with arbitrary topology cycles

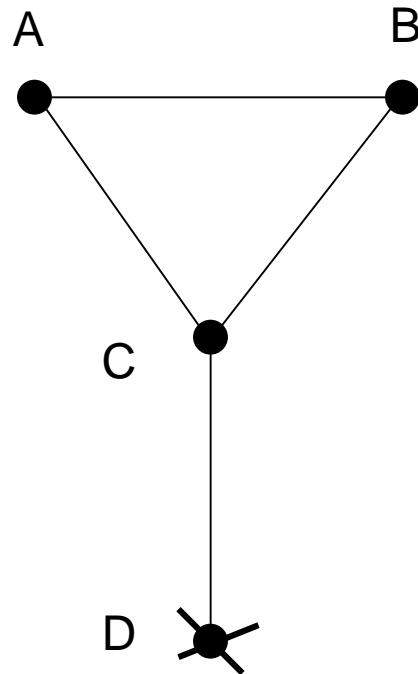
# Split Horizon with Poisoned Reverse

Report “split-horizon” routes as **infinity** to break loops on the first routing exchange.



inf.	2	3	4	B learns A is dead
inf.	<u>inf.</u> → 2	3	4	B reports to C that A's metric is inf.
inf.	inf.	3	4	After 1 exchange
inf.	inf.	inf.	4	After 2 exchanges
inf.	inf.	inf.	inf.	After 3 exchanges

# Split Horizon Failure

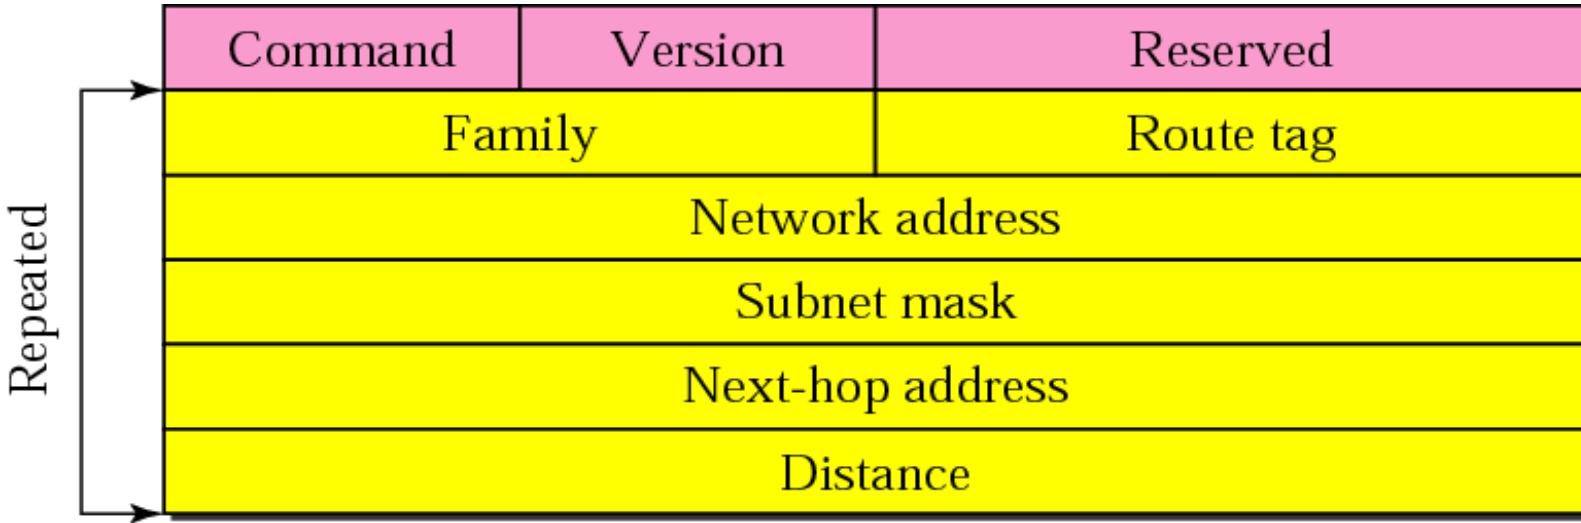


If D goes down, A and B will still count to infinity.

Split-Horizon infinity messages are sent from A->C and B->C, not A<->B

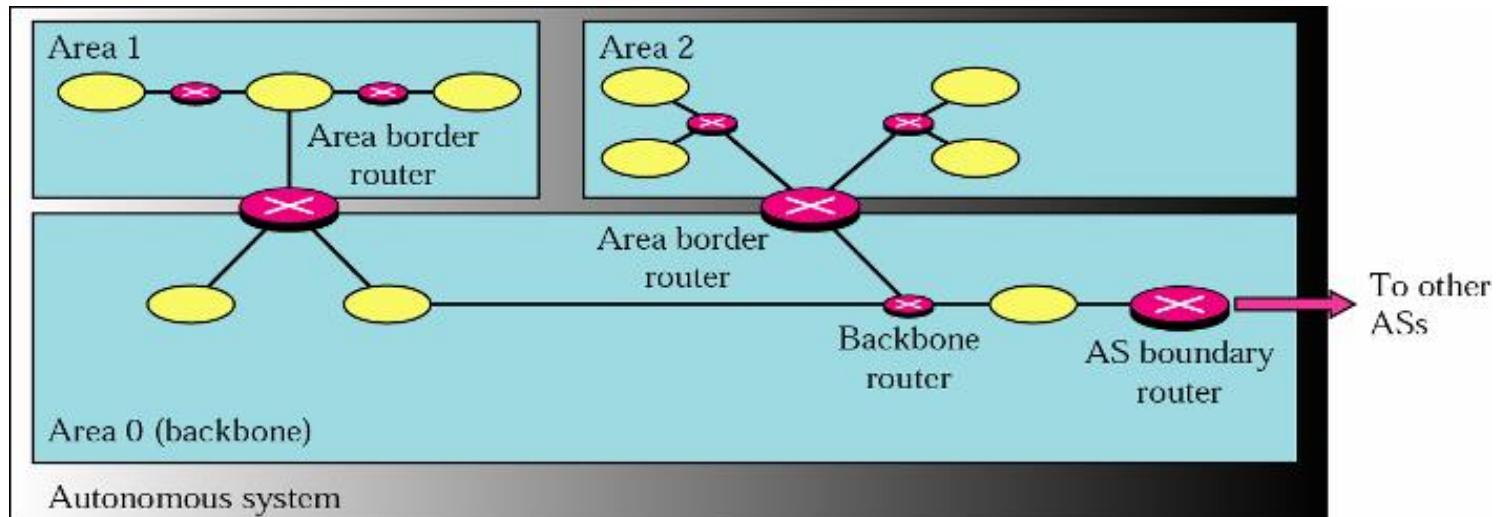


# RIPv2 Message Format



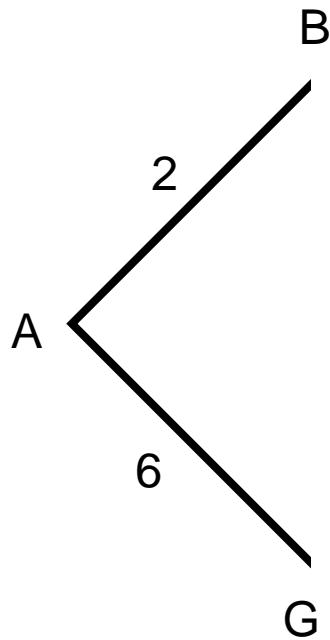
# Link State Routing

- Sharing knowledge about neighbourhood
- Sharing with every other router
- Uses Dijkstra's Shortest-Path Algorithm to calculate the routing table
- Open Shortest Path First (OSPF)
  - Divides autonomous system into areas
  - Border routers summarize information for an area



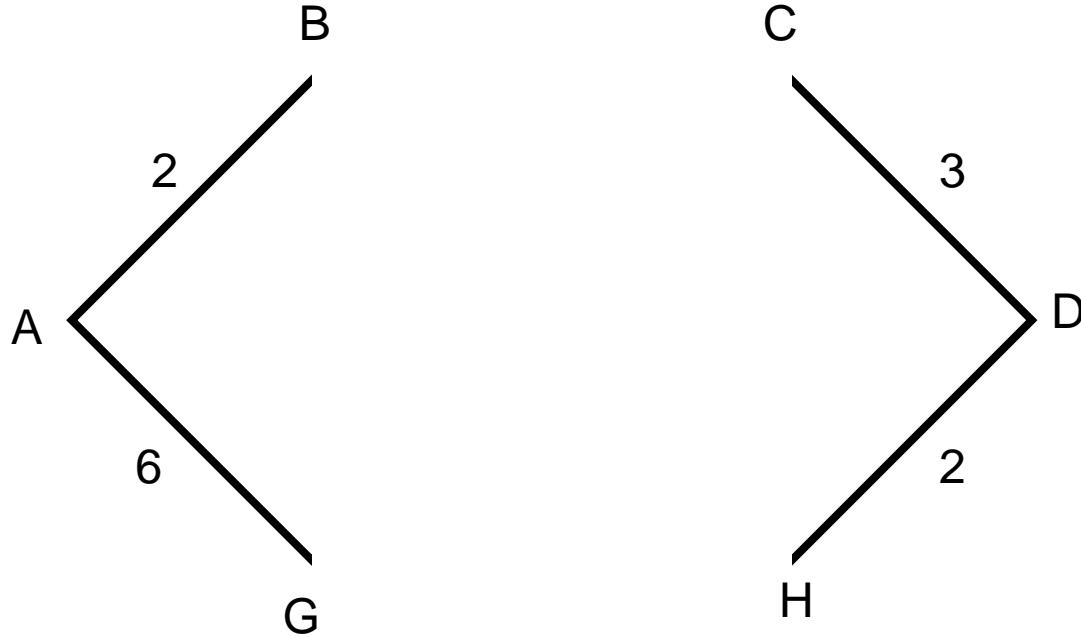
## Example: Building a network graph

A discovers 2 neighbours B and G in distance 2 and 6 respectively



## Example: Building a network graph

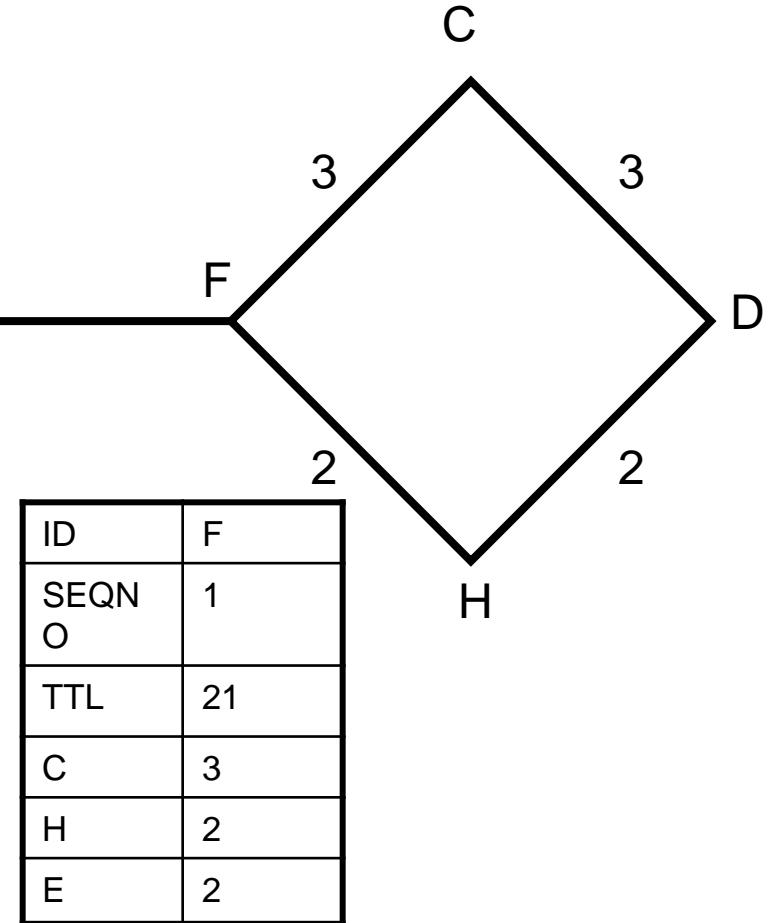
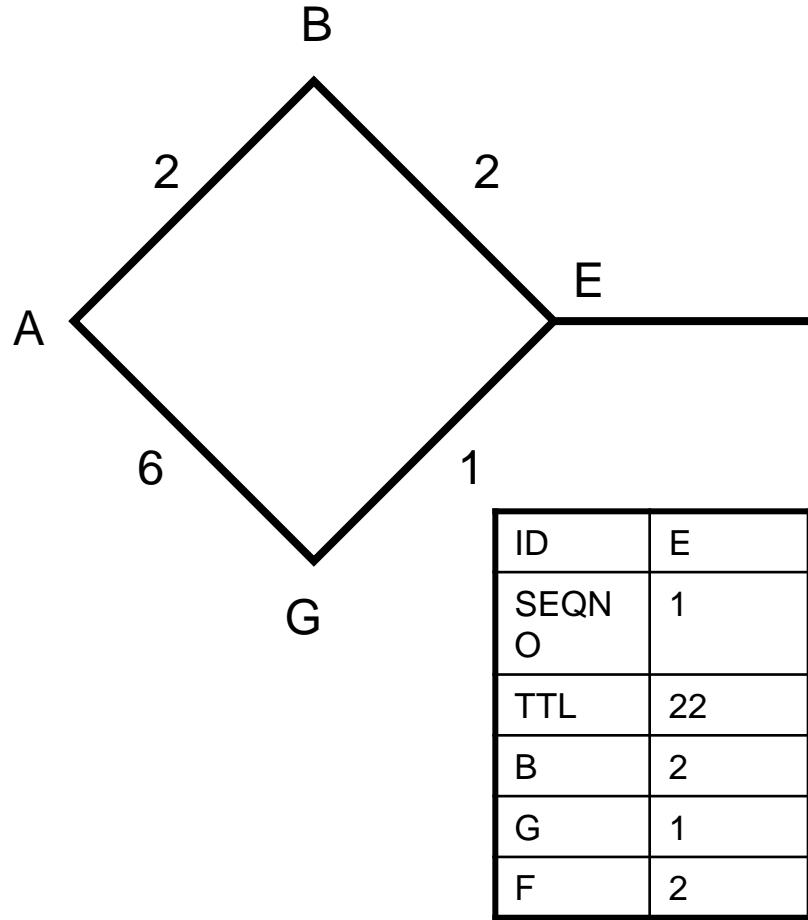
A receives an update from D



ID	D
SEQN	0
TTL	21
C	3
H	2

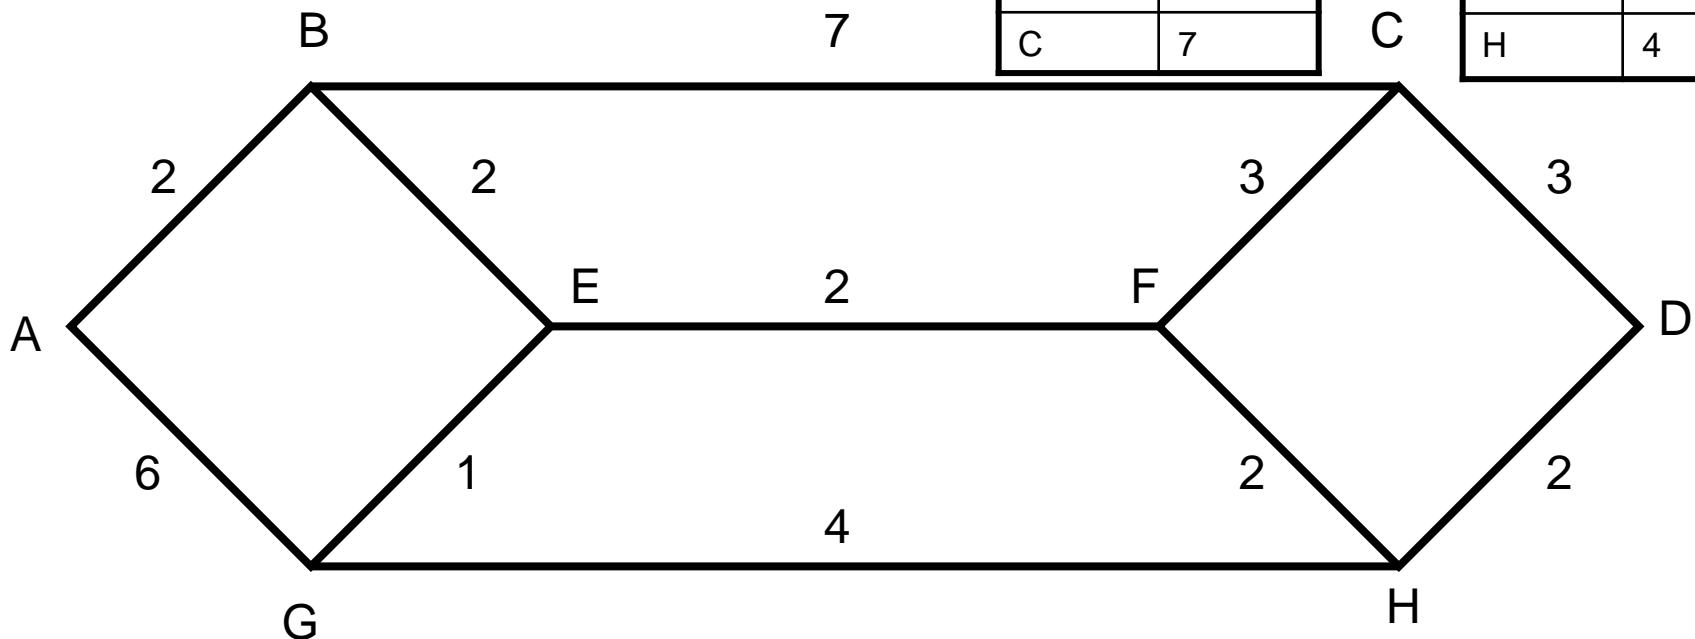
# Example: Building a network graph

A receives updates from E and F



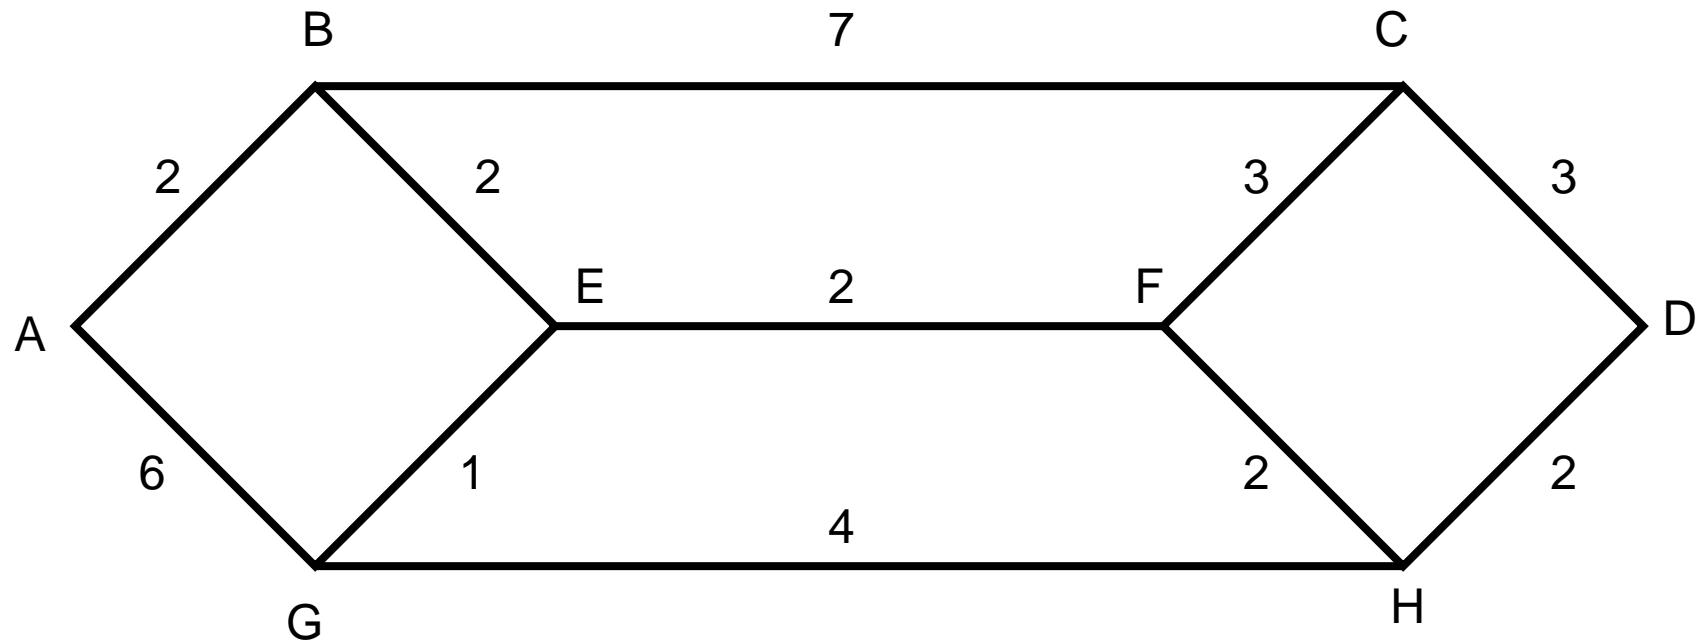
# Example: Building a network graph

A receives updates from B and G



## Example: Building a network graph

Final graph



# Dijkstra's Algorithm

Let

- $N$  denotes set of nodes in the graph
- $l(i, j)$  denotes non-negative cost (weight) for edge  $(i, j)$
- $s$  denotes initial node (source)
- $M$  denotes the set of nodes incorporated so far
- $C(n)$  denotes cost of the path from  $s$  to node  $n$

```
M = {s}  
for each n in N - {s}  
    C(n) = l(s, n)  
while (N != M)  
    M = M ∪ {w} such that C(w) is the minimum for  
        all w in (N - M)  
    for each n in (N - M)  
        C(n) = MIN(C(n), C(w) + l(w, n))
```

# Dijkstra's Algorithm

1. Start with the local node (router): the root of the tree.
2. Assign a cost of 0 to this node and make it the first permanent node.
3. Examine each neighbour node of the node that was the last permanent node.
4. Assign a cumulative cost to each node and make it tentative.
5. Among the list of tentative nodes
  - a) Find the node with the smallest cumulative cost and make it permanent.
  - b) If a node can be reached from more than one direction
    - I. Select the direction with the shortest cumulative cost.
6. Repeat steps 3 to 5 until every node becomes permanent.

# Example: Dijkstra's Algorithm

Routing Table


Tentative Nodes

A	0	-



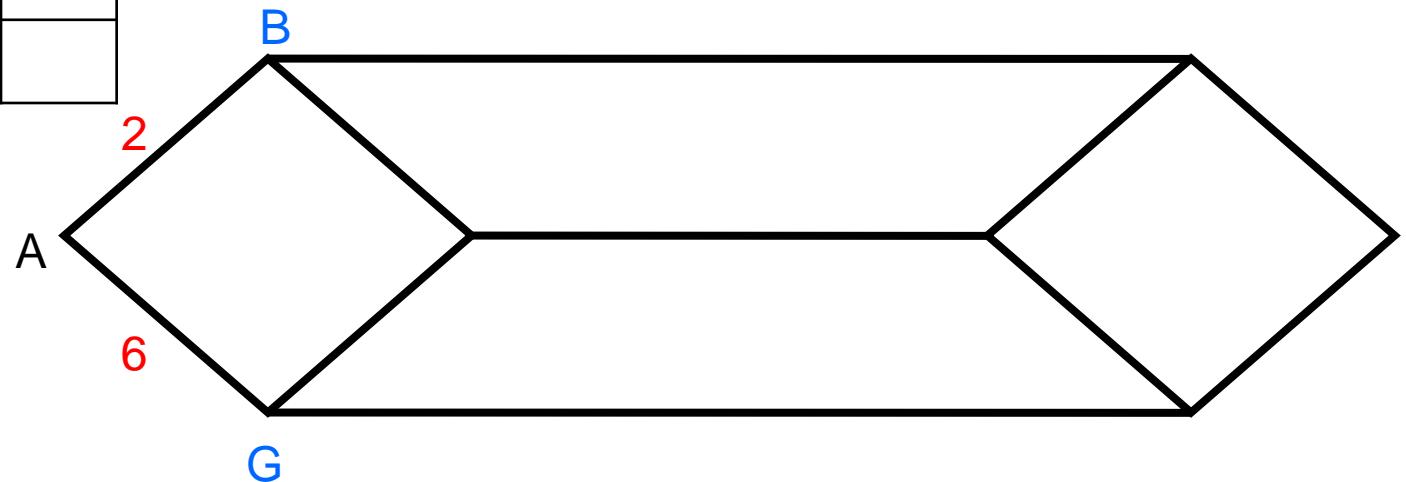
# Example: Dijkstra's Algorithm

Routing Table

A	0	-

Tentative Nodes

B	2	A
G	6	A



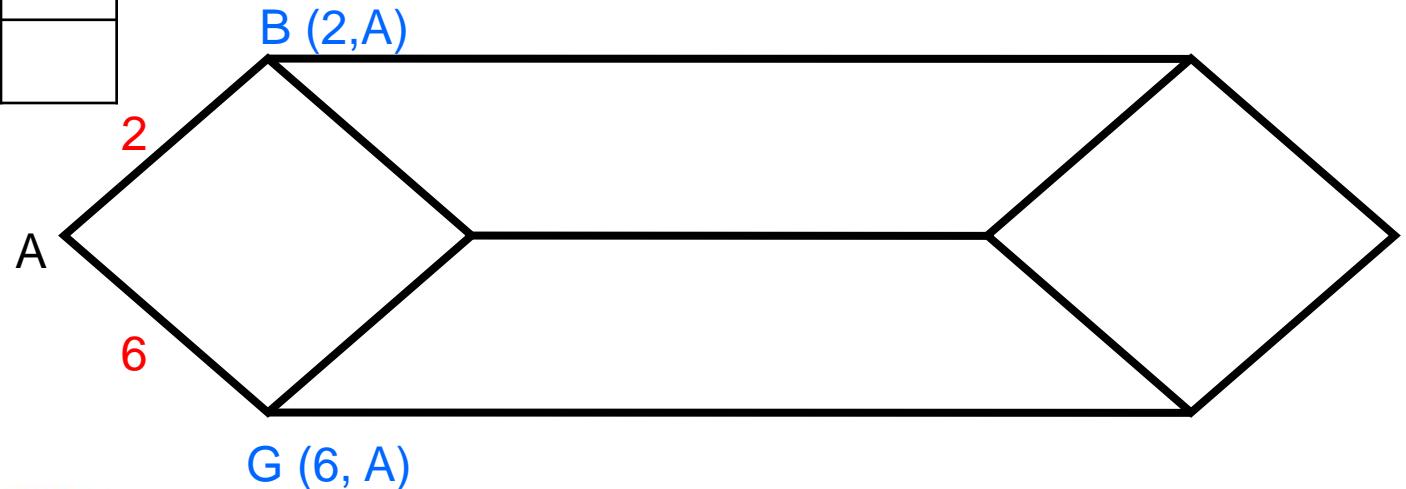
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A

Tentative Nodes

G	6	A
---	---	---



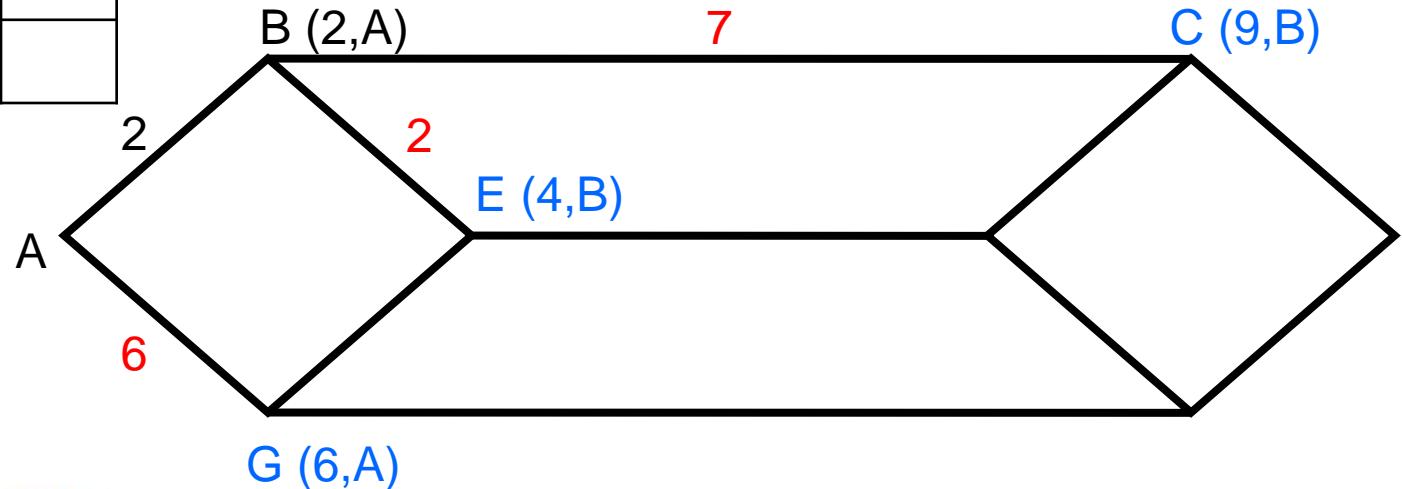
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A

Tentative Nodes

G	6	A
E	4	B
C	9	B



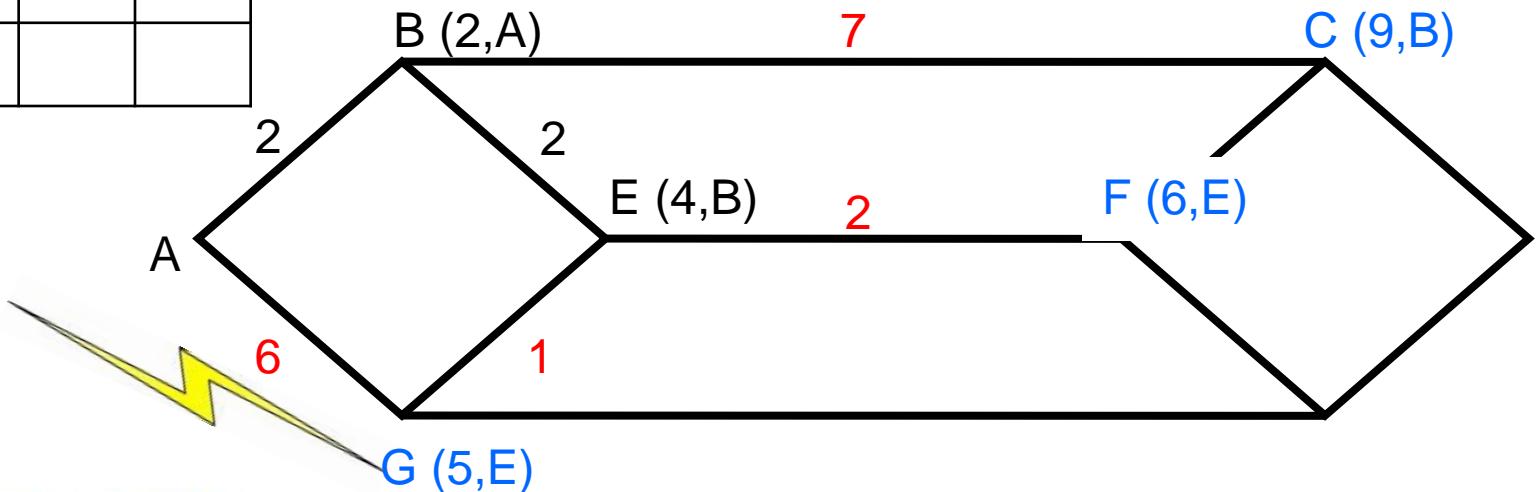
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A
E	4	B

Tentative Nodes

G	0	A
C	9	B
F	6	E
G	5	E



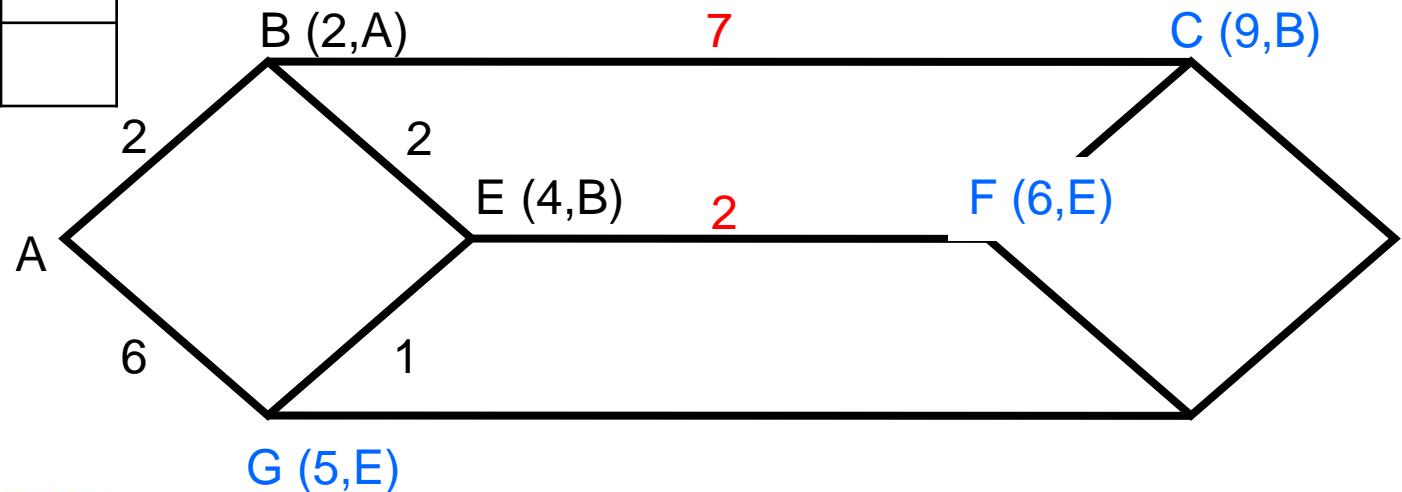
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A
E	4	B
G	5	E

Tentative Nodes

C	9	B
F	6	E



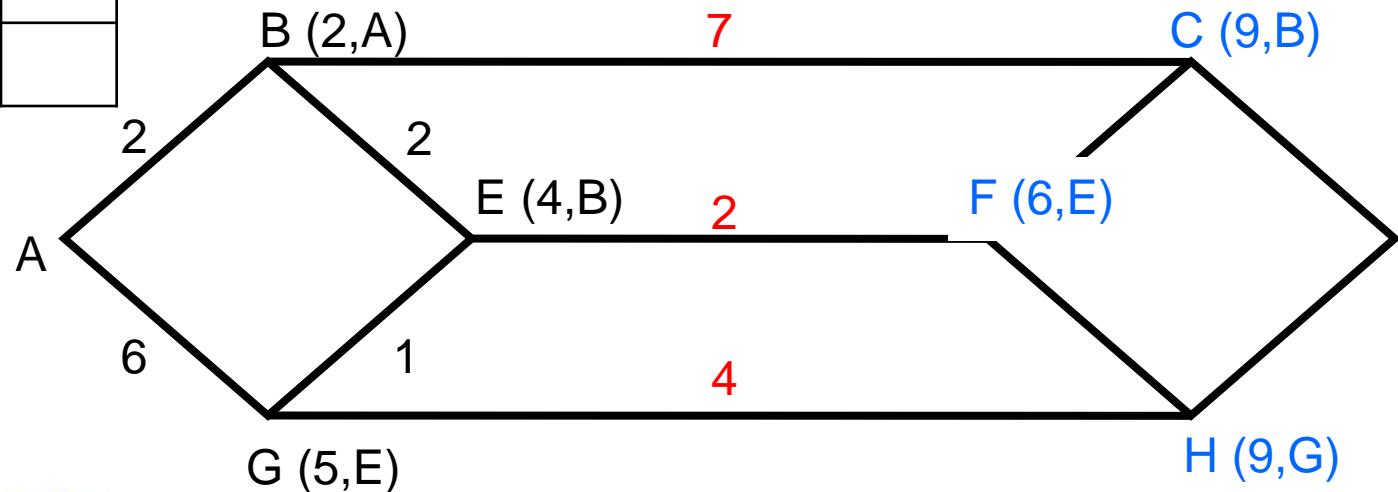
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A
E	4	B
G	5	E

Tentative Nodes

C	9	B
H	9	G
F	6	E



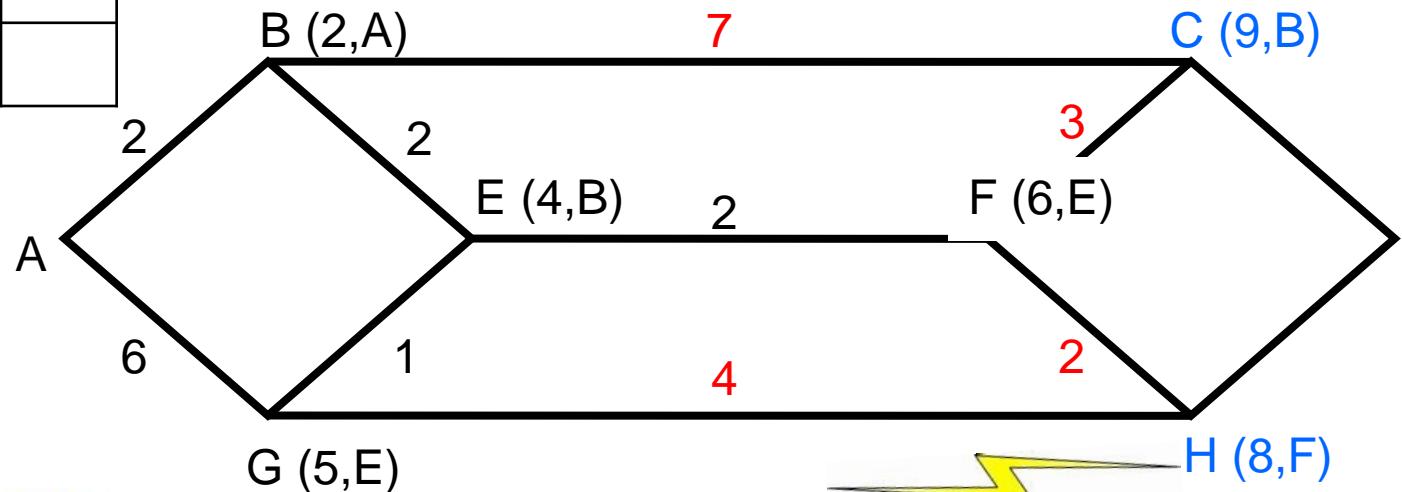
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A
E	4	B
G	5	E
F	6	E

Tentative Nodes

C	9	B
H	0	G
H	8	F



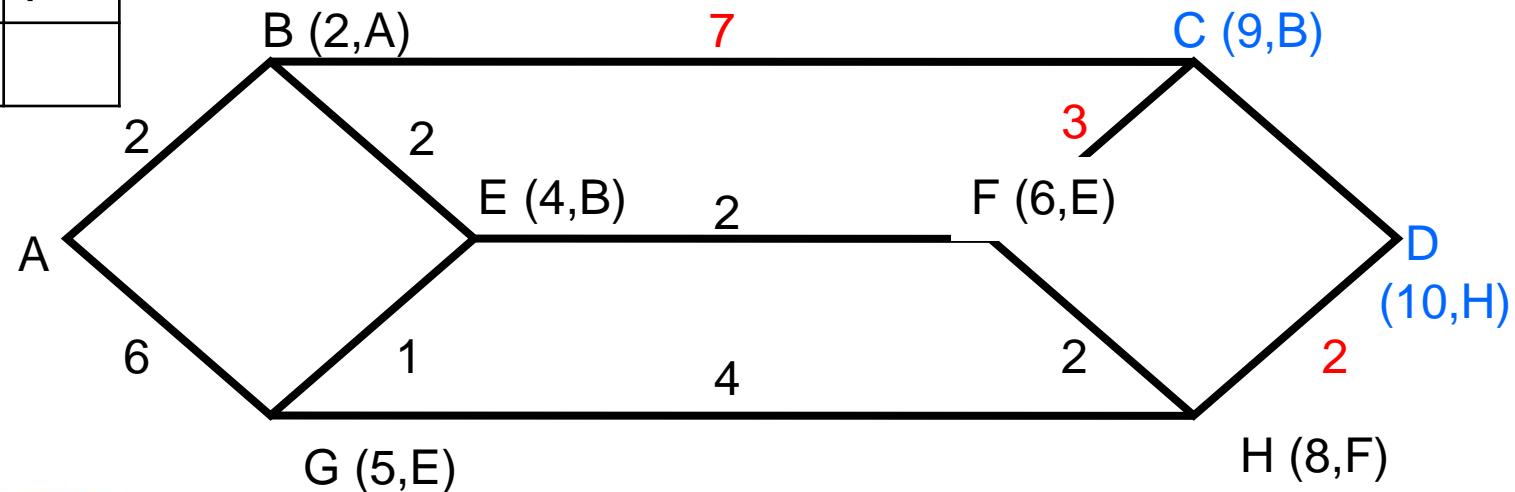
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A
E	4	B
G	5	E
F	6	E
H	8	F

Tentative Nodes

C	9	B
D	10	H



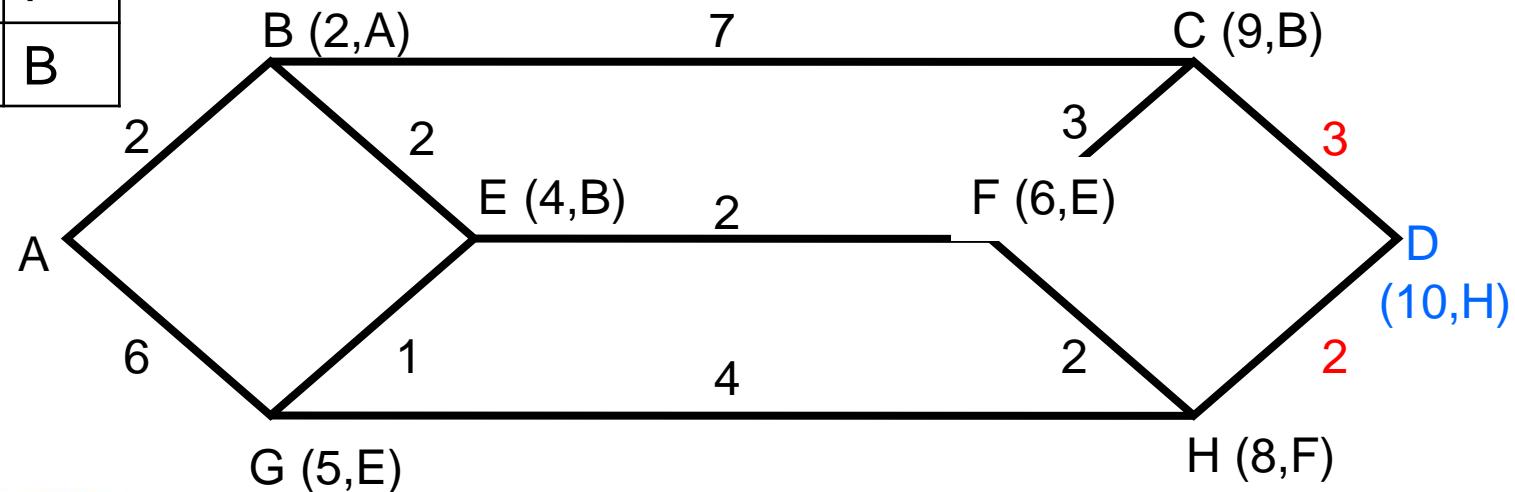
# Example: Dijkstra's Algorithm

Routing Table

A	0	-
B	2	A
E	4	B
G	5	E
F	6	E
H	8	F
C	9	B

Tentative Nodes

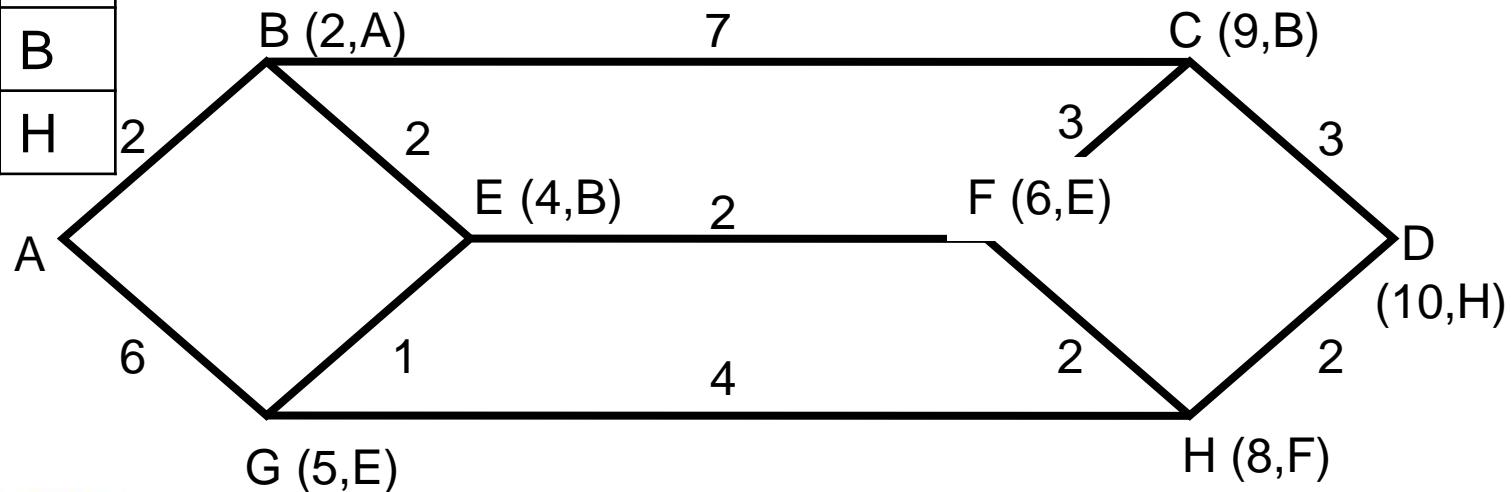
D	12	C
D	10	H



# Example: Dijkstra's Algorithm

Routing Table

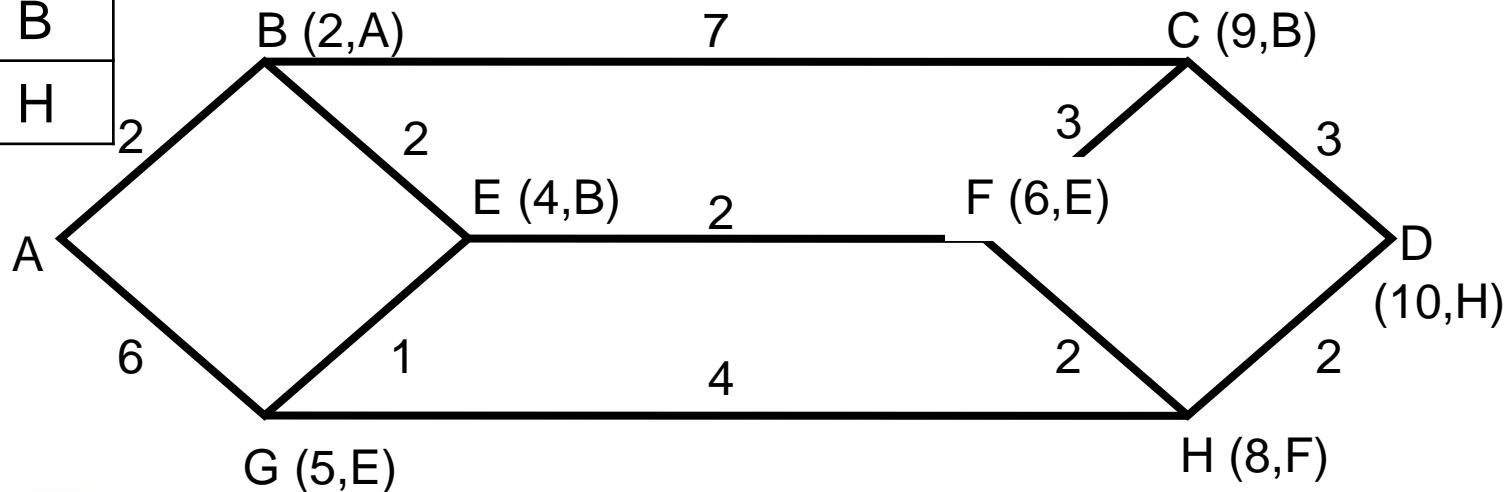
A	0	-
B	2	A
E	4	B
G	5	E
F	6	E
H	8	F
C	9	B
D	10	H



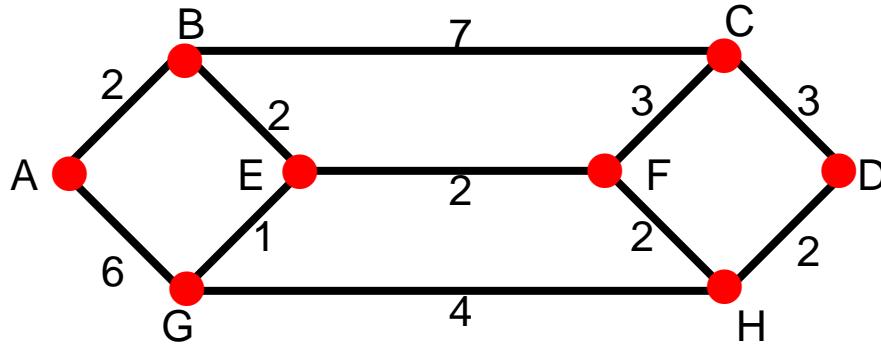
# Example: Dijkstra's Algorithm

A	0	-
B	2	A
E	4	B
G	5	E
F	6	E
H	8	F
C	9	B
D	10	H

- Shortest connection from A to any other node
- Greedy algorithm! Follows minima!



# Informal Description of Shortest Path Algo.



**origin** = starting point (e.g. point A)

**destination** = point that has a distance from the origin associated with itself (e.g. C, distance 9)

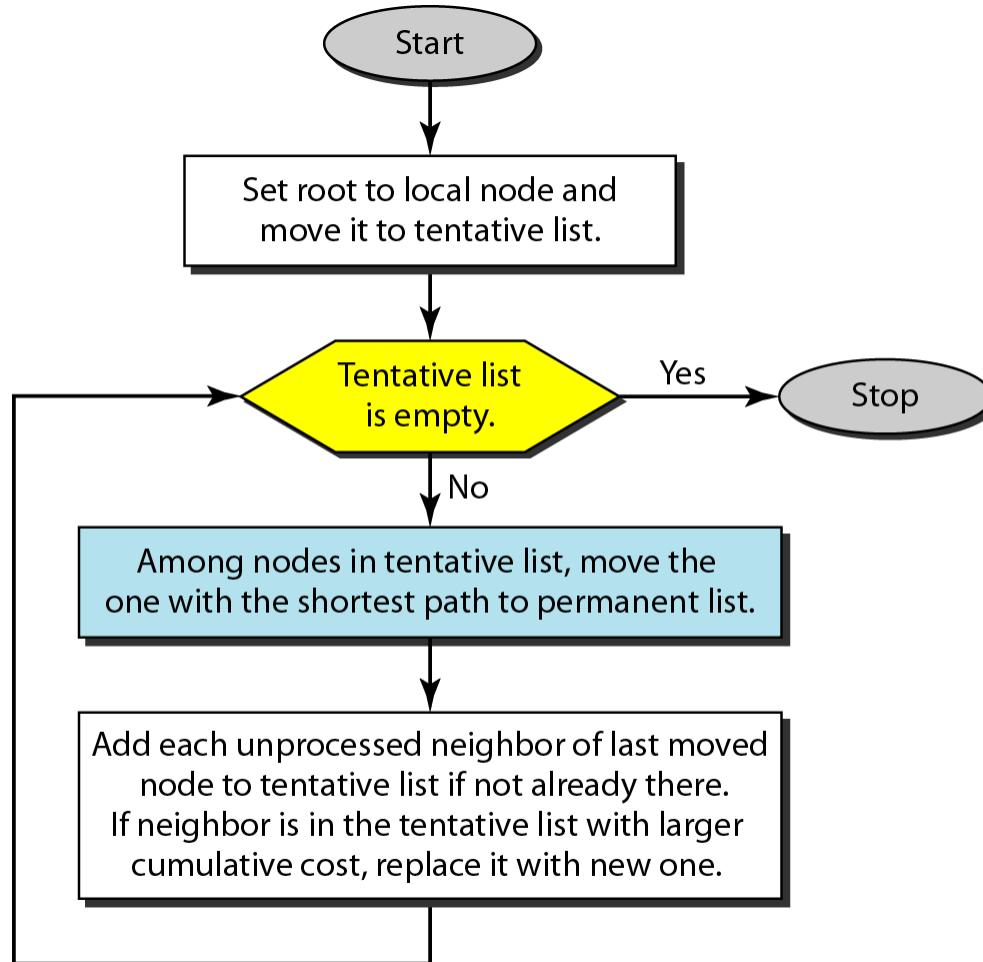
**neighbour** = destination that has a direct connection to another destination (e.g. B's neighbours A, E and C)

**table** = table of destinations with a known distance from the origin

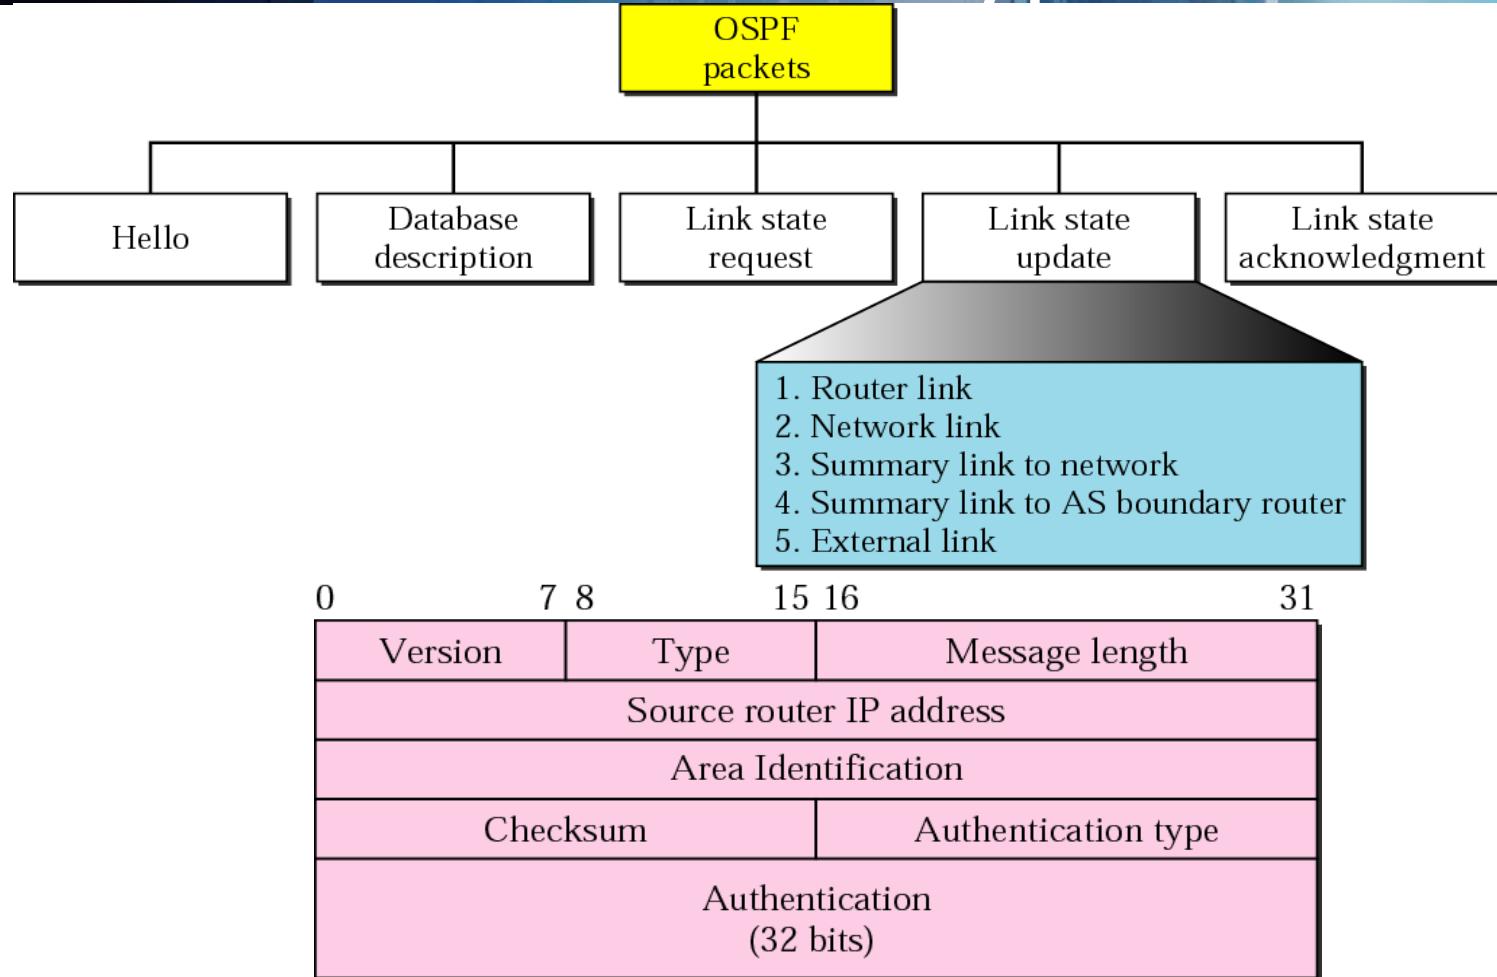
**bag** = bag of destinations to be resolved

# Informal Description of Shortest Path Algo.

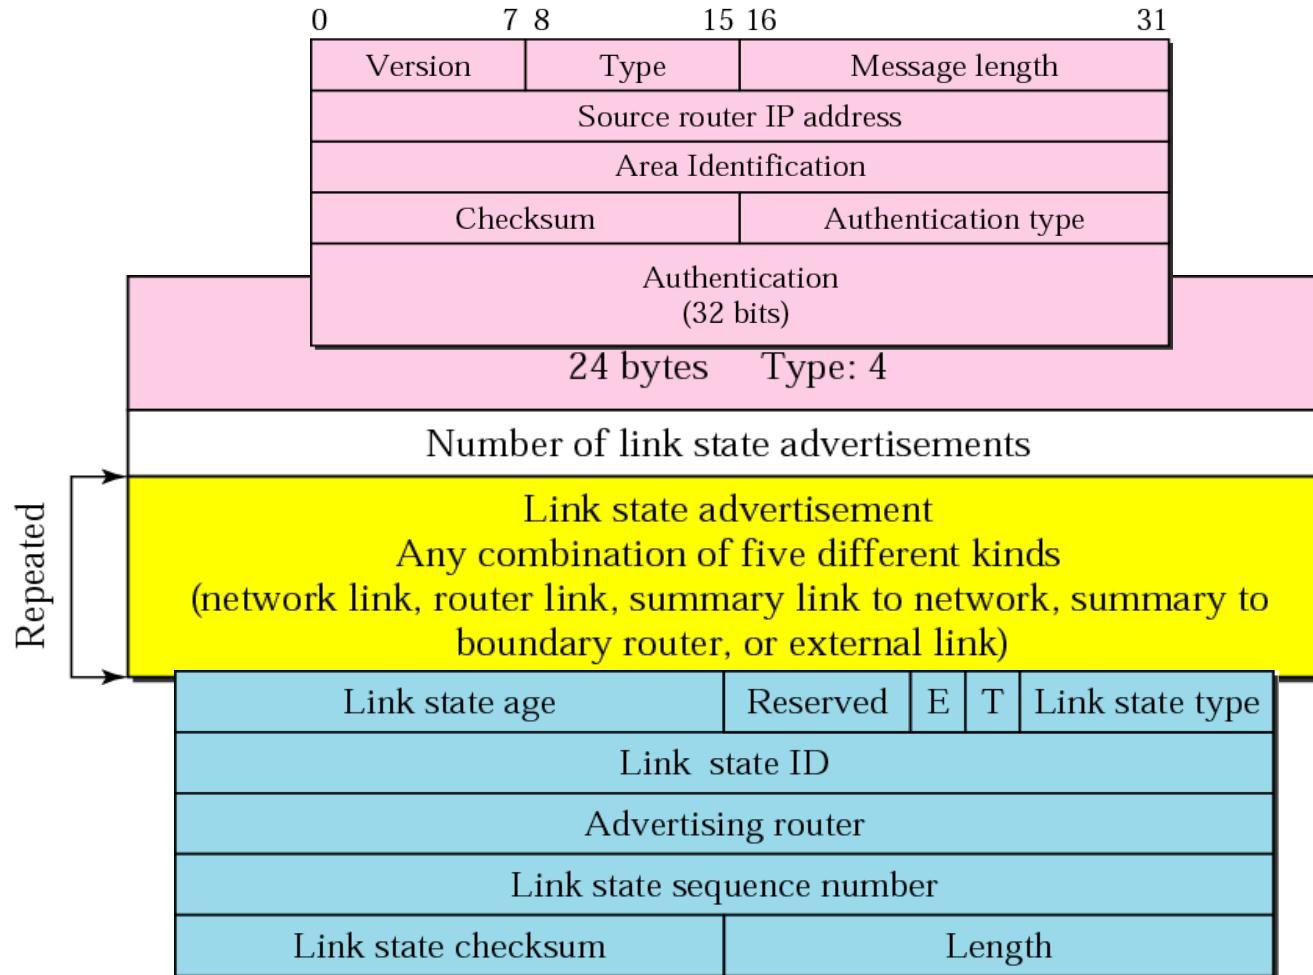
- Table and bag are empty at the beginning
- Put the origin as a destination with the distance 0 into the bag
- As long as there are destinations in the bag
  - Take the destination with the smallest distance out of the bag and put it into the table (if the distance is smaller than an already existing entry for the same destination)
  - Put all the neighbours of that destination with their distance to the origin (=distance to the destination + distance from the destination to the neighbour) into the bag



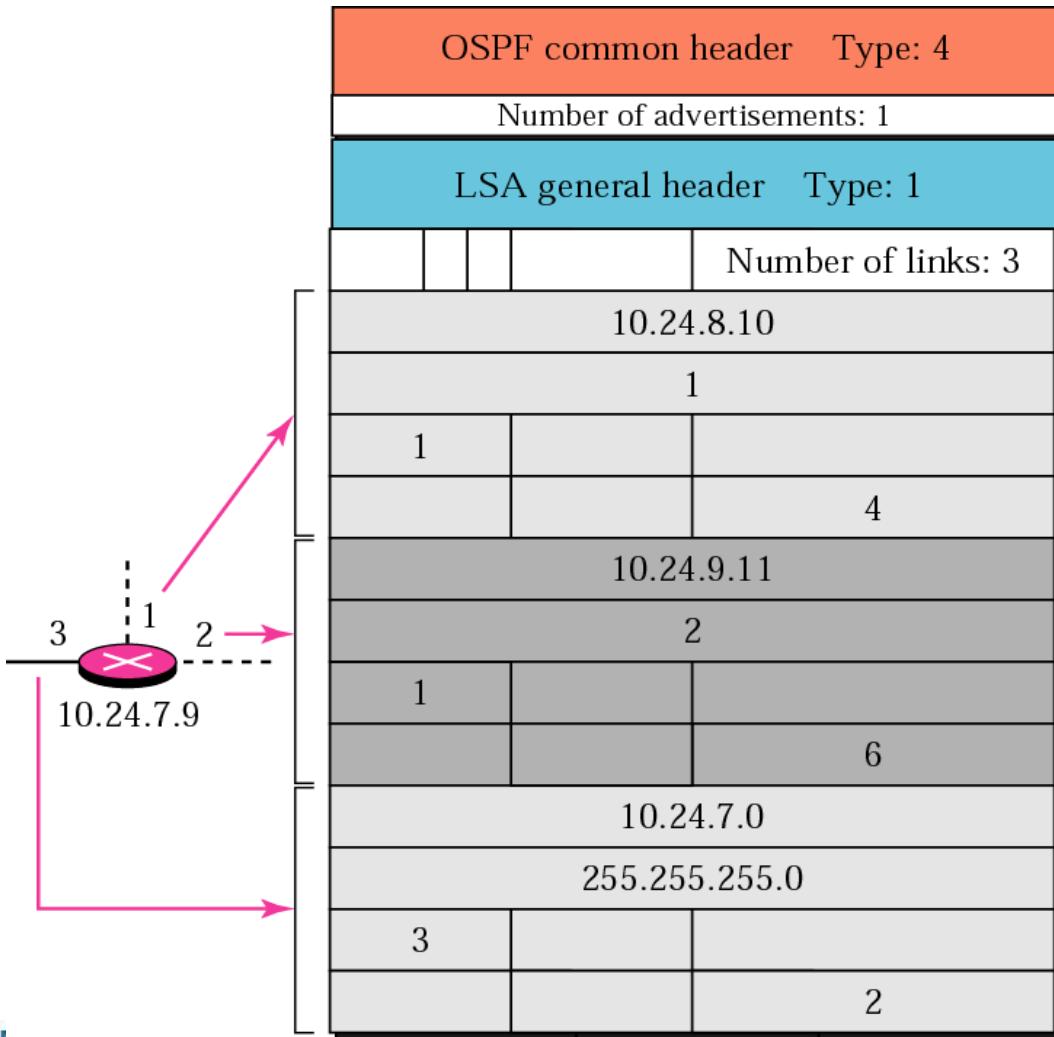
# OSPF Packet Types



# OSPF Link State Advertisement



# OSPF LSA Example

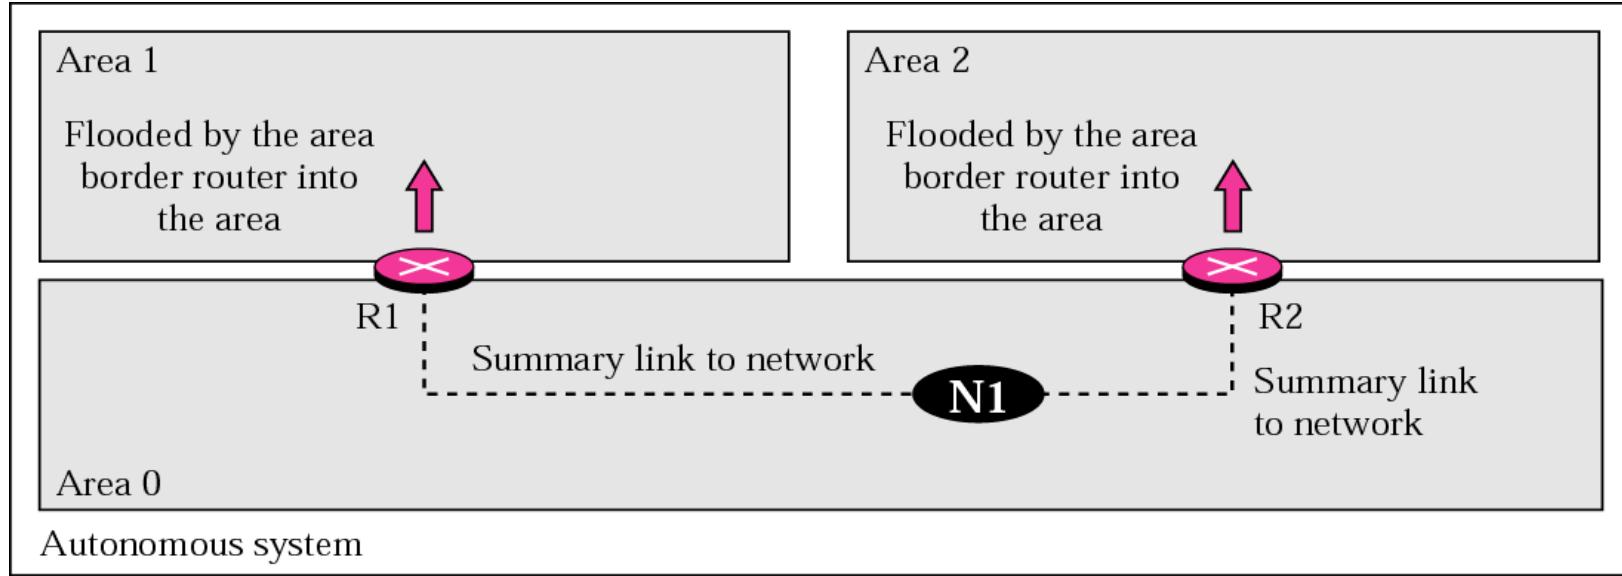


<i>Link Type</i>	<i>Link Identification</i>
Type 1: Point-to-point	Address of neighbor router
Type 2: Transient	Address of designated router
Type 3: Stub	Network address
Type 4: Virtual	Address of neighbor router

<i>Link Type</i>	<i>Link Data</i>
Type 1: Point-to-point	Interface number
Type 2: Transient	Router address
Type 3: Stub	Network mask
Type 4: Virtual	Router address

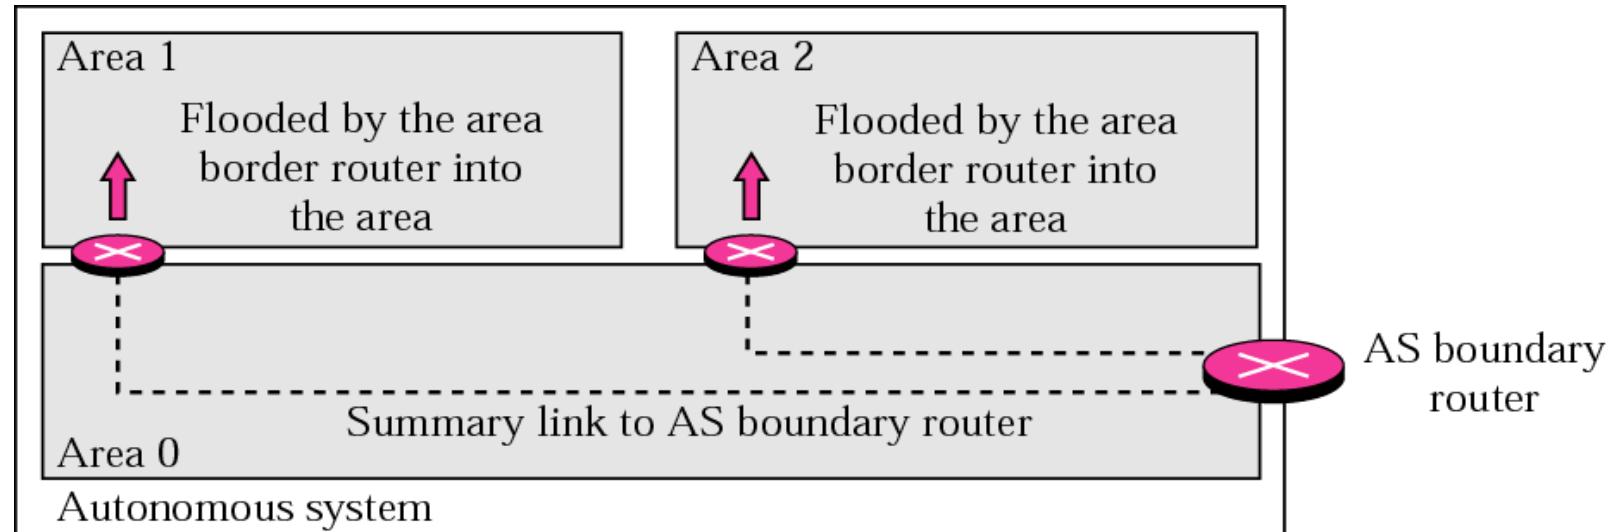
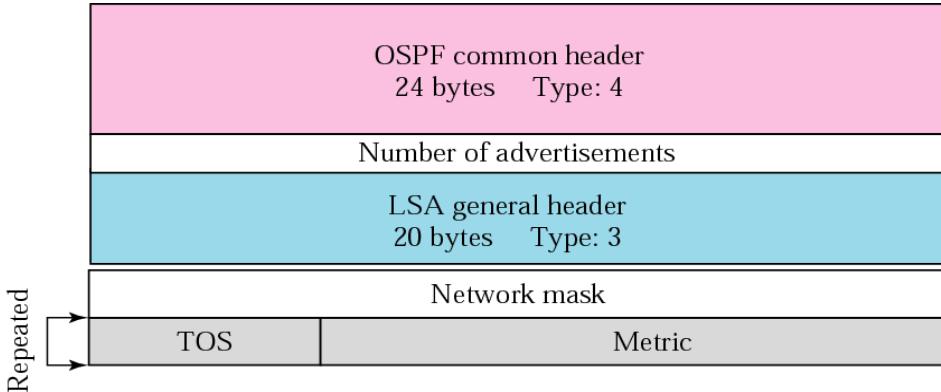
One router link advertisement

# Areas in OSPF



- Limits flooding of advertisements to areas

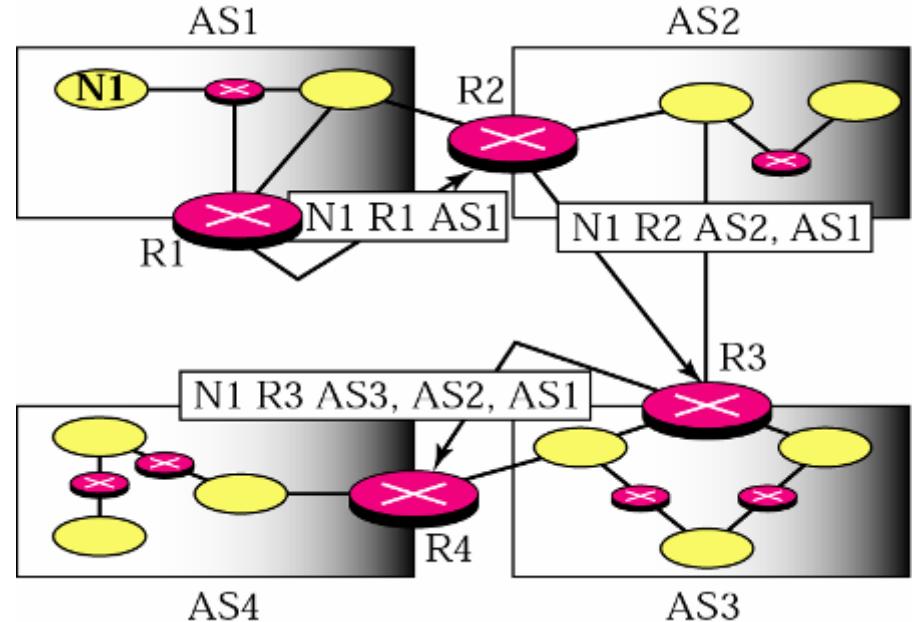
# OSPF Summary Link



# Border Gateway Protocol (BGP)

- Uses Path Vector Routing
- Advertisements include complete path to destination
- Router that forwards advertisement adds itself to the list
- Path can be checked for loops
- Policies are applied when incorporating new routes

Network	Next Router	Path
N01	R01	AS14, AS23, AS67
N02	R05	AS22, AS67, AS05, AS89
N03	R06	AS67, AS89, AS09, AS34
N04	R12	AS62, AS02, AS09



# Tables at Autonomous Systems

Dest. Path

A1	AS1
A2	AS1
A3	AS1
A4	AS1
A5	AS1

A1 Table

AS 1

Dest. Path

C1	AS3
C2	AS3
C3	AS3

C1 Table

AS 3

Dest. Path

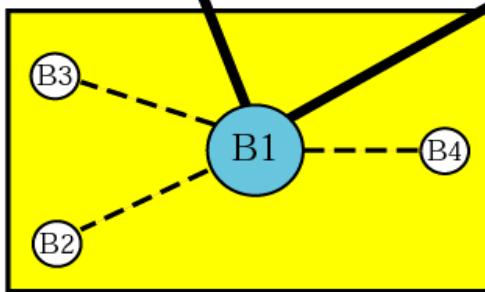
D1	AS4
D2	AS4
D3	AS4
D4	AS4

D1 Table

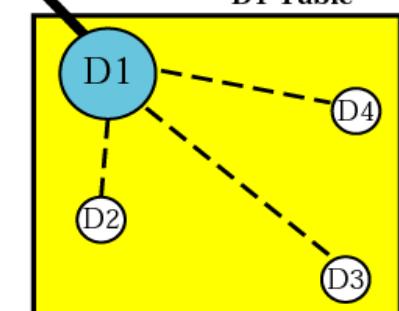
Dest. Path

B1	AS2
B2	AS2
B3	AS2
B4	AS2

B1 Table



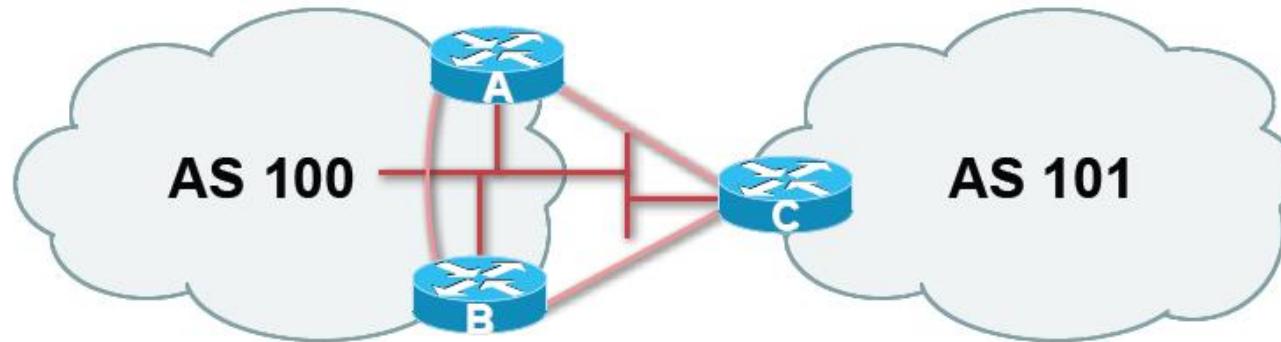
AS 2



AS 4



# BGP Example



Router A in AS100

```
interface ethernet 5/0
 ip address 102.102.10.2 255.255.255.240
!
router bgp 100
 Local ASN
 network 100.100.8.0 mask 255.255.252.0
 neighbor 102.102.10.1 remote-as 101
 Remote ASN
 neighbor 102.102.10.1 prefix-list RouterC in
 neighbor 102.102.10.1 prefix-list RouterC out
!
```

ip address of Router C  
ethernet interface

ip address on  
ethernet interface

Router C in AS101

```
interface ethernet 1/0/0
 ip address 102.102.10.1 255.255.255.240
!
router bgp 101
 Local ASN
 network 100.100.8.0 mask 255.255.252.0
 neighbor 102.102.10.2 remote-as 100
 Remote ASN
 neighbor 102.102.10.2 prefix-list RouterA in
 neighbor 102.102.10.2 prefix-list RouterA out
!
```

ip address on  
ethernet interface

ip address of Router A  
ethernet interface

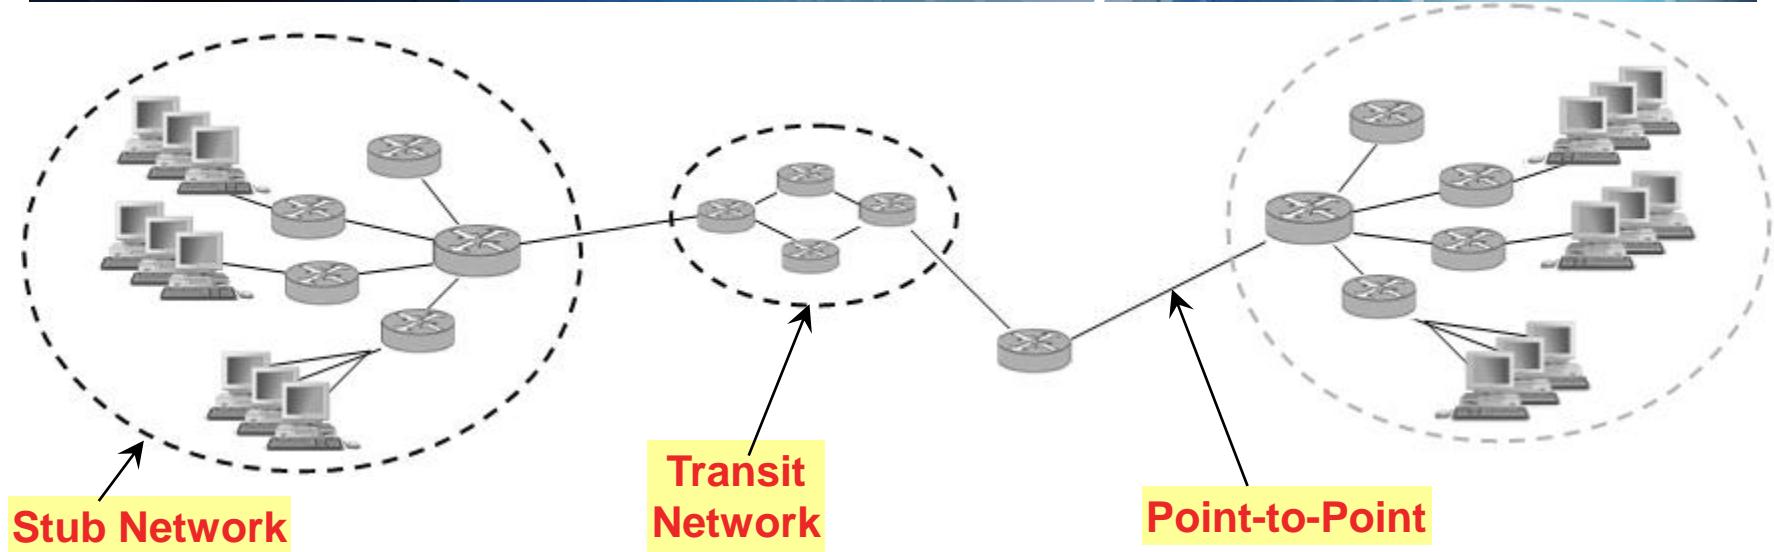
ip address on  
ethernet interface



# BGP-4: Border Gateway Protocol

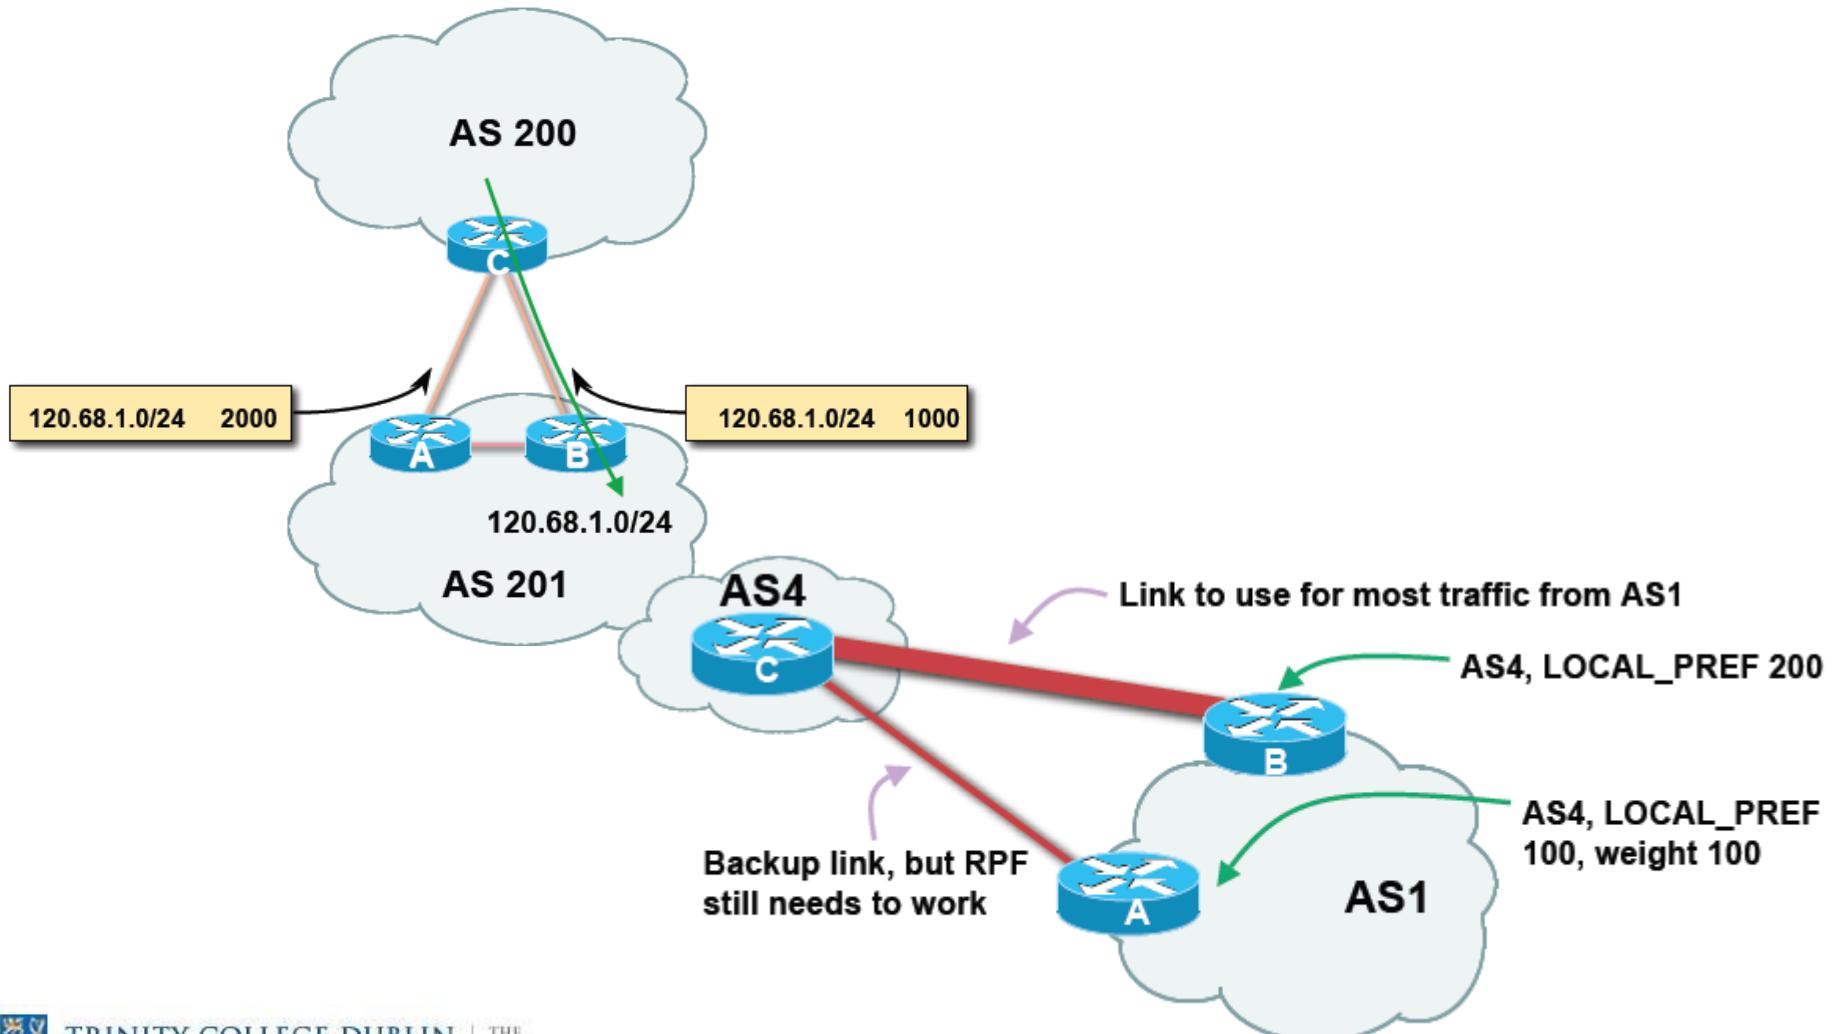
- AS Types
  - Stub AS: has a single connection to one other AS
    - Carries local traffic only
  - Multihomed AS: has connections to more than one AS
    - Refuses to carry transit traffic
  - Transit AS: has connections to more than one AS
    - Carries both transit and local traffic
- Each AS has:
  - One or more border routers
  - One BGP *speaker* that advertises:
    - Local networks
    - Other reachable networks (transit AS only)
    - Gives *path* information

# Autonomous Systems

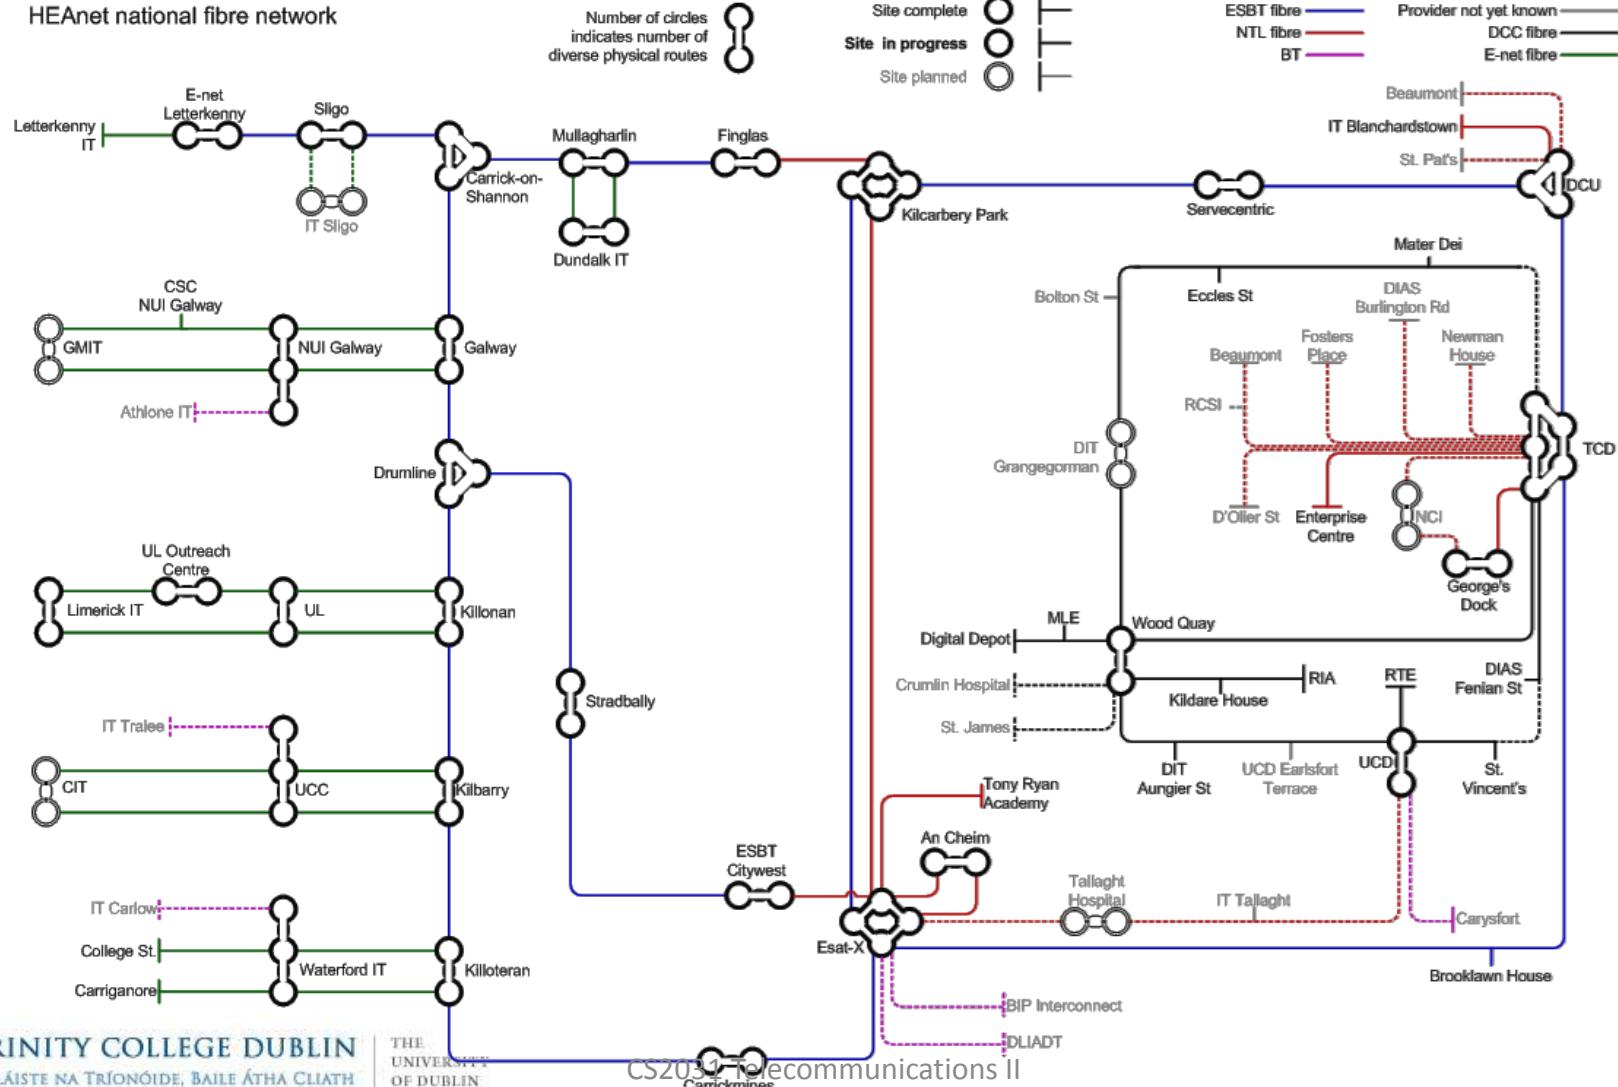


- Stub network
  - Network that does not forward to other network
- Transit network
  - Network that forwards traffic between other networks
- Point-to-point link

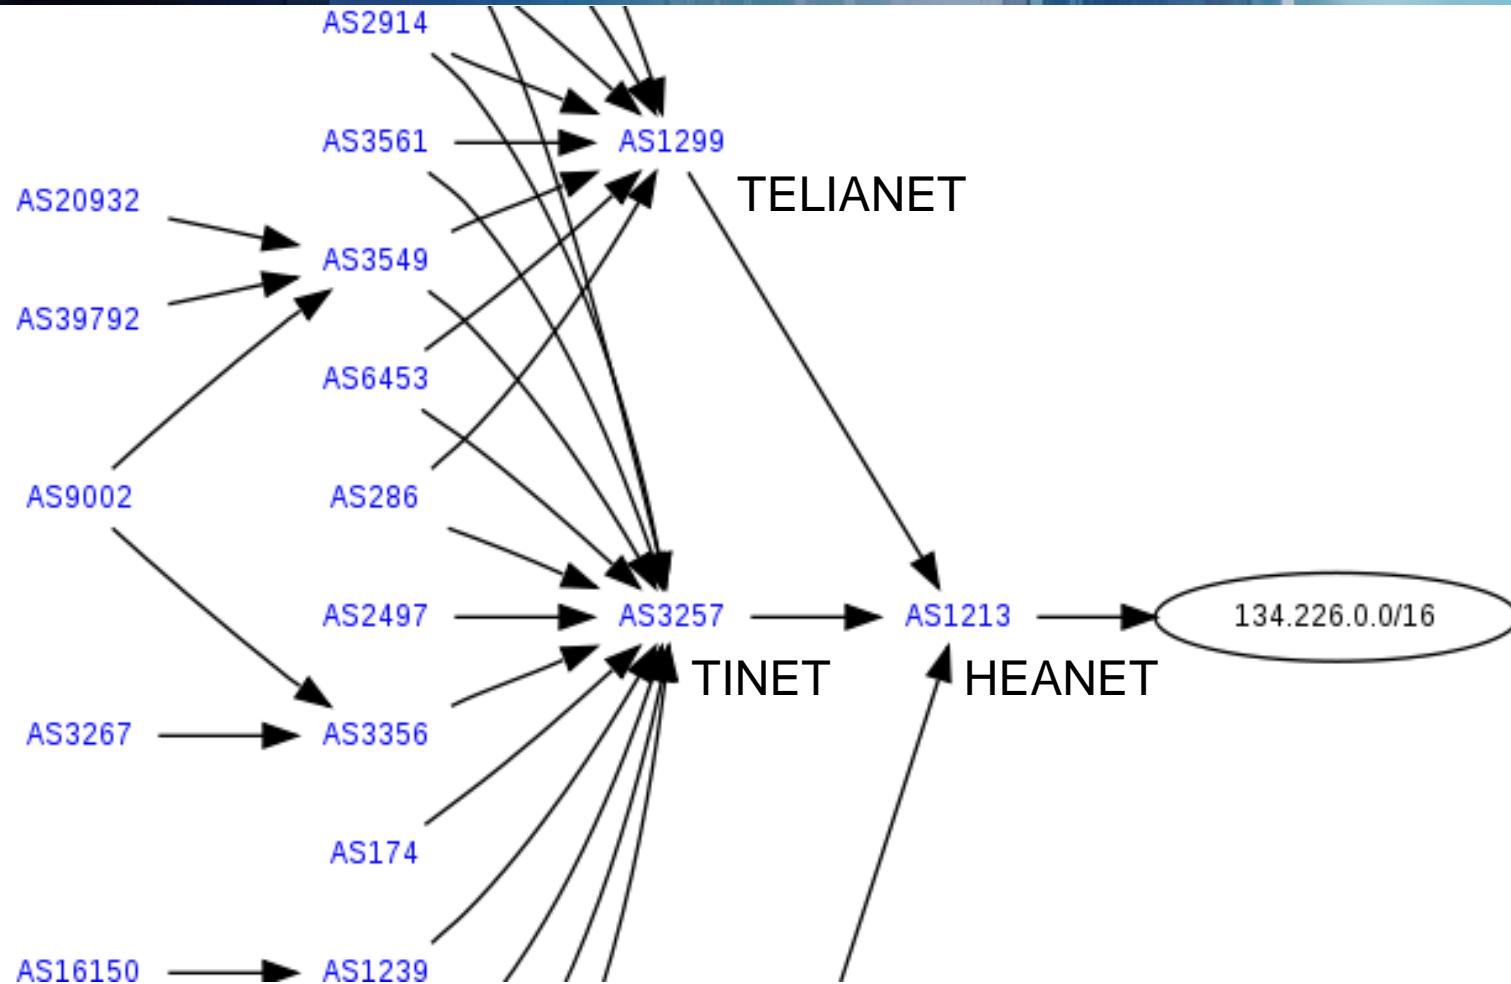
# Multiple Exits & Backup Links



# HEAnet Fibre Network

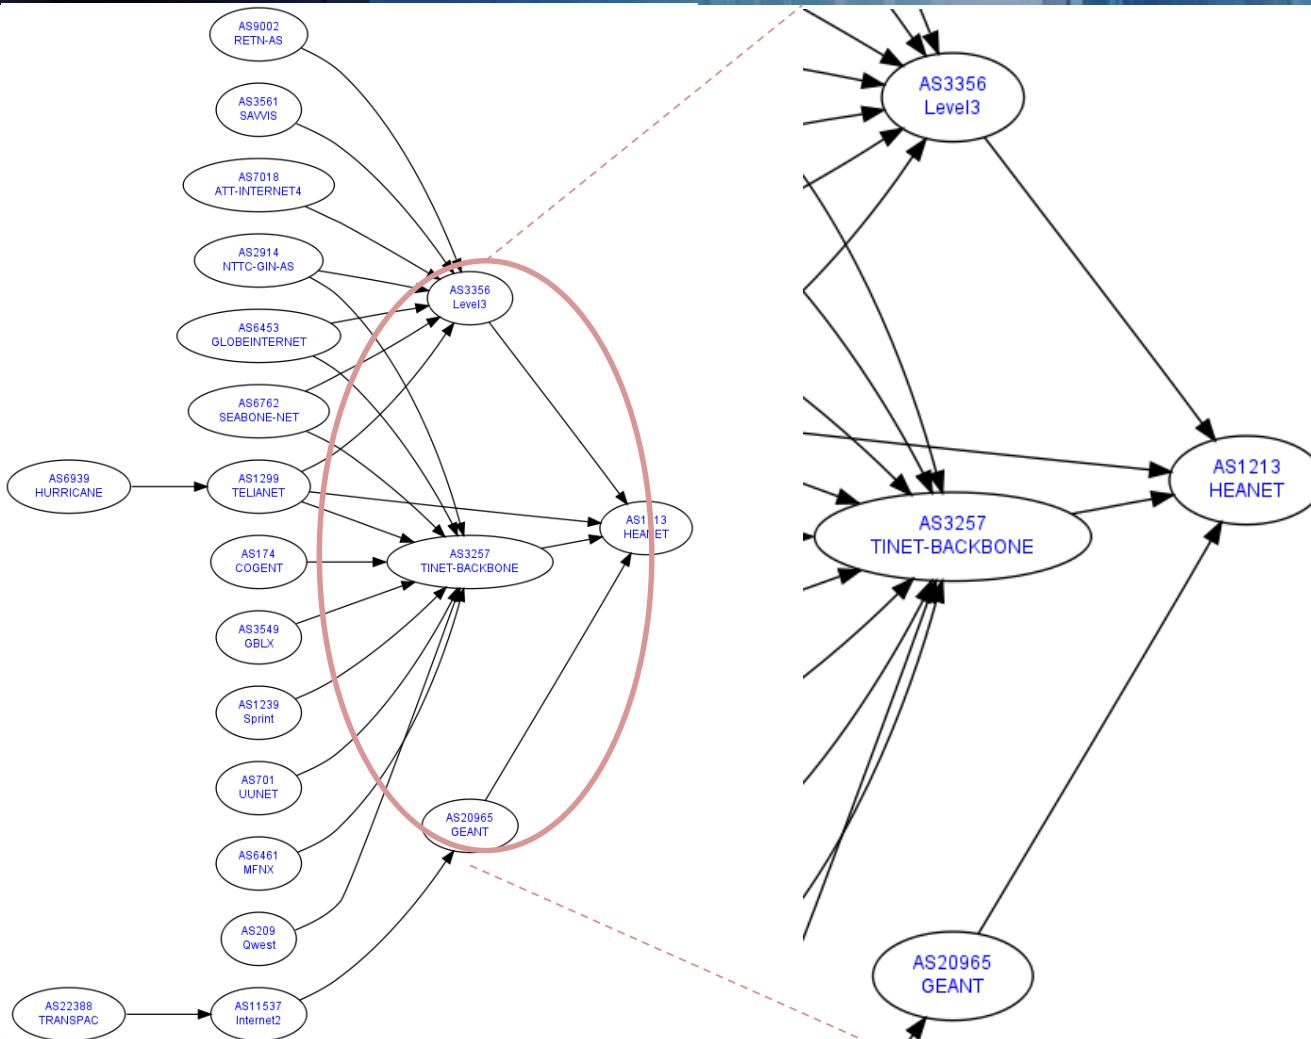


# AS1213 - HEANET



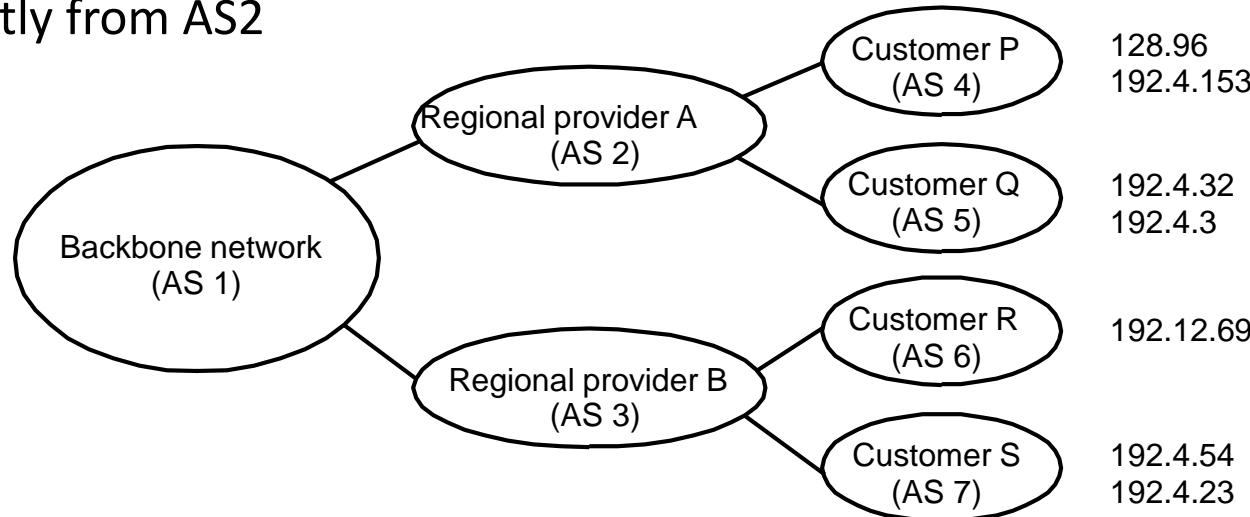
<http://www.robtex.com/route/134.226.0.0-16.html>

# AS1213 - HEANET



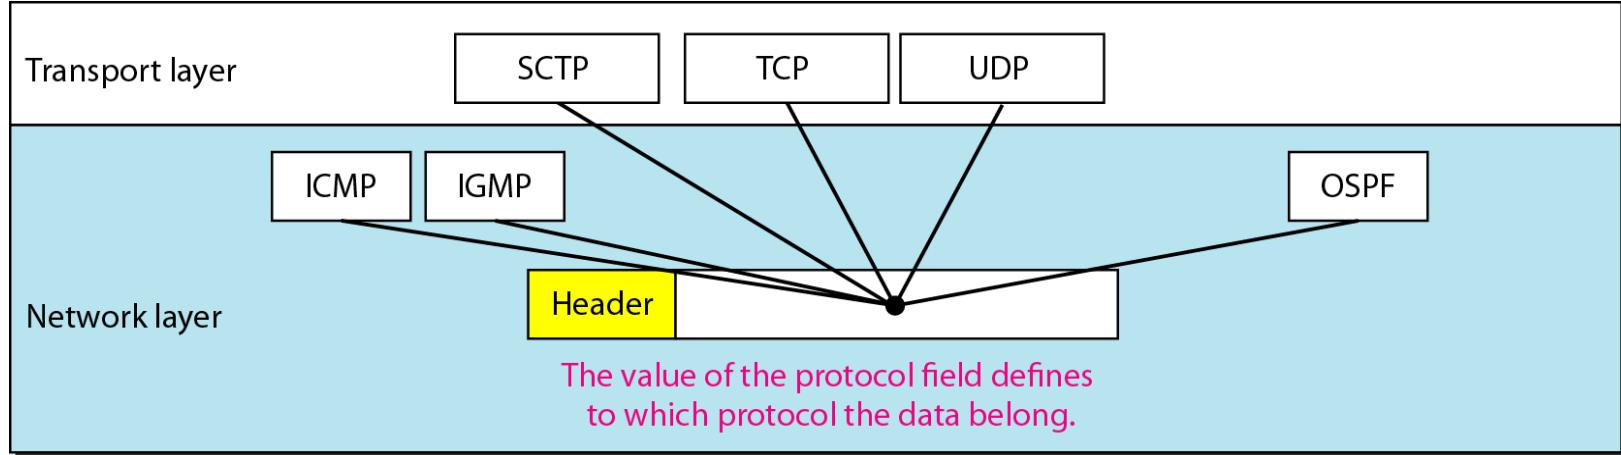
# BGP Example

- Speaker for AS2 advertises reachability to P and Q
  - network 128.96, 192.4.153, 192.4.32, and 192.4.3, can be reached directly from AS2



- Speaker for backbone advertises
  - networks 128.96, 192.4.153, 192.4.32, and 192.4.3 can be reached along the path (AS1, AS2).
- Speaker can cancel previously advertised paths

# Routing Protocols in the Stack



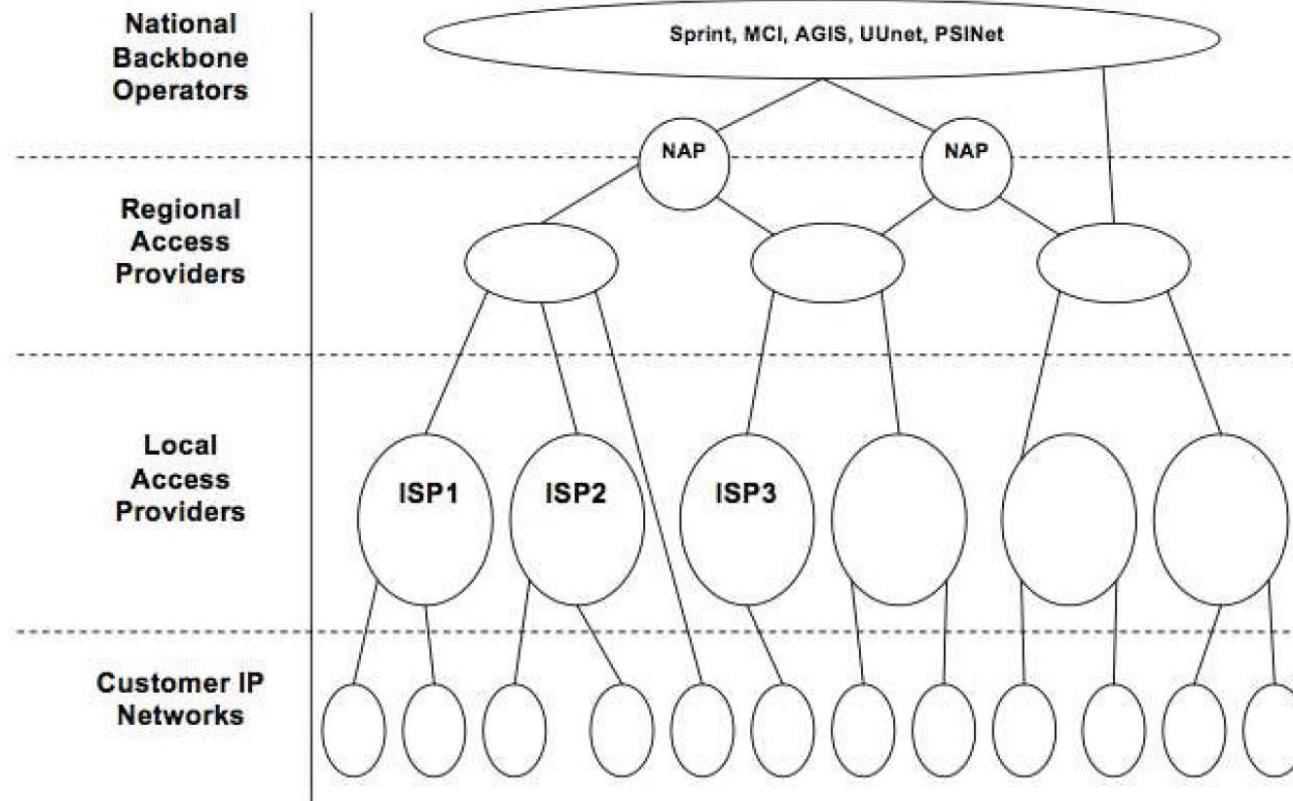
Protocol	Underlying Protocol	Protocol ID or Port
OSPF	IP	89
RIPv2	UDP	520
BGP	TCP	179

# Summary: Routing

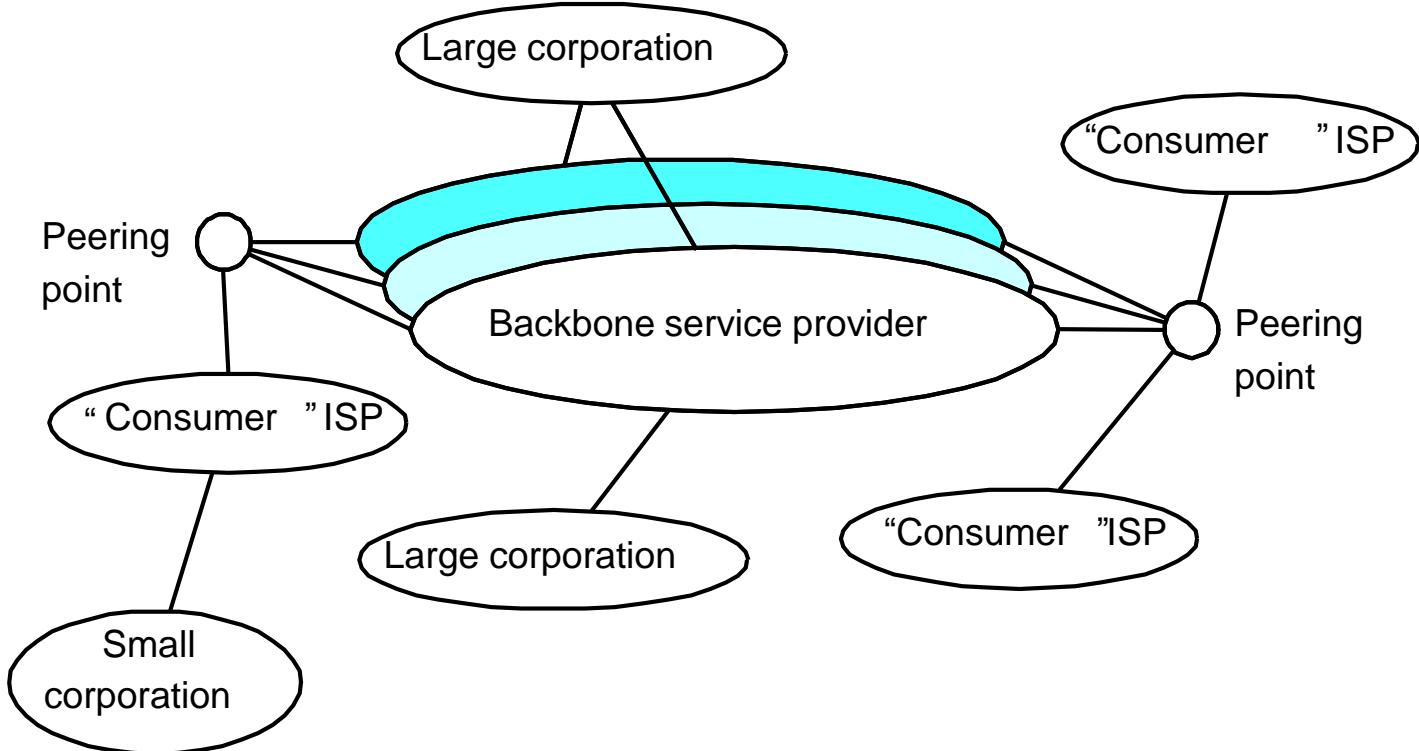
- Autonomous Systems
  - Stub network
  - Transient network
  - Point-to-point link
- Distance Vector routing
  - Share complete information with neighbours
  - Count-to-Infinity problem
  - Example: Routing Information Protocol (RIP)
- Link State routing
  - Share information about neighbours with everyone
  - Dijkstra's Shortest-Path Algorithm
  - Example: Open Shortest Path First (OSPF)



# Traditional Logical Internet Topology

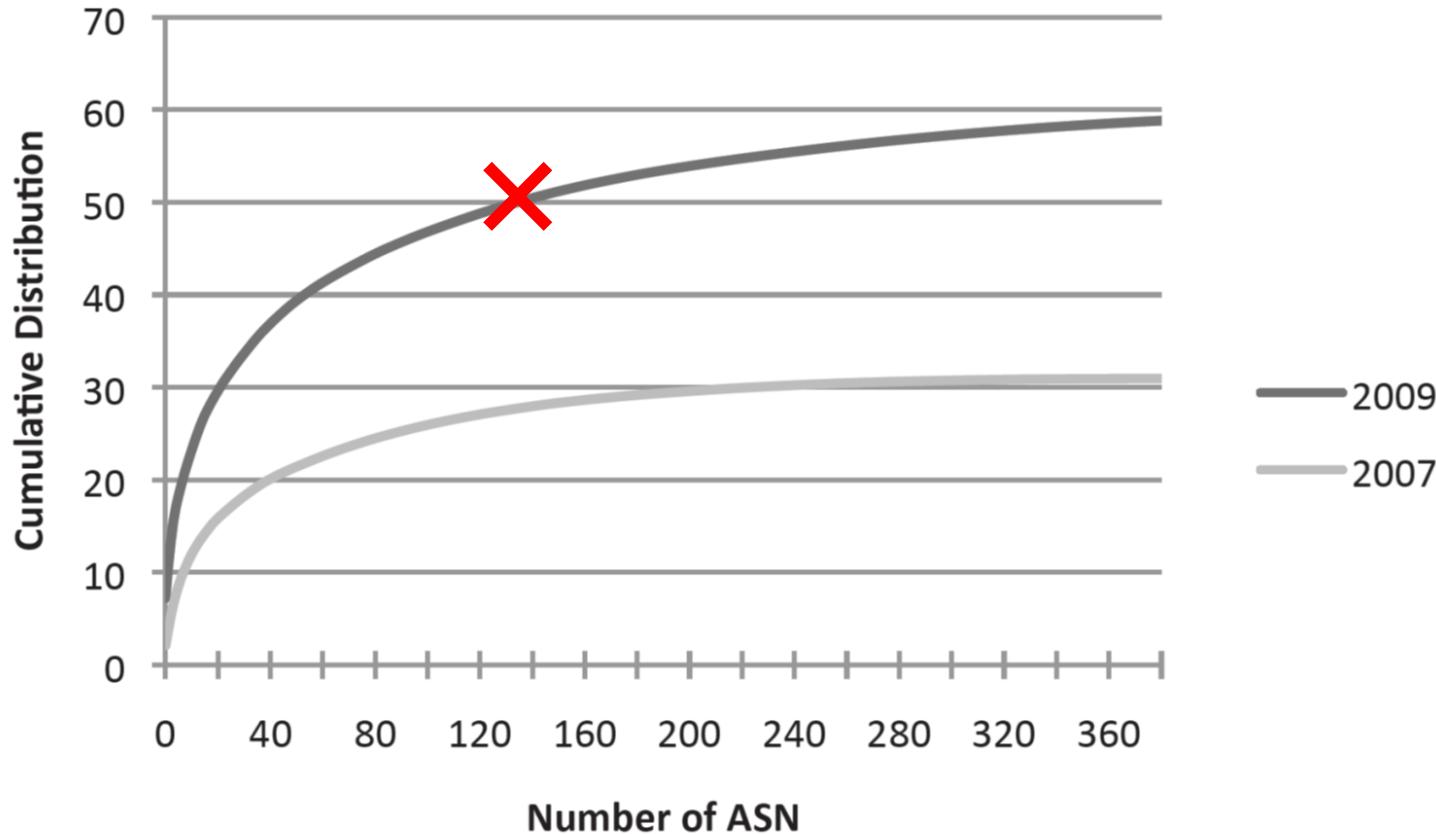


# Structure of the Internet



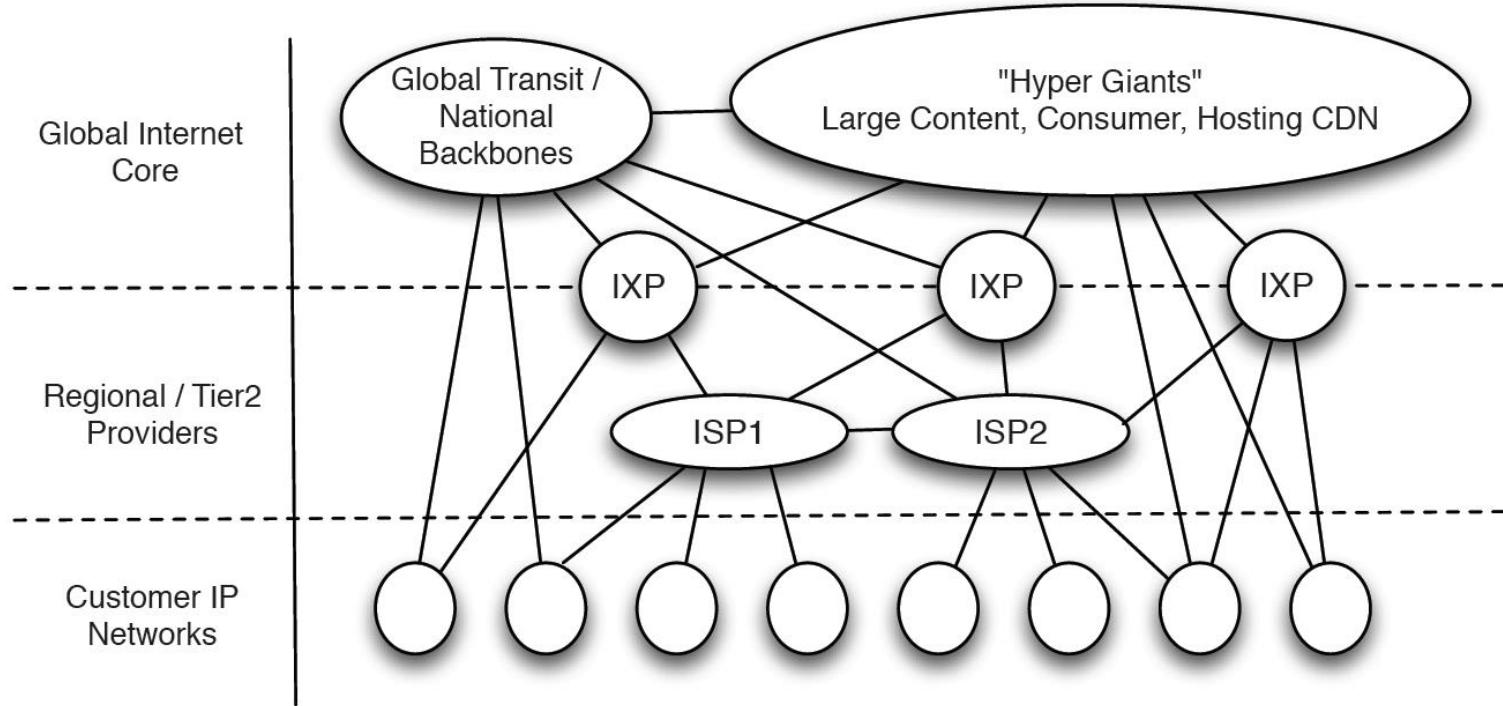
- Tier 1 provider eg. Verizon (WorldCom/MCI), Global Crossing
- Peering points – Exchange points between ISPs

# Cumulative Inter-Domain Traffic



150 ASN in 2009 – According to C. Labovitz now 30

# Emerging Logical Internet Topology



# Internet = Network of Networks

